

CAN LANGUAGE MODELS FORECAST INDUCTIVELY?

Anonymous authors

Paper under double-blind review

ABSTRACT

Judgmental forecasting requires inferring predictive relationships from sparse observations and contextual knowledge. We ask whether language models recover and use these relationships. Across prediction markets and synthetic worlds, models selectively update forecasts, infer relationships from context, and revise them as direct evidence accumulates. Larger tested models infer arbitrary conventions; intervening on one model’s internal representation changes its forecast. Training improves forecasts across domains, but an eight-seed replication weakens structural-transfer claims. A supervised pilot learns response-form selection in distribution where the tested RL recipe fails; its boundary is arbitrary labels, not a new domain. Accuracy gains on compositions of familiar mechanisms exceed a fixed prior, but a gain intervention does not establish transferred evidence use. Historical-market training lowers mean error on unseen event families. These findings support context-sensitive inductive forecasting while distinguishing in-domain learnability, transfer, and episode-specific evidence use.¹

1 INTRODUCTION

Judgmental forecasting is a demanding form of reasoning central to decisions about political, social, and economic futures. Unlike settings with abundant historical data for statistical extrapolation, judgmental forecasting requires integrating sparse, noisy evidence with contextual knowledge to form calibrated beliefs about future outcomes (Tetlock & Gardner, 2016; Mellers et al., 2014; Lawrence et al., 2006). This creates an inductive challenge: forecasters must infer from limited observations the relationships that govern how evidence impacts outcomes, and use those relationships to anticipate how events may unfold (Tenenbaum et al., 2006; Griffiths et al., 2010; Tenenbaum et al., 2011; Lake et al., 2017). Figure 1 illustrates this identification problem in a low-stakes setting. A researcher observes a carnival coin-flipping machine and, after seeing two heads, must estimate the probability that the next flip will be heads. A judgmental forecast may incorporate contextual knowledge (e.g., whether the game is rigged) to infer the process generating the observed outcomes.

Language models (LMs) now approach crowd accuracy in judgmental forecasting (Halawi et al., 2024; Schoenegger et al., 2024), but accurate predictions need not imply recovery of the underlying task structure. Models recover latent structure in some settings (Li et al., 2023), yet can predict accurately using relationships that fail to generalize in others (Vafa et al., 2025). This distinction matters when decision-makers ask how and whether a forecast should change under alternative evidence or events. Answering such questions requires predictive relationships that generalize across changes in context.

FORECASTING WITH A WORLD STRUCTURE




Figure 1: **Motivation.** Judgmental forecasts combine observations with contextual knowledge about the process that generated them. Frontier models, asked to forecast this question, average $P(H) \approx 0.68$ (see App. A).

¹ 🍌 Study dataset and artifacts 🌐 Code repository Links anonymized for review.

Present work. We study whether LMs exhibit inductive reasoning in forecasting tasks through a series of progressively more stringent behavioral tests spanning both real-world prediction markets and controlled synthetic environments. First, we evaluate whether LMs respond selectively to new evidence, distinguishing outcome-relevant information from topically related but orthogonal information (Exp. 1). Second, we test whether LMs infer predictive relationships from described context before statistical evidence accumulates. We construct synthetic forecasting problems in which the numerical observations are identical across conditions and only the described setting varies, isolating the contribution of background knowledge to the LM’s forecast (Exp. 2). Finally, we test whether relationships learned during fine-tuning transfer to new settings. We fine-tune LMs with reinforcement learning from verifiable rewards (RLVR) on both synthetic tasks and historical prediction market data, diagnose component failures with a supervised pilot, and evaluate on held-out structures and domains (Exp. 3, Exp. 4).

Results. All nine systems move $1.9\text{--}7.5\times$ farther for outcome-relevant than topical evidence, conditional on updating (Exp. 1). In Coin City, forecasts track context-selected relationships ($r = .52\text{--}.70$) and revise them as direct evidence arrives; larger models infer episode-local conventions, with a causal intervention linking the induced representation to Qwen3-14B’s forecast (Exp. 2). Training transfers across domains, but its advantage over an answer-matched shuffled control under joint domain and structure shift is uncertain at eight seeds (Qwen3-8B: .126, 95% CI $[-.002, .265]$). A supervised pilot learns response-form selection in distribution and, trained natively, fails under labels but not domain. A separate composition benchmark improves accuracy over a fixed prior and an answer-matched shuffled control, without clear evidence of transferred gain tracking (Exp. 3). On 318 held-out markets, training lowers mean error for every tested base and several checkpoints reach crowd performance, although the Qwen3-32B improvement is uncertain (Exp. 4).

2 WHAT DOES IT MEAN TO FORECAST INDUCTIVELY?

Forecasting is a useful testbed for inductive reasoning because it makes otherwise hidden choices observable. A forecaster must decide which evidence matters, which population provides the relevant comparison, how strongly that evidence should affect an outcome, and when direct observations should override an existing belief. Each of these choices implies a different numerical forecast. An accurate prediction alone does not tell us whether the forecaster made those choices for the right reasons.

This distinction matters because a forecast is rarely the last question. We often want to know what evidence should change the prediction, whether the same evidence should matter differently in another population, and what new observation should trigger revision. A model that reaches the right number through a familiar base rate, a topical cue, or some other surface regularity may still fail under exactly these changes. Average accuracy can therefore conceal an important weakness: two models can give the same forecast while relying on very different relationships between evidence and outcomes.

We use *inductive forecasting* to describe the stronger behavior we care about here. An inductive forecaster should not only identify whether evidence is relevant, but also infer the relationship that determines how that evidence should affect the outcome. When context changes which comparison is appropriate, the forecast should change with it. When direct evidence becomes available, that evidence should be able to revise the earlier contextual inference. And when the same underlying relationship appears in an unfamiliar setting, the forecaster should be able to use it without depending on the original surface form.

These requirements extend beyond forecasting. A persuasion model should respond not only to a message, but to who delivers it. A coordination model should distinguish shared incentives from what participants know about one another. A policy model should distinguish similar language from the institutional pathway through which an intervention takes effect. In each case, the conditional relationship is often the object of interest: who trusts whom, which population responds to which signal, or which mechanism connects an action to an outcome. Matching an average outcome while missing these dependencies can produce a convincing static prediction and a misleading response when the setting changes (Park et al., 2023; Aher et al., 2023; Argyle et al., 2023; Bisbee et al., 2024).

This suggests a stronger way to evaluate forecasting systems. Rather than asking only whether a model reaches the correct number, we can ask whether its forecast changes when the relevant




Figure 2: **When should a forecast change?** An agent forecasts an unresolved question, then receives one fictional update with search off. We ask whether it moves when the news matters, and whether it moves the right way.

evidence, comparison, or relationship changes. We therefore proceed through increasingly demanding tests: first whether models distinguish relevant evidence from topical information, then whether they infer how strongly evidence should matter from context, and finally whether those relationships can be revised and carried into new settings.

3 DO FORECASTS CHANGE FOR THE RIGHT REASON?

A key forecasting skill is to appropriately incorporate new evidence as it arrives; for example, classic work on judgmental forecasting identifies repeated belief revision as a hallmark of strong human forecasters (Tetlock, 2005). We argue that effective updating requires a forecaster to determine: (1) which evidence is relevant in predicting an outcome and (2) the extent to which the new evidence should affect the forecast. We assess these capacities in two experiments. First, we pair real-world prediction-market questions with controlled synthetic evidence, allowing us to test whether models update selectively and in the direction the evidence warrants. Second, we use a controlled synthetic environment in which the same evidence has different predictive effects across contexts, allowing us to test whether models can infer the relationship that determines how strongly the evidence should affect their forecast.

3.1 EXP. 1: CAN MODELS DIRECTIONALLY UPDATE?

Design. We ask nine models to forecast 100 unresolved binary Polymarket questions with web search, then to revise each forecast after one fictional news update with search disabled (Fig. 2). Updates support the outcome, oppose it, or stay on topic without bearing on it; these labels are hidden from the model. To remove lexical shortcuts by construction rather than by measurement, we also froze 80 finite-rule instances whose evidence, prior and named actor are identical across three stated relations, with three distinct, exactly computable posteriors (App. B).

Finding: models update in the implied direction. Across 30,094 valid revisions, pro-outcome evidence raises forecasts, anti-outcome evidence lowers them, and orthogonal topical evidence changes little (Fig. 3a). Conditional on moving, models move $1.9\text{--}7.5\times$ farther for directional updates (Fig. 3b). Uncertainty and robustness analyses are in App. B.3.4; independent human review covers the original packets, not the frozen cohort.

Finding: lexical cues do not explain the result. Bag-of-words classifiers separate directional from orthogonal evidence almost perfectly (.989) but are far worse at telling support from opposition (.688), where models reach .887–.990. In the frozen cohort a predictor ignoring the stated relation answers identically within a triple, so it can win no triple at all: four such predictors, one allowed the best constant revision while knowing every answer, score exactly 320/960 items and 0/320 triples, while GPT-5.6 reverses 48/48 triples and matches the exact posterior on 144/144 items (App. B.8).

Finding: it depends on inference-time reasoning. Holding weights and prompts fixed, Qwen3-32B wins 29/320 triples with reasoning disabled and 297/320 with it enabled (1/960 against 811/960 items), and both arms see identical items, so the gap is not difficulty or leakage. Counting every planned call, Claude Opus 5 declined 84 of 576, blocking 153 dependent updates, and retains 17/48 triples once those count as failures (App. B.8).

Takeaway. Lexical cues identify original-packet relevance (.989 accuracy). A frozen follow-up uses identical news across relational contexts in 80 finite-rule instances spanning eight mechanism classes. Qwen3-32B reverses both directions in 14/80 pairs without thinking and 77/80 with thinking; GPT-5.6 reaches 24/24 on a fixed subset, with 0.0-point broken-link movement (base wording; all planned failures count). Wording, control drift, and failure analyses appear in App. B.8. This formally checked diagnostic has no new human validation and does not establish natural-news coverage. Experiment 2 tests contextual relationship selection under a controlled generative process.

3.2 EXP. 2: CAN MODELS INFER HOW STRONGLY EVIDENCE SHOULD MATTER?

Design. Experiment 1 tests which evidence warrants an update. Forecasting also requires inferring how strongly that evidence should matter. In the opening coin-flip example (Fig. 1), the same observation warrants different predictions depending on the process believed to have generated it.

We systematize this ambiguity in *Coin City*: the same campaign news moves polls strongly in one reference city and weakly in another. Context identifies the better analogue for target City C, while zero to four City C cases provide direct evidence that can confirm or override it. Each city’s polling response follows

$$\Delta p_{ij} = g_j x_{ij} + \epsilon_{ij}, \quad \epsilon_{ij} \sim \mathcal{N}(0, 5^2), \quad (1)$$

where x_{ij} is signed campaign news and g_j is the city’s stable sensitivity, which is never stated: the model must infer it from examples.

We compare target cases alone, reference cases without target context, correct context, and opposite context, holding numerical evidence fixed across the paired arms (Fig. 4). A fifth arm replaces familiar descriptions with episode-varying labels whose meaning must be inferred from the references. Forecast error tests whether context helps; the forecast-implied response to news tests which reference the model selects (App. C.5).

Finding: the forecast follows the contextual clue. Figure 5 shows that, when City C provides few direct observations, context guides which reference city’s response is used to forecast. The correct context lowers forecast error, while the wrong analogue increases error. Moreover, the effect implies the model is quantifying the importance of the contextual clue: after controlling for the underlying response regime, the forecast-implied slope correlates $r = .51-.86$ with the fitted slope of the cue-selected reference city, while its correlation with the other reference remains near zero. As City C observations accumulate, context matters less, as forecasts increasingly rely on direct local evidence.




Figure 3: **Models move much more for evidence than for topical chatter, and usually move in the implied direction.** (a) Average revision after supporting, contrary, and topical updates. (b) How much farther models move for directional evidence, conditional on moving (Sens.), and how often a move of at least three points has the right sign (EHC $_{\geq 3}$). Abstention is decomposed separately in Table 9, which disagrees with movement about which deployments are selective; intervals are in App. B.3.4.

Shared prompt in every arm

Each row is a separate polling case rather than a time series. Positive net news favors the candidate and negative net news harms the candidate. Poll change is the ending poll minus the starting poll. Within a city, typical responsiveness to net news is stable across cases, although individual end polls remain noisy. Rows with the same Pair label have equal-sized news in opposite directions.

Query: A new City C case begins with a poll of 53.9 and has net news +5. What poll do you predict at the end of this new case? Think carefully about the expected poll, but do not provide step-by-step working. Respond with one JSON object only: {"rationale": "one short sentence", "predicted_poll": <number from 0 to 100>}

Illustrated arm: References + target context**CITY A**

Background: Residents generally encounter campaign developments through **national news coverage**, where campaign stories receive sustained attention and tend to produce larger immediate poll movements.

Pair	Start	News	Delta	End
1	45.4	-8	-6.5	38.9
1	45.4	+8	+12.2	57.6
2	55.1	-6	+0.9	56.0
2	55.1	+6	+11.4	66.5

CITY B

Background: Residents generally encounter campaign developments through **local news coverage**, where campaign stories compete with local issues and tend to produce smaller immediate poll movements.

Pair	Start	News	Delta	End
1	35.0	-10	-5.6	29.4
1	35.0	+10	+9.3	44.3
2	51.8	-10	-5.5	46.3
2	51.8	+10	+7.2	59.0

CITY C - TARGET CONTEXT

Background: City C residents generally encounter campaign developments through **national news coverage**.

Pair	Start	News	Delta	End
1	57.8	+8	+9.4	67.2
1	57.8	-8	-1.8	56.0
2	47.4	+10	+8.1	55.5
2	47.4	-10	-10.6	36.8

Figure 4: **Coin City turns contextual reasoning into a paired comparison.** Every arm shares the query and numerical cases. The context arm adds a description connecting City C to one displayed media environment; the opposite-context arm changes only that sentence. Target cases then accumulate, allowing local evidence to confirm or correct the contextual choice.




Figure 5: **Context improves forecasts before target cases and helps assign the response pattern across six systems.** Giving the correct analogue lowers error before City C has a history; naming the wrong analogue raises it. As direct City C cases accumulate, forecasts rely less on either contextual cue. Full results by model are in App. C.5.

3.2.1 INFER A CONVENTION WITHOUT STABLE SEMANTICS

However, a different explanation for the relational mapping above is that the model may already associate familiar cues, like “national news,” with a stronger response and “local news” with a weaker one, rather than actually inferring the mapping from the reference cases. We remove this shortcut by replacing these cues with the arbitrary labels KIV and ZOR. Their meanings are randomized across episodes, so neither label has any stable association with strong or weak responsiveness. Twelve of fifteen tested systems recover the episode-local mapping before any target cases, including all tested local checkpoints above 8B (App. C.5.3), although recovery does not always improve forecast error. When we further challenge this inference by making the two response regimes harder to distinguish, recovery weakens, but sufficiently large models are still able to recover the relationship, with the threshold varying by model family (App. C.8).




Figure 6: **Larger models infer an episode-local convention, then yield to direct evidence.** On 88 new episodes, four larger Qwen and Llama checkpoints (14–72B) encode the changing KIV/ZOR mapping; three smaller checkpoints (4–8B) do not reliably encode it. The decoded state follows context before City C has a history and follows the target relationship as local cases accumulate.

This pattern matches what we would expect if models are inferring a latent mapping: doing so becomes harder as the underlying regimes become less distinguishable. *But when a model does recover that mapping, what has it actually inferred?* We ask whether the episode-local KIV/ZOR relationship can be decoded from the model’s hidden states. In larger open models, the decoded representation follows the context-selected relationship before City C has its own history, and then shifts toward City C’s observed behavior as direct evidence accumulates (Fig. 6).

But decoding still leaves open whether this representation actually matters for the forecast. Figure 7 tests this directly by swapping the label-state representation between otherwise matched Qwen3-14B prompts. Before City C provides any direct observations, the intervention changes both the analogue the model names and its numerical forecast. After four City C cases, the same swap has no reliable effect. Across episodes, the intervention produces a regime-aligned symmetric forecast shift of $+0.333$ before target cases and -0.035 after four; the latter falls within the sign-flip null range. For Qwen3-14B, this shows that the induced mapping helps determine the forecast when target evidence is sparse, but gives way once direct observations reveal City C’s own response relationship (App. C.7.7).

Takeaway. Models use context to select a predictive relationship and revise that selection from local evidence. Arbitrary labels remove a familiar-semantic shortcut; decoding reveals an episode-local representation, and intervention links it to Qwen3-14B’s forecast when target evidence is absent. Experiment 3 asks which parts of this behavior training carries into new domains and response structures.

4 CAN RELATIONSHIP USE SURVIVE A CHANGE OF WORLD?

Finetuning may teach evidence use or only typical answers. We test whether training transfers relationships to new worlds and improves unseen-event forecasts.

4.1 EXP. 3: WHAT TRAINING TRANSFERS, AND WHAT IT DOES NOT

Why matching matters. The parent experiment compares byte-identical prompts, identical target-answer multisets, and equal training budgets, changing only case–answer pairing. Figure 8 crosses a domain change—campaign news and candidate support become shipping bulletins and cargo flow—with a change from direct to persistent responses (App. F).

Finding: domain transfer is stronger than structural transfer. Training generalizes across the surface change. Across both arms and 8 seeds, Qwen3-




Figure 8: **Transfer grid.** Training uses shading; all four are evaluated. Columns vary domain; rows vary response mechanism.




Figure 7: **Changing the induced convention changes the forecast when evidence is sparse.** In the illustrated pair, swapping the internal representation moves the forecast from 51.3 to 54.0 and changes which reference city the model names; reversing the swap reverses the movement. As more City C behavior is revealed, the same label-state swap has no reliable effect. Full controls in Table 51.

8B reduces domain-shift response error by 5.360 points [5.266,5.454], with every seed agreeing. The matched arm reaches 1.89 versus the base’s 7.30. Because unparsed draws take the full range penalty, we report base contrasts twice: the joint-cell gain is 1.129 [1.063,1.196] fail-closed and .677 [.611,.744] on parsed draws.




Figure 9: **Qwen3-8B Coin City transfer at round 300 under five-draw stochastic decoding, eight training seeds.** Bars show task-averaged response MAE over four evidence depths; trained arms average eight seeds. Each cell prints the paired seed-first prior-minus-matched contrast, bold when its 95% interval excludes zero; at eight seeds only the in-distribution cell does. The structure-shift row is a control, not a transfer test: training contains no structural variation, so the trained policy emits no delayed response anywhere (Section 4.1).

The finer contrast between episode-matched and shuffled-pair training does not survive replication. Estimated at three seeds and extended to eight under a pre-registered amendment, the Qwen3-8B joint-shift contrast falls from .251 [.058,.470] to .126 [-.002,.265]; all four of Qwen3-4B’s cells excluded zero at three seeds and all span it at eight. A wrapper crossover retains a similar three-seed mean (.251 versus .255; App. D.1.6). One cell survives every test: Qwen3-8B in distribution, .060 [.018,.112], 8/8 seeds, clearing the conservative t [.005,.115], a wild cluster bootstrap ($p = .0089$) and an exact sign test ($p = .0039$). It does not replicate on Qwen3-4B, and Llama-3.1-8B’s estimate has the opposite sign.

Response-form selection is learnable, but transfer is limited. The original structure-shift row is a control: training contains only the direct response, and every trained checkpoint emits zero delayed response. We therefore rebuilt the task with both forms in training: two references have identical horizon-1 responses but differ in persistence, and a cue names which the target follows. Three-seed Qwen3-8B RL learns marginal persistence without reliable cue following (App. D.1.7).

In a fixed component diagnostic crossing both domains and semantic or arbitrary labels, all eight checkpoints identify the named reference perfectly. Supplying exact response patterns does not rescue the 8B forecasts; an untrained 32B model shows partial numerical ability but often omits the resting baseline (App. D.1.8). An exploratory supervised pilot on 4,800 prompts yields 97.9–100.0% structure accuracy in Coin City/semantic at four seeds versus 50.0% for matched RL, falling to chance under arbitrary labels or in Coin Harbor. Trained natively it reaches 100.0% in Harbor/semantic yet stays at chance under arbitrary labels: the barrier is the label mapping, not the domain. Rate and budget differ, limiting this to in-domain learnability, not a general SFT advantage or transferable induction (App. D.1.9).

Composition accuracy and evidence use are distinct. On the mechanism-family benchmark, episode-matched Qwen3-4B training beats a *fixed population-prior* control on held-out compositions: response-MAE advantages are .337 [.278,.404] with the mechanism described and .357 [.328,.387] with it inferred, both clearing all four estimators at 5 seeds. A coherent gain intervention (two retained seeds) tracks better than the prior on seen mechanisms but not on held-out ones (App. D.1.10). A distribution-matched shuffled-target control at 5 seeds now replaces the fixed prior: matched training keeps a smaller advantage on the pre-registered stochastic endpoint (.094 [.032,.188] inferred). The shuffle leaves episode level near-intact while destroying response shape, so the gain falls where the manipulation bites; held-out gain tracking stays unestablished (App. D.1.11).

4.2 EXP. 4: EXTERNAL GROUNDING IN UNSEEN HISTORICAL MARKETS

Historical markets test external forecasting accuracy without labeled mechanisms. We train on resolved markets using only information available before each outcome, then evaluate 318 unseen questions from different event families. Lower Brier scores mean less squared error between forecast probabilities and realized outcomes (App. E).

Training lowers error at every tested size, but the Qwen base-relative gap shrinks from 1.7B to 32B without steady improvement in trained accuracy (Fig. 10). At 32B, trained mean Brier is .1265 versus .1266 for the base, the smallest observed gain; the trained-minus-base interval $[-.00045, .00029]$ includes zero. A Llama 3B/8B replication shows the same qualitative convergence pattern (App. Figure 16). Several trained checkpoints reach contemporaneous Polymarket performance. Hosted systems provide context, not additional points on the local scaling curve.

What Exp. 4 establishes. Outcome-based training improves mean forecasting accuracy on held-out event families, with uncertain gains at the strongest Qwen base. This comparison is trained versus base; without a shuffled-outcome control, it does not isolate the contribution of learning




Figure 10: **Training improves the tested bases, but scale alone does not guarantee a better trained forecaster.** Across Qwen3 checkpoints from 1.7B to 32B, several trained models reach crowd-level performance, while the gain narrows at the strongest base. Local model size is shown on a log parameter axis; hosted systems at right are reference points. The Llama replication is in App. Figure 16.

episode-specific relationships. It complements the controlled transfer tests with external predictive validity.

5 RELATED WORK: WHY ACCURACY IS NOT ENOUGH

Judgmental forecasting. Strong human forecasters combine base rates, decomposition, diverse evidence, and repeated revision (Tetlock, 2005; Mellers et al., 2014; Tetlock & Gardner, 2016). LM benchmarks increasingly use time-indexed or leakage-controlled questions, and retrieval or ensembling can bring systems near crowd accuracy (Zou et al., 2022; Karger et al., 2025; Halawi et al., 2024; Schoenegger et al., 2024). These results establish what probabilities models can produce. They do not tell us what would make those probabilities change.

Induction and relational reasoning. Transformers can infer functions from examples, and work on analogy and world models asks whether learned representations capture relations or dynamics (Garg et al., 2022; Chan et al., 2022; Webb et al., 2023; Yasunaga et al., 2024; Ha & Schmidhuber, 2018; LeCun, 2022; Li et al., 2023). At the same time, shortcut research shows that surface success can fail under structural change (Geirhos et al., 2020; Mitchell, 2023; Vafa et al., 2025). A forecast can therefore be right for the wrong reason: two models may report the same number while relying on different populations, evidence, or pathways.

Social agents and conditional validity. Generative agents can reproduce plausible individual and collective behavior (Park et al., 2023; Aher et al., 2023; Argyle et al., 2023), yet matching a population average is not the same as preserving the relationships within that population (Bisbee et al., 2024). For systems intended to inform decisions, we care about both the *probability reported* and the *relationship that makes the evidence predictive*. Our experiments therefore move through increasingly stringent tests: relevance and direction, contextual relationship selection, revision from local evidence, induction of a novel convention, and transfer across worlds. Each step rules out a weaker account of the same forecast without assuming a general internal world model.

6 LIMITATIONS: WHAT THESE STUDIES DO AND DO NOT ESTABLISH

These controlled tasks demonstrate context-dependent predictive relationships, not human-like understanding or a general world model. Real settings contain strategic interaction, unreliable cues, changing variables, and longer feedback horizons absent from our market packets and constructed domains.

The internal intervention concerns one checkpoint and one convention; the historical study covers one market cohort and training recipe. Additional architectures, later cohorts, interactive settings, and misleading evidence would test generality.

Performance is heterogeneous. Smaller checkpoints often struggle to induce the KIV/ZOR meaning. The episode-matching advantage under joint domain and mechanism shift does not survive the eight-seed replication; the benefit for compositions of familiar mechanisms is checkpoint-specific and survives a distribution-matched shuffled-target control only under stochastic decoding. Structure diagnostics use no target observations; the supervised pilot replicates at four seeds but differs from RL in update budget and, at that setting, learning rate. Semantic cue-to-pattern learning could explain its in-domain success; native training places its boundary at labels, not domain. The five-seed gain intervention establishes evidence use on seen but not held-out mechanisms. Historical-market gains narrow at the strongest Qwen base. The nonmonotonic scale comparison uses one recipe; larger models may require different optimization or data budgets.

7 CONCLUSION

Forecast accuracy alone does not establish inductive reasoning. Models select relevant evidence and revise contextual relationships, while training improves forecasts across domains and held-out markets. A supervised pilot learns response-form selection in distribution; its boundary is arbitrary labels, not domain. Composition accuracy exceeds a shuffled-target control on seen mechanisms without establishing transferred evidence use. These contrasts distinguish predictive accuracy from reliable relationship use without presuming a general world model.

AI USE STATEMENT

Language models produced the forecasts and experimental responses analyzed in this work and assisted with software and manuscript preparation. They also generated the context-reversal development vignettes. The authors verified numerical claims against saved experimental records and executable analyses, checked citations against original sources, reviewed the code, and take responsibility for the submission and its artifacts.

ETHICS STATEMENT

The reported analyses use public market metadata and price histories, model outputs, synthetic worlds, and an unpaid adult review of synthetic materials. The review collected no identifying or sensitive information. Under our institutional guidelines, it does not constitute human-subjects research; consequently, no IRB review was sought. Counterfactual evidence packets were evaluated offline and were not posted to markets or used for trading. We report aggregate model behavior rather than recommend deployment.

REPRODUCIBILITY STATEMENT

Appendices A–G specify the prompts, data-generating processes, sample construction, censoring, estimators, deployment settings, missingness, hyperparameters, baselines, and fail-closed checks. Builders, manifests, checksums, and saved generations support each result. Experiment 3 reports the original three-seed Qwen3-8B, Qwen3-4B, and Llama-3.1-8B roster, eight-seed Qwen extensions, a five-seed mechanism-family study with its five-seed shuffled-target control, a three-seed structure-selection pilot, and exploratory component, four-seed and native-cell supervised, and five-seed gain-intervention diagnostics. Parent-roster endpoints use round 300 and five draws under a common stochastic decoder. The appendix discloses the pilot’s interim look at training monitors and the outcome-informed supervised-launch amendment. Experiment 4 evaluates Qwen and Llama checkpoints on family-disjoint historical markets and reports hosted deployments as categorical references. Evaluation examples are excluded from training; selection rules and exploratory analyses are documented in the appendix.

ACKNOWLEDGMENTS

Acknowledgments are redacted for double-blind review.

REFERENCES

- Gati V. Aher, Rosa I. Arriaga, and Adam Tauman Kalai. Using large language models to simulate multiple humans and replicate human subject studies. In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pp. 337–371. PMLR, 2023. URL <https://proceedings.mlr.press/v202/aher23a.html>.
- Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of one, many: Using language models to simulate human samples. *Political Analysis*, 31(3):337–351, 2023. doi: 10.1017/pan.2023.2.
- James Bisbee, Joshua D. Clinton, Cassy Dorff, Brenton Kenkel, and Jennifer M. Larson. Synthetic replacements for human survey data? the perils of large language models. *Political Analysis*, 32(4):401–416, 2024. doi: 10.1017/pan.2024.5.
- A. Colin Cameron, Jonah B. Gelbach, and Douglas L. Miller. Bootstrap-based improvements for inference with clustered errors. *The Review of Economics and Statistics*, 90(3):414–427, 2008. doi: 10.1162/rest.90.3.414.
- Stephanie C. Y. Chan, Adam Santoro, Andrew K. Lampinen, Jane X. Wang, Aaditya Singh, Pierre H. Richemond, James L. McClelland, and Felix Hill. Data distributional properties drive emergent in-context learning in transformers. In *Advances in Neural Information Processing Systems*, volume 35, 2022. doi: 10.52202/068431-1371. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/77c6ccacf9962e2307fc64680fc5ace-Abstract-Conference.html.
- Shivam Garg, Dimitris Tsipras, Percy Liang, and Gregory Valiant. What can transformers learn in-context? a case study of simple function classes. In *Advances in Neural Information Processing Systems*, volume 35, 2022. doi: 10.52202/068431-2217. URL https://proceedings.neurips.cc/paper_files/paper/2022/hash/c529dba08a146ea8d6cf715ae8930cbe-Abstract-Conference.html.
- Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, and Felix A. Wichmann. Shortcut learning in deep neural networks. *Nature Machine Intelligence*, 2(11):665–673, 2020. doi: 10.1038/s42256-020-00257-z.
- Thomas L Griffiths, Nick Chater, Charles Kemp, Amy Perfors, and Joshua B Tenenbaum. Probabilistic models of cognition: Exploring representations and inductive biases. *Trends in cognitive sciences*, 14(8):357–364, 2010.
- David Ha and Jürgen Schmidhuber. World models. *arXiv preprint arXiv:1803.10122*, 2018.
- Danny Halawi, Fred Zhang, Chen Yueh-Han, and Jacob Steinhardt. Approaching human-level forecasting with language models. *Advances in Neural Information Processing Systems*, 37: 50426–50468, 2024.
- R. E. Kalman. A new approach to linear filtering and prediction problems. *Journal of Basic Engineering*, 82(1):35–45, 1960. doi: 10.1115/1.3662552.
- Ezra Karger, Houtan Bastani, Chen Yueh-Han, Zachary Jacobs, Danny Halawi, Fred Zhang, and Philip E. Tetlock. Forecastbench: A dynamic benchmark of AI forecasting capabilities. In *International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=lfPkGWXLlf>.
- Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th Symposium on Operating Systems Principles*, pp. 611–626, 2023. doi: 10.1145/3600006.3613165.
- Brenden M. Lake, Tomer D. Ullman, Joshua B. Tenenbaum, and Samuel J. Gershman. Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40:e253, 2017. doi: 10.1017/S0140525X16001837.

- Michael Lawrence, Paul Goodwin, Marcus O'Connor, and Dilek Önköl. Judgmental forecasting: A review of progress over the last 25 years. *International Journal of Forecasting*, 22(3):493–518, 2006.
- Yann LeCun. A path towards autonomous machine intelligence. *OpenReview*, 2022.
- Kenneth Li, Aspen K. Hopkins, David Bau, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Emergent world representations: Exploring a sequence model trained on a synthetic task. In *Proceedings of the Eleventh International Conference on Learning Representations*, 2023. URL https://openreview.net/forum?id=DeG07_TcZvT.
- Zichen Liu, Changyu Chen, Xinyi Wan, Chao Du, Wee Sun Lee, and Min Lin. OAT: A research-friendly framework for LLM online alignment. <https://github.com/sail-sg/oat>, 2024.
- Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. In *International Conference on Learning Representations*, 2019. URL <https://openreview.net/forum?id=Bkg6RiCqY7>.
- Barbara Mellers, Lyle Ungar, Jonathan Baron, Jaime Ramos, Burcu Gurcay, Katrina Fincher, Sydney E Scott, Don Moore, Pavel Atanasov, Samuel A Swift, et al. Psychological strategies for winning a geopolitical forecasting tournament. *Psychological science*, 25(5):1106–1115, 2014.
- Microsoft. Build agents using the responses API in microsoft foundry, 2026. URL <https://learn.microsoft.com/en-us/azure/foundry/agents/quickstarts/responses-api>.
- Melanie Mitchell. Ai’s challenge of understanding the world. *Science*, 382(6671):eadm8175, 2023. doi: 10.1126/science.adm8175. URL <https://www.science.org/doi/10.1126/science.adm8175>.
- Ollama. Ollama, 2026. URL <https://github.com/ollama/ollama>.
- Joon Sung Park, Joseph C. O’Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology*, 2023. doi: 10.1145/3586183.3606763.
- Samyam Rajbhandari, Jeff Rasley, Olatunji Ruwase, and Yuxiong He. ZeRO: Memory optimizations toward training trillion parameter models. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 2020. doi: 10.1109/SC41405.2020.00024.
- Philipp Schoenegger, Indre Tuminauskaitė, Peter S Park, Rafael Valdece Sousa Bastos, and Philip E Tetlock. Wisdom of the silicon crowd: Llm ensemble prediction capabilities rival human crowd accuracy. *Science Advances*, 10(45):eadp1528, 2024.
- Joshua B Tenenbaum, Thomas L Griffiths, and Charles Kemp. Theory-based bayesian models of inductive learning and reasoning. *Trends in cognitive sciences*, 10(7):309–318, 2006.
- Joshua B. Tenenbaum, Charles Kemp, Thomas L. Griffiths, and Noah D. Goodman. How to grow a mind: Statistics, structure, and abstraction. *Science*, 331(6022):1279–1285, 2011.
- Philip E. Tetlock. *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton University Press, 2005.
- Philip E Tetlock and Dan Gardner. *Superforecasting: The art and science of prediction*. Random House, 2016.
- Keyon Vafa, Peter G. Chang, Ashesh Rambachan, and Sendhil Mullainathan. What has a foundation model found? using inductive bias to probe for world models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pp. 60727–60747. PMLR, 2025. URL <https://proceedings.mlr.press/v267/vafa25a.html>.

Taylor Webb, Keith J. Holyoak, and Hongjing Lu. Emergent analogical reasoning in large language models. *Nature Human Behaviour*, 7:1526–1541, 2023. doi: 10.1038/s41562-023-01659-w.

Ben Wilson. Analysis of automated prompt engineering for forecasting. Metaculus, 2025. URL <https://www.metaculus.com/notebooks/38421/results-of-automated-prompt-engineer-for-forecasting/>.

Michihiro Yasunaga, Xinyun Chen, Yujia Li, Panupong Pasupat, Jure Leskovec, Percy Liang, Ed H Chi, and Denny Zhou. Large language models as analogical reasoners. In *International Conference on Learning Representations*, 2024. URL <https://arxiv.org/abs/2310.01714>.

Andy Zou, Tristan Xiao, Ryan Jia, Joe Kwon, Mantas Mazeika, Richard Li, Dawn Song, Jacob Steinhardt, Owain Evans, and Dan Hendrycks. Forecasting future world events with neural networks. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2022.

APPENDICES

The supplement is organized by experiment. A descriptive coin-flip elicitation precedes the Experiment 1 evidence-response details. Experiment 2 documents contextual relationship selection and revision, including a linear-probe representation analysis; Experiments 3–4 document controlled and historical transfer, replication, structure-selection diagnostics, supervised learnability, and mechanism evidence use. Shared training details and the artifact map close the appendix.

A	Descriptive coin-flip forecast elicitation	15
B	Experiment 1: prospective forecast updating	15
B.1	Study design, sample, and agent roster	16
B.2	Forecast, intervention, and validation protocol	17
B.3	Estimands, record processing, and aggregation	19
B.4	Extended results	21
B.5	Study limitations	23
B.6	Reproducibility and artifact map	30
B.7	Exploratory context-reversal development pilot	32
B.8	Fresh finite-rule context reversal: reasoning and wording	37
C	Experiment 2: contextual relationship selection and revision in Coin City	42
C.1	Design and data-generating process	42
C.2	Prompt arms and information available	43
C.3	Model execution and response processing	44
C.4	Estimands, benchmarks, and statistical inference	46
C.5	Response-inference diagnostics	47
C.7	Linear-probe analysis of cue-conditioned representations	54
C.9	Missingness, scope, and artifact map	65
D	Experiment 3: controlled relationship transfer	67
D.1	Design and evaluation	67
D.1.6	Later eight-seed precision extension	74
D.1.7	Structure-selection experiment	75
D.1.8	Component diagnosis with paired cue flips	75
D.1.9	Supervised learnability and its transfer boundary	76
D.1.10	Mechanism-family evidence use under coherent gain changes	78
E	Experiment 4: historical-market training and held-out evaluation	81
E.3	Held-out results	82
E.5	Qwen3-4B evidence-updating evaluation	84
F	Experiments 3–4: shared training configuration and runtime	85
F.1	Coin City A-to-B factorial	86
G	Experiments 3–4: reproducibility and artifact map	88

A DESCRIPTIVE COIN-FLIP FORECAST ELICITATION

This appendix reports the forecasts used to position the model markers in Figure 1; it is not one of the four experiments. We queried GPT-5.4, Claude Opus 4.8, and DeepSeek V4-Pro eight times each (24 calls total) at temperature 0.7 and a 600-token output ceiling. All responses parsed. The figure places each model’s logo at its eight-sample mean; the green $\approx .68$ annotation rounds the pooled mean of .679.

The exact user prompt was:

You are at a carnival game. A coin you have never seen before is flipped twice in front of you, and it comes up heads both times. These two flips are everything you know about this coin.

What is the probability that the next flip comes up heads?

Think about it intuitively, the way you actually would in the moment -- do NOT use any formula, rule, or calculation, and do not work out any numbers in your reasoning. Just reason in plain words about what you’d expect and why. Then END your reply with your gut answer as a JSON object on its own line:

```
{"rationale": "one short sentence, no math",
  "p_heads": <number between 0 and 1>}
```

Table 1: Forecasts underlying the three model markers in Figure 1. Each row summarizes eight independent temperature-0.7 calls.

Model	Mean	Median	Range
GPT-5.4	.644	.670	[.60, .67]
Claude Opus 4.8	.681	.675	[.65, .75]
DeepSeek V4-Pro	.712	.750	[.60, .75]
Pooled ($n = 24$)	.679	.670	[.60, .75]

What this probe does and does not separate. It has one condition. The 0.5 and 1.0 endpoints in Figure 1 are reference policies rather than measurements: no fair-coin arm and no no-carnival arm were collected, so the probe cannot attribute the departure from 0.5 to the carnival framing. It also does not separate a prior about carnival coins from ordinary sequential updating on an unknown bias. Laplace’s rule of succession gives $3/4$ after two heads from a coin with an unknown bias, using no world knowledge at all, and the observed .679 sits below that rather than above it. Finally, the instruction not to use any formula or calculation discourages explicitly deriving either the 0.5 or 0.75 reference policy, changing the elicitation mode and complicating comparison to both. We report the probe as a descriptive illustration of the distinction the paper draws, not as evidence for it; the four experiments carry the argument. `eval/run_coin_probe.py` implements three further arms—a stipulated fair coin, an unfamiliar coin with no carnival, and the carnival prompt without the no-calculation instruction—which would supply the missing contrast.

The prompt, temperature, output limit, parsed probabilities, rationales, and raw completions are stored in `exp2_simulated_worlds/biased_news/data/coin_probe/coin_probe_2026-06-30.jsonl`; the collection and parsing code is `exp2_simulated_worlds/biased_news/eval/run_coin_probe.py`. The prompt names neither a fair nor a rigged coin. It requests qualitative reasoning while retaining a numeric forecast for the plotted quantity.

B EXPERIMENT 1: PROSPECTIVE FORECAST UPDATING

Experiment 1 tests whether forecast revisions follow evidence direction and whether directional packets produce larger changes than orthogonal packets. All reported counts refer to parseable records in the frozen consistency report unless stated otherwise.

The main comparison reports the nine deployments with scale-valid archived revisions. This appendix retains Qwen 2.5-7B for directional-compliance diagnostics and documents why its magnitude-based statistics are excluded.

B.1 STUDY DESIGN, SAMPLE, AND AGENT ROSTER

Protocol summary. Experiment 1 crosses ten forecasting agents with a frozen set of unresolved binary Polymarket questions. The study has four stages: (1) freeze the market sample; (2) obtain one or more initial forecasts from each agent; (3) generate nine synthetic evidence packets for each of the 100 core markets (three pro- H_1 , three anti- H_1 , and three outcome-orthogonal); and (4) ask the same agent to revise each parseable initial forecast after receiving one packet at a time. The unit of analysis is an *agent* $\times$ *market* $\times$ *initial run* $\times$ *counterfactual packet* update. Requested repeat counts were five initial runs for hosted agents and three for local agents; realized valid counts are uneven because of provider failures, parsing failures, and targeted top-ups. Section B.3 specifies how those records are normalized and aggregated.

B.1.1 SAMPLING FRAME AND MARKET SELECTION

The primary sampling frame is a Polymarket API snapshot taken on June 9, 2026. Selection was deterministic given that snapshot. We retained markets that were binary, active and unresolved, scheduled to resolve 14–180 days after sampling, priced strictly between 0.10 and 0.90 on the YES contract, and supported by at least \$1,000 in USD volume. A keyword-based relevance filter retained politics, government, elections, economics, technology, and geopolitics while excluding sports, cryptocurrency, entertainment, and other out-of-scope topics. When several contracts belonged to the same event, we retained the highest-volume contract.

The recorded filter funnel makes the sampling loss explicit. The snapshot contained 60,989 binary records; 37,386 were active, 8,639 passed the topic filter, 3,570 passed the resolution-window filter, 1,403 passed the price filter, 839 passed the volume filter, and event-level deduplication left 436 candidates. A diversity pass selected the 100-market core using a cosine-similarity threshold of 0.25 and a maximum of 16 markets per topical bucket. The frozen core contains 16 global-election, 16 other, 15 Iran-nuclear, 12 U.S.-federal-policy, 10 AI/technology, 9 U.S. state-election, 7 Russia-Ukraine, 5 Brazil, 5 U.S.-midterm, 3 macroeconomic, and 2 China-Taiwan markets.

Counterfactual generation and Stage 3 updating target this 100-market core for every model, but per-model coverage of it is not uniform. Llama 3.1-70B has 99 analyzable markets because one market has no parseable initial forecast, and Claude Opus 4.8 has usable updates on all three packet arms in 73 markets; Section B.5.3 reports the coverage table and the resulting robustness checks.

Table 2: Representative markets from the frozen 100-market core. Snapshot mid is the stored YES midpoint on June 9, 2026 rather than an eventual outcome.

Question	Topic	Snapshot mid
Will Renan Santos win the 2026 Brazilian presidential election?	Brazil	0.17
US–Iran nuclear deal by June 30?	Middle East	0.20
US $\times$ Iran permanent peace deal by July 31, 2026?	Iran	0.32
Will a new Gemini flagship be released by June 30, 2026?	AI/technology	0.89

B.1.2 AGENT ROSTER AND EXECUTION MODES

The roster comprises three hosted API systems and seven open-weight local systems (Table 3). The hosted models ran through Azure AI Foundry and used its native search-supported Responses workflow (Microsoft, 2026). Local models were served by Ollama (Ollama, 2026) on cluster GPUs; their research turn was supplied by Tavily search. All agents received the same substantive three-step contract: build an event model, gather current evidence, and emit a structured forecast. Provider-specific search orchestration and model runtimes were not identical.

Table 3: Experiment 1 agent roster and requested initial-forecast repeats. Requested k is the collection target rather than an analysis denominator; exact realized coverage appears in Table 6.

Model	Runtime	Research search	Nominal params	Requested k
GPT-5.4	Azure Foundry	Foundry native	—	5
Claude Opus 4.8	Azure Foundry	Foundry native	—	5
DeepSeek V4-Pro	Azure Foundry	Foundry native	—	5
Qwen 2.5-7B	Ollama	Tavily	7B	3
Qwen 2.5-14B	Ollama	Tavily	14B	3
Qwen 2.5-32B	Ollama	Tavily	32B	3
Qwen 2.5-72B	Ollama	Tavily	72B	3
Llama 3.1-8B	Ollama	Tavily	8B	3
Llama 3.1-70B	Ollama	Tavily	70B	3
Llama 3.3-70B	Ollama	Tavily	70B	3

For local models, reasoning calls used the frozen configuration temperature of 0.1 and the final JSON-formatting call used temperature 0.0. Search queries, tool responses, token counts, parse errors, and provider errors were retained with each run. Some smaller models returned probabilities on a 0–100 rather than 0–1 scale. The initial-forecast parser converted those values, but the frozen Stage 3 data retain a mixed-scale Qwen 2.5-7B subset documented in Section B.3.1.

B.2 FORECAST, INTERVENTION, AND VALIDATION PROTOCOL

B.2.1 INITIAL FORECAST ELICITATION

Design rationale. The elicitation asks each agent to restate the resolution criterion, construct explicit YES and NO hypotheses, identify a status-quo trajectory, research current evidence, and map its final YES probability to the posterior assigned to H_1 . Question restatement and status-quo anchoring follow prior forecasting-prompt evidence (Halawi et al., 2024; Wilson, 2025).

Executed three-turn sequence. Every frozen initial forecast followed the same substantive sequence, with provider-specific tool plumbing:

1. **Event model.** The agent received the question, full resolution description, days to resolution, and category. It restated what would count as YES versus NO; identified actors, mechanisms, and latent variables; formed H_1 (YES) and H_2 (NO); and described the status-quo outcome.
2. **Current-evidence research.** The agent gathered evidence for and against each hypothesis, checked whether the required chain of events was plausible within the remaining time, and revised its assessment. Hosted systems used Azure Foundry’s native search workflow; local systems generated queries that were executed through Tavily and received the returned evidence in context.
3. **Structured forecast.** A final call converted the accumulated analysis to the common JSON contract below. A run entered downstream processing only when this object was parseable.

The research turn included this calibration instruction:

Given the days remaining, what chain of events is required for H_1 , and is that pace of change plausible? Give extra weight to the status quo; depart substantially only when concrete directional evidence warrants it.

The common output contract was:

```
{
  "event_model": {
    "key_actors": [str],
    "key_mechanisms": [str],
    "latent_variables": [str]
  },
  "hypotheses": [
    {"id": "H1", "description": str,
```

```

    "prior_probability": float, "posterior_probability": float,
    "supporting_evidence": [str], "contradicting_evidence": [str]],
    {"id": "H2", ...}
  ],
  "hypothesis_to_forecast_mapping": str,
  "yes_prob": float,
  "confidence": "low"|"medium"|"high",
  "rationale": str,
  "evidence_sources": [str]
}

```

The contract makes the intended invariant explicit: $\hat{p} = P(H_1)$. The ICS metric tests whether the returned fields obey that invariant after an intervention. The substantive prompt contract was common, but prompts were not byte-identical across runtimes because the hosted and local search interfaces require different orchestration.

B.2.2 COUNTERFACTUAL PACKET CONSTRUCTION AND VALIDATION

For each of the 100 core markets, GPT-5.4 generated nine synthetic evidence packets: three pro- H_1 , three anti- H_1 , and three orthogonal. Each generation call received the market question, resolution description, and a canonical parseable initial event model and forecast. It produced:

- **evidence_headline**: one sentence that states the new development;
- **evidence_text**: a 2–4 sentence synthetic wire-style item;
- **mechanism_targeted**: the causal pathway by which the item should affect, or remain irrelevant to, resolution;
- **direction**: pro- H_1 , anti- H_1 , or orthogonal;
- **slot_type**: one of the six roles in Table 4; and
- **plausibility_rating**: the generator’s 1–5 self-rating.

Within each direction, sequential exclusion required a new packet to target a different actor and mechanism from earlier packets. Search was disabled so the packets remained controlled synthetic premises rather than paraphrases of live reporting.

Table 4: Counterfactual slot taxonomy. Each market receives one packet of every direction–slot combination shown, for nine packets total.

Slot type	Direction	Role in the intervention
DATA	pro, anti	A measurement, price signal, statistic, or poll that shifts the likelihood of H_1 .
DECISION	pro, anti	An actor decision, official statement, policy announcement, commitment, or refusal bearing on H_1 .
STRUCTURAL	pro, anti	A technical, logistical, legal, or contextual change that makes H_1 more or less likely.
PROCEDURAL	orthogonal	A same-domain administrative or process update without a direct causal path to H_1 versus H_2 .
PERSONNEL	orthogonal	A staffing or leadership change involving a relevant actor but an issue unrelated to H_1 versus H_2 .
BACKDROP	orthogonal	Contextual economic, logistical, or social information that sets the scene without shifting the outcome odds.

Packet-level integrity checks. The frozen packet file contains 900/900 parseable records: 100 markets multiplied by nine packets, with exactly 300 pro- H_1 , 300 anti- H_1 , and 300 orthogonal records. Every packet has a nonempty headline, and the logs contain zero search calls. All 900 generator-assigned plausibility ratings equal 4/5. That field documents the generator’s compliance claim but provides neither variation nor independent validation; the separate human materials review is reported below. These checks verify structural completeness; they do not establish that directional strength or difficulty is perfectly matched across packets.

As an illustration, for the Hormuz market the pro- H_1 headline *Oman says Iran agrees daily protected windows for commercial Hormuz transits* elicited the largest single update reported in the viewer: Claude Opus 4.8 revised from 10% to 45% ($\Delta\hat{p} = +35$ pp). The orthogonal IMO incident-reporting item produced negligible movement.

B.2.3 STAGE 3 UPDATE AND PAIRING

Each packet was presented separately with the agent’s own parseable initial structured forecast. The agent was told to reason only from that forecast and the packet and to return the same structured forecast with revised H_1 , H_2 , and YES probabilities. Search was disabled.

One deployment received an extra instruction. The Stage 3 template was not common across the roster. GPT-5.4 alone was run with the runner’s `novelty` option, which appends a rule asking for a `novelty_assessment`: *a single sentence stating whether this evidence was likely already reflected in your initial forecast as background knowledge (suggesting a small update) or is genuinely new and surprising (warranting a larger revision)*. The frozen records show the asymmetry exactly: 3,838 of 3,921 GPT-5.4 updates carry the field and no other deployment carries it in any record. The instruction tells a model to modulate revision size by perceived novelty, which is close to the construct Section B.5.3 estimates, and GPT-5.4 is the deployment that behaves differently on exactly that dimension: it returns the prior unchanged on 8.4% of orthogonal packets against 71.4% and 44.4% for the other two hosted systems, writing a small revision where they abstain. We cannot separate that difference from the prompt difference, so GPT-5.4’s abstention rate is not comparable to the rest of the roster and we do not read the hosted ordering on abstention as a model difference. Its conditional magnitude, directional correctness, and AUC use the same scoring as every other deployment.

For an agent–market cell with k parseable initial runs, the intended crossed design contains $9k$ update records. Each update is conditionally independent in prompt context: it contains one initial run and one packet, never another packet or another agent’s answer. The novelty field is not a reported endpoint and should not be viewed as a validated correction.

Provider interruptions, parse failures, and subsequent top-ups make realized cells uneven. We therefore report exact valid record and market denominators by model in Table 6 rather than treating requested k or the theoretical $9k$ as observed sample sizes.

B.3 ESTIMANDS, RECORD PROCESSING, AND AGGREGATION

B.3.1 RECORD CONSTRUCTION, DEDUPLICATION, AND MISSINGNESS

The frozen analysis artifact is `expl_prospective/data/results/consistency_report_2026-06-17.json`; tables in this paper use its stored denominators rather than a fresh scan of the evolving collection directories. During report construction, date suffixes are stripped from task identifiers so daily collection files join to the same market. Counterfactuals are deduplicated by `cf_id`; updates are deduplicated by `update_id`, with a parseable record preferred over an error record when both exist. Files are read newest first for the initial-forecast and market-price lookup. Malformed JSON lines are skipped. Missing values are handled per estimand: a record may therefore contribute to ICS but not EHC, or vice versa, depending on which fields parse.

Stage 3 update fields were evaluated on their stored scale in the frozen report. For Qwen 2.5-7B, 207 of 8,636 deduplicated records store `updated_yes_prob` above 1.5; no other model does. Its directional sign metrics remain usable, but its absolute revision magnitudes and sensitivity ratio are scale-contaminated and are omitted from the paper tables.

B.3.2 CONSISTENCY ESTIMANDS

For update record i , let $d_i \in \{+1, -1\}$ denote the intended direction (+1 for pro- H_1 , −1 for anti- H_1). Orthogonal packets do not have an intended sign and are excluded from EHC and HFC. The 3 pp movement threshold excludes cases in which a model effectively leaves the forecast unchanged.

Evidence–hypothesis consistency (EHC). Let $\Delta p(H_1)_i = p(H_1)_i^{\text{post}} - \hat{p}_i^{\text{initial}}$. For directional records with $|\Delta p(H_1)_i| \geq 0.03$,

$$\text{EHC}_i = \mathbf{1}[\Delta p(H_1)_i d_i > 0]. \quad (2)$$

EHC asks whether the updated hypothesis field moves in the direction implied by the packet; it does not score the size of that movement.

Hypothesis–forecast consistency (HFC). Let $\Delta \hat{p}_i = \hat{p}_i^{\text{post}} - \hat{p}_i^{\text{initial}}$. For directional records with $|\Delta \hat{p}_i| \geq 0.03$,

$$\text{HFC}_i = \mathbf{1}[\Delta \hat{p}_i d_i > 0]. \quad (3)$$

HFC asks the same directional question of the headline YES probability. Because the output contract requires this probability to equal $P(H_1)$, we retain HFC in the appendix as a compliance diagnostic rather than as a separate main-text result.

Internal coherence score (ICS). ICS is defined whenever both updated fields parse and is unconditional on movement:

$$\text{ICS}_i = \mathbf{1}[|\hat{p}_i^{\text{post}} - p(H_1)_i^{\text{post}}| < 0.02]. \quad (4)$$

The inequality is strict: a discrepancy of exactly 2 pp does not pass.

Directional sensitivity. Packet selectivity is summarized by

$$\text{Sensitivity} = \frac{|\Delta \hat{p}|_{\text{pro+anti}}}{|\Delta \hat{p}|_{\text{orthogonal}}}. \quad (5)$$

A value near one indicates comparable movement for directional and orthogonal packets; larger values indicate preferential response to directional packets. This ratio is descriptive. It cancels a uniform multiplicative scale within a model, but does not repair mixed scales across records, hence the Qwen 2.5-7B caveat above.

B.3.3 AGGREGATION AND UNCERTAINTY

For EHC, HFC, and ICS, the evaluator first pools eligible records within each model–market cell to obtain a market-level rate. The reported model score is the unweighted mean of those market rates, so a market with more successful repeat runs does not receive greater weight. Reported standard errors are the sample standard deviation of market-level rates divided by the square root of the number of markets.

Revision magnitudes and sensitivity retain the record-pooled point estimands. Uncertainty uses 20,000 percentile cluster-bootstrap draws over represented markets. All nine packets and every repeated initial-run update from a selected market remain together, and sensitivity is recomputed from pooled directional and orthogonal movement within each draw. This accounts for dependence among records from the same market while preserving the published estimand.

B.3.4 OUTPUT COMPLIANCE AND THRESHOLD ROBUSTNESS

Output-field coupling. The final-output schema explicitly instructs the agent that $\hat{p} = P(H_1)$, and the Stage 3 prompt requests the same structure after updating. EHC and HFC are therefore not independent tests of an inferred hypothesis-to-forecast link. ICS is a schema-adherence diagnostic, while close EHC–HFC agreement reflects a mixture of directional updating and compliance with the prompted equality. Among the 26,010 relevant records with both updated fields present, 25,984 (99.90%) satisfy the equality exactly after scale normalisation and 25,987 (99.91%) satisfy it within the 2-pp ICS tolerance. We therefore do not treat EHC–HFC parity as independent evidence for a latent representation or for spontaneous propagation between separately elicited quantities.

Unconditional estimand. For threshold t , Table 5 retains a valid directional update when $|\Delta| \geq t$ and reports the fraction for which $d_i \Delta_i > 0$. At $t = 0$, every valid pro- and anti- H_1 update enters the denominator and a zero change is scored incorrect. We first compute correctness within market and then average markets equally, matching the primary analysis. The *Below 3 pp* column is the record-weighted fraction and count excluded by the published 3-pp conditional metric. The table’s n is the complete pre-threshold denominator: records require a parseable initial probability and the applicable updated field, but are not selected on the direction or magnitude of movement.

Table 5: Unconditional directional correctness and robustness to 0-, 1-, 3-, and 5-percentage-point movement thresholds. The 0-pp column is unconditional over all valid directional updates and scores zero movement as incorrect; positive-threshold columns condition on the stated magnitude. *Below 3 pp* gives the percentage and count omitted from the corresponding published 3-pp EHC or HFC denominator. Rates are macro-averaged across markets.

Model	Metric	n valid	Directional correctness				Below 3 pp
			0 pp	1 pp	3 pp	5 pp	
Claude Opus 4.8	EHC	2,119	98.2%	98.6%	99.0%	99.5%	18.1% (383)
	HFC	2,119	98.3%	98.6%	99.0%	99.5%	18.1% (383)
GPT-5.4	EHC	2,678	98.4%	98.4%	98.7%	98.8%	6.9% (185)
	HFC	2,684	98.4%	98.4%	98.7%	98.8%	6.9% (185)
DeepSeek V4-Pro	EHC	2,929	97.3%	97.6%	97.7%	98.0%	12.7% (371)
	HFC	2,929	97.3%	97.6%	97.7%	98.0%	12.6% (370)
Qwen 2.5-72B	EHC	1,795	96.3%	96.3%	96.3%	96.4%	0.8% (15)
	HFC	1,795	96.3%	96.3%	96.3%	96.4%	0.8% (15)
Llama 3.3-70B	EHC	1,788	96.4%	96.4%	96.5%	97.4%	0.7% (12)
	HFC	1,788	96.4%	96.4%	96.5%	97.4%	0.7% (12)
Llama 3.1-70B	EHC	1,659	96.1%	96.1%	96.1%	96.7%	0.6% (10)
	HFC	1,659	96.1%	96.1%	96.1%	96.7%	0.6% (10)
Qwen 2.5-32B	EHC	1,795	96.4%	96.4%	96.4%	96.5%	1.9% (34)
	HFC	1,795	96.4%	96.4%	96.4%	96.5%	1.9% (34)
Qwen 2.5-14B	EHC	1,800	93.1%	94.4%	94.3%	94.3%	2.9% (52)
	HFC	1,800	93.1%	94.4%	94.3%	94.3%	2.9% (52)
Llama 3.1-8B	EHC	3,601	88.1%	88.7%	88.7%	88.9%	1.7% (62)
	HFC	3,604	88.4%	88.7%	88.7%	88.9%	1.4% (52)
Qwen 2.5-7B	EHC	5,759	90.2%	90.2%	90.6%	90.8%	2.2% (129)
	HFC	5,759	89.9%	90.2%	90.5%	90.8%	2.4% (140)

Threshold-sweep results. The threshold sweep preserves the qualitative conclusion but narrows its wording. Hosted-model unconditional correctness is 97.3–98.4%; at the published 3-pp threshold it is 97.7–99.0%. The 3-pp rule excludes 6.9% of GPT-5.4, 12.6–12.7% of DeepSeek V4-Pro, and 18.1% of Claude Opus 4.8 directional updates. Thus, the near-1.0 EHC/HFC values do not result from discarding a majority of cases, and the conclusion is stable across 0, 1, 3, and 5 pp. The precise statement is *99% directionally correct among updates of at least 3 pp*. It is not *99% coherent updating*. Small local models have fewer sub-3-pp updates (0.6–2.9%) but lower unconditional correctness (88.1–96.4%), so their errors are not hidden primarily by the movement filter.

Normalisation and reproducibility. Before recomputing deltas, the analysis independently normalises the initial headline probability, updated headline probability, and updated H_1 posterior from the documented $[0, 1]$ or $[0, 100]$ formats. This matters for 166 relevant Llama 3.1-8B records in which H_1 was emitted as a percentage while `yes_prob` was emitted as a proportion. The complete eligible record and market counts at every threshold, invalid-field counts, market-level standard errors, input-file SHA-256 hashes, and full-precision estimates are in `expl_prospective/data/results/threshold_robustness.json`. The table is regenerated by `expl_prospective/agent/evaluate_threshold_robustness.py`; neither script nor table filters items using realized outcomes.

B.4 EXTENDED RESULTS

B.4.1 ANALYSIS COVERAGE

Table 6 gives the exact Stage 3 record and market counts. The market column counts markets with at least one valid update. Coverage of all three packet arms is not uniform across models: Llama 3.1-70B has 99 analyzable markets because one market has no parseable initial forecast, and Claude Opus 4.8

has all three arms in 73 of 100. Section B.5.3 gives the per-arm coverage and assesses whether missingness is arm-selective.

Table 6: Valid Stage 3 update records and represented core markets by agent. Hosted systems first; open-weight rows in descending parameter count.

Model	Type	Records	Markets
Claude Opus 4.8	Hosted	3,125	100
GPT-5.4	Hosted	3,913	100
DeepSeek V4-Pro	Hosted	4,395	100
Qwen 2.5-72B	Local 72B	2,692	100
Llama 3.3-70B	Local 70B	2,682	100
Llama 3.1-70B	Local 70B	2,490	99
Qwen 2.5-32B	Local 32B	2,693	100
Qwen 2.5-14B	Local 14B	2,700	100
Llama 3.1-8B	Local 8B	5,404	100
Qwen 2.5-7B	Local 7B	8,636	100
Total		38,730	

B.4.2 PROSPECTIVE SETTLEMENT AUDIT

The frozen-cohort evaluation includes a dated settlement audit that leaves the directional-updating estimand unchanged. On August 24, 2026, we queried the public Polymarket market endpoint for the exact 100 frozen IDs. We admitted a label only when the market was closed and the YES/NO token prices were exactly $[1, 0]$ or $[0, 1]$. This strict rule yielded 60 scoreable markets (21 YES and 39 NO). One additional closed market had an invalid 0.5/0.5 settlement and 39 remained unresolved at the audit cutoff; all 40 were excluded.

For each agent–market cell, the forecast is the median of its parseable repeated initial YES probabilities, matching the maintained Experiment 1 viewer. The crowd benchmark is the YES price stored in the frozen June 9 selection record. We report threshold accuracy (YES iff $p > 0.5$), Brier score, and natural-log loss. Proper-score contrasts are paired agent-minus-market differences; negative values favor the agent. Percentile intervals use 20,000 market-level bootstrap resamples.

Table 7: Descriptive accuracy on the first 60 strictly settled markets. A/M denotes agent/June 9 market on the same rows; Correct is the number of thresholded calls matching the outcome. Lower Brier and log loss are better, and each Δ is agent minus market. Llama 3.1-70B lacks a parseable initial forecast for one settled market, so its paired denominator is 59.

Model	n	Correct A/M	Brier A/M	Δ Brier [95% CI]	Log A/M	Δ log [95% CI]
Claude Opus 4.8	60	41/39	.242/.246	-.004[-.063, .054]	.738/.700	.038[-.129, .210]
GPT-5.4	60	40/39	.225/.246	-.021[-.074, .030]	.663/.700	-.037[-.178, .103]
DeepSeek V4-Pro	60	40/39	.257/.246	.011[-.050, .069]	.793/.700	.093[-.093, .282]
Qwen 2.5-72B	60	32/39	.265/.246	.019[-.031, .066]	.728/.700	.028[-.097, .144]
Llama 3.3-70B	60	30/39	.266/.246	.020[-.038, .078]	.732/.700	.032[-.113, .177]
Llama 3.1-70B	59	33/38	.258/.249	.008[-.058, .075]	.712/.707	.004[-.155, .164]
Qwen 2.5-32B	60	32/39	.264/.246	.018[-.044, .078]	.728/.700	.028[-.126, .178]
Qwen 2.5-14B	60	37/39	.277/.246	.031[-.031, .092]	.827/.700	.127[-.049, .308]
Llama 3.1-8B	60	31/39	.266/.246	.020[-.047, .085]	.741/.700	.041[-.127, .202]
Qwen 2.5-7B	60	38/39	.230/.246	-.016[-.075, .041]	.651/.700	-.049[-.195, .090]

The raw correct-call counts are close for the hosted agents and crowd: Claude Opus 4.8 gets 41/60 correct, GPT-5.4 and DeepSeek V4-Pro each get 40/60, and the June 9 market side gets 39/60. GPT-5.4 has the lowest Brier score (.225 versus .246 for the market), a paired difference of $-.021 [-.074, .030]$; its log-loss difference is $-.037 [-.178, .103]$. Qwen 2.5-7B has the lowest log loss (.651 versus .700), but gets 38/60 thresholded calls correct. No agent’s paired Brier or log-loss interval excludes zero.

This audit describes the outcome-ready slice of the frozen cohort. Settlement by the audit date censors the cohort toward shorter-horizon and early-resolving questions, so the 60 are not a random subset of the frozen 100. Moreover, the benchmark is the June 9 snapshot whereas agent forecasts

were collected from June 10 through June 17; later-running agents may have had additional public information. The result therefore does not establish superiority to the crowd. It documents the prospective check; the directional-update estimates do not, by themselves, imply a detectable accuracy advantage.

B.4.3 REVISION MAGNITUDES AND PACKET SELECTIVITY

Table 8 reports mean absolute forecast revision $|\Delta\hat{p}|$ by packet direction. For the nine models whose updates are consistently stored on 0–1, hosted-system sensitivity is 8.2–24.0 $\times$, compared with 1.9–6.4 $\times$ for local systems. These are the unconditional ratios; Table 9 separates the abstention and magnitude factors behind them, and the main text reports those two rather than the product.

Qwen 2.5-7B has mixed 0–1 and 0–100 stored updates (Section B.3.1); we therefore omit its absolute magnitudes and sensitivity ratio rather than present them as comparable estimates.

Table 8: Mean absolute forecast revision by counterfactual direction. Hosted systems first; open-weight rows in descending parameter count. Sensitivity is mean directional divided by mean orthogonal movement; n is the valid Stage 3 record count. Values are percentage points for models with consistently stored 0–1 updates.

Model	Type	pro- H_1	anti- H_1	orthogonal	Sensitivity	n
Claude Opus 4.8	Hosted	12.7 pp	9.6 pp	0.5 pp	24.0 $\times$	3,125
DeepSeek V4-Pro	Hosted	13.9 pp	11.4 pp	1.4 pp	8.9 $\times$	4,395
GPT-5.4	Hosted	10.6 pp	8.5 pp	1.2 pp	8.2 $\times$	3,913
Qwen 2.5-72B	Local 72B	15.0 pp	10.2 pp	3.0 pp	4.2 $\times$	2,692
Llama 3.3-70B	Local 70B	14.8 pp	14.3 pp	2.3 pp	6.4 $\times$	2,682
Llama 3.1-70B	Local 70B	14.8 pp	15.5 pp	3.3 pp	4.6 $\times$	2,490
Qwen 2.5-32B	Local 32B	15.0 pp	10.3 pp	4.0 pp	3.2 $\times$	2,693
Qwen 2.5-14B	Local 14B	14.8 pp	8.6 pp	4.1 pp	2.9 $\times$	2,700
Llama 3.1-8B	Local 8B	11.1 pp	12.0 pp	5.9 pp	1.9 $\times$	5,404
Qwen 2.5-7B	Local 7B	—	—	—	—	8,636

Qwen 2.5-7B magnitude cells are omitted because 207 deduplicated updates are stored above 1.5, mixing probability scales within the row. Table 8 documents the magnitude contrast; Figure 11 then places directional consistency and packet selectivity on a common visual scale.




Figure 11: **Directional consistency and packet selectivity by deployment.** *Left:* EHC is 87–90% for 7–8B local models and 98–99% for hosted systems. *Right:* sensitivity is 1.9–6.4 $\times$ for local models and 8.2–24.0 $\times$ for hosted systems; Qwen 2.5-7B is omitted from this panel because its archived absolute revisions mix 0–1 and 0–100 scales and therefore do not yield a comparable ratio. The panel plots the unconditional ratio, which Table 9 decomposes.

B.5 STUDY LIMITATIONS

B.5.1 SAMPLING AND OUTCOME SCOPE

The sample is a filtered slice of liquid, medium-horizon Polymarket questions rather than a probability sample of forecasting tasks. Topic exclusions, the 0.10–0.90 price window, the 14–180 day horizon, event deduplication, and the diversity cap shape its difficulty distribution. The primary Experiment 1 estimand evaluates conditional updating on questions that were unresolved at collection. The settlement audit in Section B.4.2 is naturally censored to the first 60 strict settlements and is descriptive. Experiment 4 separately evaluates accuracy on a locked sample of resolved historical markets with synchronized forecast cutoffs and market-price baselines.

B.5.2 WHAT THE UPDATE PROMPT WITHHOLDS

The Stage 3 template fills in one packet field, `evidence_text`. The packet record also holds the intended direction, `target_hypothesis`, `expected_hypothesis_shift`, `slot_type`, `evidence_headline`, `mechanism_targeted`, and a free-text rationale. None of these reaches the agent; all are read only when the update is scored. The agent sees its own prior forecast, the snippet, and the fixed instruction.

We checked this two ways. Over the 12,728 frozen records that kept their prompt string, spanning all three hosted systems, the evidence block is identical to the packet’s `evidence_text` and contains no private packet field. Prompt persistence was a per-runner logging choice, so this covers a subset of records by logging coverage rather than by outcome; the runners that stored nothing use the same one-field template. Turning to the snippets, none of the 900 names H_1 , H_2 , the market, or its resolution. Twenty-two (2.4%, 10 pro- H_1 and 12 anti- H_1) use any likelihood wording—*odds*, *chance of*, *moreless likely*, *probability*, *likelihood*—and each time it describes something in the world, a simulation percentage or a betting line, not the market’s resolution. The sign is not stated; it must be inferred from the relation between the reported event and the market outcome. Reproduced by `expl_prospective/agent/audit_packet_polarity.py`.

B.5.3 ALTERNATIVE EXPLANATIONS FOR THE SELECTIVITY GAP

Here we state the competing accounts of the movement gap and what we find for each. All checks come from `expl_prospective/agent/audit_selectivity_robustness.py`, `audit_orthogonal_abstention.py`, `audit_lexical_sign_baseline.py`, and `audit_group_contrast.py`.

Where the numbers come from. The checks read the frozen June 17 report, which enumerates every scored run, so they use exactly the data behind the published numbers. The enumeration matches the report’s own per-arm counts for seven of the nine models; Claude Opus 4.8 and GPT-5.4 carry 14 and 4 extra runs in markets absent from the per-market join, 18 of 38,730. Reproduced sensitivity ratios agree with the published ones to within 0.2. Three checks need packet text and so join runs to the June 10 packet set; 346 runs (1.1%) do not join because they come from a June 7–8 pilot against a superseded 63-packet generation, and they are dropped from those three checks only.

Attrition is much higher for the hosted systems. Claude Opus 4.8 made 9,313 calls and yielded 3,232 usable updates, losing 3,656 to broken pipes, 914 to HTTP 401s after a key expired mid-run, and 1,481 to unparseable responses. Of those, 1,461 come from one malfunctioning June 15 retry pass that failed at 91%; the main June 10 collection parsed at 99.3%, so the published compliance figures are not survivorship. DeepSeek V4-Pro made 6,338 calls for 4,598 usable updates, losing 1,629 to HTTP 401s; GPT-5.4 made 4,269 for 4,057. The six local models lose almost nothing—four record no failed calls, the other two lose 8 and 15 responses to parsing. Content-based exclusion is confined to 28 GPT-5.4 calls across 2 markets. Recorded failures are predominantly transport or parsing failures; the coverage check below assesses arm and market imbalance, not content-specific selection. Separately, two update streams were corrupted by concurrent writes and repaired in August 2026, after the freeze; the original bytes and unrecoverable fragments sit in `expl_prospective/data/_recovery_audit/`. Nothing in this section reads those files.

The generator saw one agent’s world model. All packets were built from a single canonical initial forecast, and it is DeepSeek V4-Pro’s: its June 10 YES probability matches the stored generation prior on 100 of 100 markets, where no other agent matches on more than 28. For the other eight models the evidence came from a different system’s event model. The seed model gains nothing by

it—among hosted systems it is second on sensitivity ($8.9\times$ against Claude’s $24.0\times$), third on AUC (.967), and last on EHC (.977)—but its estimates are the least independent in the roster.

Coverage is uneven across arms. Table 10 gives per-arm market coverage. Seven models cover 99 or 100 markets on every arm. Claude Opus 4.8 covers 82 on the orthogonal arm and 77 on the directional arms, 73 in common, so its pooled ratio divides means taken over slightly different market sets. Restricting to its 73 common markets gives $23.7\times$ against $24.2\times$ pooled. As a proxy check, every other model has similar sensitivity on the markets Claude covers and those it misses (8.4 vs. 7.6 for GPT-5.4, 9.7 vs. 7.2 for DeepSeek V4-Pro, 4.3 vs. 4.2 for Qwen 2.5-72B, 1.9 vs. 2.1 for Llama 3.1-8B).

The ratio has an unstable denominator. Claude’s orthogonal mean is 0.5 pp, so its $24\times$ comes from a small divisor rather than unusually large directional movement, and its market-clustered interval is correspondingly wide, [17.3, 36.9]. Table 11 reports three statistics that avoid that division. The difference of means is 5.6–12.3 pp and does not order the models the way the ratio does. The within-market probability that a directional packet moves the forecast more than an orthogonal one runs from .791 for Llama 3.1-8B to .994 for GPT-5.4, putting GPT-5.4 above Claude Opus 4.8 and reversing the ratio ranking. Every model moves further on directional packets in every market it covers.

The denominator is mostly zeros. That instability has a specific cause, and naming it replaces the ratio with two statistics that are easier to read. A revision of exactly zero—the model reproduces the probability shown to it—enters the mean like any other record, so the ratio silently mixes how often a deployment declines to move with how far it moves when it does. Writing A for the share of runs with $\Delta\hat{p} = 0$ and M for the mean absolute revision among the rest, $|\Delta\hat{p}| = (1 - A)M$, and the ratio factors exactly:

$$\text{Sensitivity} = \underbrace{\frac{M_{\text{dir}}}{M_{\text{orth}}}}_{\text{magnitude}} \times \underbrace{\frac{1 - A_{\text{dir}}}{1 - A_{\text{orth}}}}_{\text{abstention}}. \quad (6)$$

Table 9 gives both factors. Directional abstention is 0.0–1.4% for every deployment, so the second factor is governed entirely by the orthogonal arm, where abstention runs from 1.1% for Llama 3.1-8B to 71.4% [63.8, 78.5] for Claude Opus 4.8. Claude’s $24\times$ is therefore $7.0\times$ [5.5, 9.1] in magnitude multiplied by 3.48 in abstention; DeepSeek V4-Pro’s $8.9\times$ is $4.9\times$ times 1.79; GPT-5.4’s $8.2\times$ is $7.5\times$ times 1.09, because GPT-5.4 writes a small revision where the other two hosted systems return the prior unchanged. Across the roster the published 1.9–24.2 $\times$ becomes 1.9–7.5 $\times$ once the zeros are separated out, and the ordering changes: GPT-5.4 rather than Claude Opus 4.8 is highest, and Llama 3.3-70B at $5.7\times$ sits above DeepSeek V4-Pro at $4.9\times$.

We report both factors rather than the product. Declining to revise at all is an appropriate response to an outcome-irrelevant premise, so the abstention rate is a selectivity measure in its own right and a better-behaved one: it is a bounded proportion, it is paired within model across arms, and it needs no division by a small number. It is also the more fragile of the two. An exact copy of a displayed prior is a discrete act adjacent to formatting, and the update prompt does instruct the agent to leave unaffected fields unchanged, so we do not read a high rate as evidence about internal computation. The no-information condition described in Section B.5 is what would separate declining to revise from copying the anchor, and it was not run. Reproduced by `expl_prospective/agent/audit_orthogonal_abstention.py`.

Table 9: Decomposition of the sensitivity ratio into magnitude and abstention factors (Equation 6). *Orthogonal abstention* is the share of orthogonal runs with $\Delta\hat{p}$ exactly zero; the directional-arm rate is 0.0–1.4% throughout. The two *given a move* columns condition on a nonzero revision, and their ratio is the magnitude factor. Brackets are 95% market-clustered percentile intervals from 20,000 draws, the scheme used for the published movement intervals. Qwen 2.5-7B is omitted for mixed probability scales.

Model	Orthogonal abstention	Mean $ \Delta\hat{p} $ given a move		Sens. given a move	Sens. published
		orthogonal	directional		
Claude Opus 4.8	71.4% [63.8, 78.5]	1.6 pp	11.2 pp	$7.0 \times$ [5.5, 9.1]	$24.2 \times$ [17.3, 36.9]
GPT-5.4	8.4% [5.2, 11.9]	1.3 pp	9.6 pp	$7.5 \times$ [6.6, 8.5]	$8.2 \times$ [7.2, 9.3]
DeepSeek V4-Pro	44.4% [38.9, 50.0]	2.6 pp	12.7 pp	$4.9 \times$ [4.4, 5.5]	$8.9 \times$ [8.0, 9.9]
Qwen 2.5-72B	24.2% [19.6, 29.0]	3.9 pp	12.6 pp	$3.2 \times$ [3.0, 3.4]	$4.2 \times$ [3.9, 4.6]
Llama 3.3-70B	10.7% [7.0, 14.8]	2.6 pp	14.6 pp	$5.7 \times$ [5.2, 6.2]	$6.4 \times$ [5.8, 7.0]
Llama 3.1-70B	2.5% [0.8, 5.0]	3.4 pp	15.2 pp	$4.5 \times$ [4.2, 4.8]	$4.6 \times$ [4.3, 4.9]
Qwen 2.5-32B	11.1% [8.0, 14.6]	4.5 pp	12.6 pp	$2.8 \times$ [2.6, 3.1]	$3.2 \times$ [2.9, 3.5]
Qwen 2.5-14B	26.1% [20.7, 31.9]	5.5 pp	11.9 pp	$2.2 \times$ [2.0, 2.3]	$2.9 \times$ [2.6, 3.2]
Llama 3.1-8B	1.1% [0.5, 1.9]	6.0 pp	11.6 pp	$1.9 \times$ [1.8, 2.0]	$1.9 \times$ [1.8, 2.1]

Abstention and movement disagree about which deployments are selective. The abstention rate is not interchangeable with leaving an orthogonal premise alone, which is why it appears here rather than beside the ratio in the main text. The quantity that answers the plain question – how far does a deployment move on a premise that cannot bear on the outcome – is mean absolute orthogonal movement, unconditional on moving, and it orders the deployments differently. That mean is 0.46 pp for Claude Opus 4.8, 1.16 for GPT-5.4, 1.43 for DeepSeek V4-Pro, 2.29 for Llama 3.3-70B, 2.98 for Qwen 2.5-72B, 3.29 for Llama 3.1-70B, 3.97 for Qwen 2.5-32B, 4.07 for Qwen 2.5-14B, and 5.92 for Llama 3.1-8B.

The two orderings conflict in three places, and the conflicts are large. GPT-5.4 abstains on 8.4% of orthogonal runs, seventh of nine, while moving less than every deployment except Opus. DeepSeek V4-Pro abstains more than five times as often at 44.4%, yet moves further on average, 1.43 against 1.16 pp. Qwen 2.5-14B ranks third on abstention at 26.1% and eighth on movement at 4.07 pp. A deployment that emits round numbers scores high on an exact-zero criterion; one that perturbs a final digit scores near zero even when it barely moves, so part of the spread reflects output granularity rather than selectivity. Reading the abstention column as a selectivity ranking would therefore invert the ordering for three of nine deployments. We report it because declining to revise is a meaningful response in its own right and because it is a factor in Equation 6, not as a measure of how far a deployment leaves an irrelevant premise alone.

The deployment contrast depends on the statistic. Comparing the three hosted systems against the six scale-valid open-weight models by market-clustered bootstrap (20,000 replicates, macro-averaged within group so a model with more runs does not dominate), the statistics disagree. On the ratio the difference is $+9.8$ [7.2, 14.2], group means $13.7 \times$ and $3.9 \times$, and the per-model ranges do not overlap. On AUC it is $+0.070$ [.062, .079], which excludes zero, but the ranges do overlap: Llama 3.3-70B at .989 and Llama 3.1-70B at .978 both beat DeepSeek V4-Pro at .967. On raw movement there is no difference at all, $+0.8$ [−0.0, 1.7] pp around means of 10.1 and 9.3.

The decomposition explains why the ratio is the one statistic that separates the groups completely. Both of its factors favour the hosted systems on average and neither separates them on its own: orthogonal abstention differs by $+0.286$ [.244, .329] around group means of .413 and .126, with four open-weight deployments at or above GPT-5.4’s 8.4%; the magnitude factor differs by $+3.10$ [2.31, 4.08] around $6.5 \times$ and $3.4 \times$, with Llama 3.3-70B above DeepSeek V4-Pro. Multiplying two partially separating factors yields a product whose ranges happen not to overlap, so we do not report that non-overlap as a finding. The two classes move about equally far on directional evidence; the hosted systems more often decline to move at all on controls, and move somewhat less far when they do. We state the deployment result that way rather than as a general selectivity ordering. Group contrasts come from `expl_prospective/agent/audit_group_contrast.py`.

The arms differ in wording. By construction directional packets take the DATA, DECISION, and STRUCTURAL slots and orthogonal packets take PROCEDURAL, PERSONNEL, and BACKDROP,

and the text differs accordingly: directional snippets average 96.8 words and 7.1 digits, orthogonal ones 85.7 and 3.7. Movement could track density rather than content. Two checks argue otherwise. Within the directional arm, where length cannot indicate sign, Spearman correlation between word count and $|\Delta\hat{p}|$ is $-.06$ to $.00$ and between digit count and $|\Delta\hat{p}|$ is $.01$ to $.10$. And the arms' word counts almost fully overlap (orthogonal 66–114, directional 66–121); restricting to that band keeps 889 of 900 packets and moves every sensitivity estimate by at most 0.1 (Table 11, last column). The arms are not matched on surface form. These checks show that word and digit counts do not explain the effect; other surface differences remain. Directional correctness is not explained by length, since length cannot say which way to move.

A shallow classifier might be enough. We tested it. Multinomial naive Bayes and L2-regularised logistic regression were fit on the 900 packet texts with five-fold cross-validation grouped by market, so no market appears on both sides of a fold. Telling directional from orthogonal packets is nearly free: naive Bayes reaches .989 against a .667 majority rate. That follows from the slot-type plan rather than revealing anything—the most distinctive orthogonal tokens are *appointed*, *appointment*, *oversee*, and *accreditation*, which name the roles the generator was told to write. So a model could route on that marking without weighing content, and the sensitivity ratio alone cannot rule it out. Matching the control arm on evidence type would need a different packet design; the checks above bound only the length part of the problem.

Recovering the *sign* is another matter. On pro- H_1 versus anti- H_1 the same classifiers reach .688 and .675 against .500 chance, and prepending the market question does not help (.665 and .680); these bag-of-words models do not combine question polarity with event valence. The tokens it keys on carry no consistent sign—*vessel*, *venue*, *signature* for pro- H_1 , *table*, *materials*, *findings* for anti- H_1 . Against this .69 lexical baseline, the human panel reaches .762 on its smaller categorical task and the hosted agents .973–.984 on probability revision. These are not matched performance comparisons. The agent-classifier gap is evidence against this bag-of-words account, not proof that no surface heuristic can recover sign; unlike the arm contrast, it stays within pro- versus anti- H_1 packets.

Does the stated reasoning track the packet? If movement were routed on surface marking, a model's written rationale would not need to bear any particular relation to what it does with the forecast. We tested two versions of this (`expl_prospective/agent/audit_mechanism_tracking.py`).

The first asks whether a rationale identifies the specific mechanism its packet was built to target, ranking the true `mechanism_targeted` against the other eight from the same market so that question, actors, and domain vocabulary are held fixed. Rationales rank it first on .278–.430 of records against 1/9 chance, but the snippet alone scores .516, so most of that is echo of the text the model just read. Removing every token the snippet supplied leaves .128–.195 under matched query length, still above chance—but the generator's own rationale for the packet, scored the same way, lands at .119, essentially chance, because it paraphrases the mechanism rather than restating it. With the reference writer at chance the metric has no usable dynamic range, and we draw no conclusion from it.

Models often declare orthogonal packets irrelevant: explicit non-relevance wording appears in 4.9–61.9% of orthogonal rationales against 0.1–2.8% of directional ones. What matters is whether saying it predicts doing it. Within the orthogonal arm, records whose rationale declares irrelevance move less than those that do not, and the market-clustered gap excludes zero for eight of the nine scale-valid deployments: 0.88 pp [0.56, 1.22] for Claude Opus 4.8, 1.70 [1.45, 1.96] for DeepSeek V4-Pro, 2.73 [2.28, 3.19] for Qwen 2.5-72B, down to 0.36 [0.21, 0.54] for GPT-5.4, whose unflagged orthogonal movement is already 1.3 pp. The exception is Llama 3.1-8B at 0.26 [−0.64, 1.06], where no coupling is detectable; it is also the weakest deployment on sensitivity and directional correctness, though we do not test whether the two are related. Qwen 2.5-7B is omitted for mixed probability scales.

For eight of nine deployments, then, the revision a model makes is not independent of the relevance it states, and the one deployment where that link is undetectable is the weakest on every other measure. Two limits: the flag is a hand-built lexicon, so it captures what a rationale declares rather than what the model understood; and consistent self-description is compatible with surface routing, so this does not show relevance was derived from the mechanism.

One deployment was prompted differently. GPT-5.4, and only GPT-5.4, received the `novelty_assessment` rule at Stage 3 (Section B.2.3). Because that rule speaks directly to

revision size, its abstention rate is not comparable to the other deployments'. The within-GPT contrast uses one template, but the extra instruction could interact with packet arm, so its contrast should not be attributed to model differences alone.

Two remaining caveats. A single generator wrote all packets and scored its own work at a uniform 4/5 for plausibility, so automated validation establishes structural completeness but not equal difficulty or independent content validity. And the local models ran through Ollama at the default tag for each name, with the served weight precision not recorded per run; rows ordered by parameter count also vary in quantization and inference stack, so we do not attribute the open-weight ordering to parameter count alone.

Table 10: Per-arm market coverage in the frozen June 17 report. Markets counts entries in the model's per-market block, which for three Qwen rows includes 20 carrying no scored update; the arm columns count markets with at least one.

Model	Markets	pro- H_1 /anti- H_1	orthogonal	all three
Claude Opus 4.8	100	77	82	73
GPT-5.4	100	99	99	99
DeepSeek V4-Pro	100	100	100	100
Qwen 2.5-72B	120	100	100	100
Llama 3.3-70B	100	100	100	100
Llama 3.1-70B	100	99	99	99
Qwen 2.5-32B	120	100	100	100
Qwen 2.5-14B	120	100	100	100
Llama 3.1-8B	100	100	100	100

Table 11: Selectivity under statistics that do not divide by the orthogonal mean. *Sens.* is the frozen ratio; *Diff.* is mean directional minus mean orthogonal absolute revision; *AUC* is the within-market probability that a directional packet moves the forecast more than an orthogonal one, averaged over markets; *Markets* counts markets where mean directional movement exceeds mean orthogonal movement; *Length-m.* recomputes the ratio on the 66–114-word band the arms share. Qwen 2.5-7B is omitted for mixed probability scales.

Model	Sens.	Diff.	AUC	Markets	Length-m.
Claude Opus 4.8	24.0×	10.7 pp	0.984	77/77	24.3×
GPT-5.4	8.2×	8.4 pp	0.994	99/99	8.1×
DeepSeek V4-Pro	8.9×	11.2 pp	0.967	100/100	8.8×
Qwen 2.5-72B	4.2×	9.6 pp	0.950	100/100	4.2×
Llama 3.3-70B	6.4×	12.3 pp	0.989	100/100	6.4×
Llama 3.1-70B	4.6×	11.9 pp	0.978	99/99	4.6×
Qwen 2.5-32B	3.2×	8.7 pp	0.906	100/100	3.2×
Qwen 2.5-14B	2.9×	7.7 pp	0.852	100/100	2.9×
Llama 3.1-8B	1.9×	5.6 pp	0.791	100/100	1.9×

B.5.4 HUMAN MATERIALS REVIEW

Nine adult reviewers submitted the 18-item materials protocol. It samples one report from each of 18 distinct markets, balanced across six pro-YES, six anti-YES, and six orthogonal packets. One reviewer met the disclosed response-pattern exclusion below, leaving eight quality-eligible reviewers. We treat the study as descriptive materials evidence. Before the synthetic report, every item shows the market question, resolution criteria, and matching summary from the frozen June 10 initial forecast. The consent screen and every item instruct reviewers to treat June 10, 2026 as the current date and use no external information or assistance. Intended directions, source IDs, generator metadata, and model responses remain hidden.

Reviewers record the report's conditional direction, confidence, clarity, real-world plausibility, and usability as an explicit hypothetical premise. We report exact agreement with the frozen direction and Cohen's κ between the panel-majority labels and the frozen key. Reviewer-level Cohen's κ , the

component ratings, and every item-level judgment appear below. Ambiguous judgments remain in every denominator.

Practice responses are stored separately and excluded from the analysis.

The analyzed sample contains eight quality-eligible reviewers (R1 and R3–R9), exceeding the six-reviewer target for this descriptive materials check. Table 12 includes every eligible reviewer.

Exact-direction agreement is 108/144 (75.0%). The panel-majority label matches the frozen direction on 17 of 18 items (94.4%; Cohen’s $\kappa = .919$): all six anti- H_1 items, all six orthogonal items, and five of six pro- H_1 items. Across individual judgments, exact agreement is 35/48 for pro- H_1 , 37/48 for anti- H_1 , and 36/48 for orthogonal packets. Reviewers marked the premise usable on 124/144 judgments (86.1%). The corresponding all-reviewer sensitivity is reported below. These descriptive results indicate that readers usually recovered the intended conditional direction and could use the reports as stipulated premises. They also bound how legible that direction is: 75.0% exact agreement and mean pairwise Cohen’s $\kappa = .47$ indicate substantial but imperfect recovery across the four response labels. The panel is not a matched baseline for the agents—reviewers emit a categorical label on 18 items, agents emit a revised probability against their own prior on all 900—so we read it as evidence that sign had to be inferred, not as a human–model comparison.

Table 12: **Stage 3 human materials judgments for eight quality-eligible reviewers.** Direction is exact agreement with the frozen conditional direction and κ is Cohen’s κ against that key over the four response options. Confidence, clarity, and plausibility (C/C/P) are reviewer-level medians on 1–5 scales. The pooled row is descriptive because reviewers rated the same items; the majority row takes the modal label per item, resolving ties to *ambiguous*.

Reviewer	Direction agreement	κ vs. key	Premise usable	C/C/P
R1	12/18 (66.7%)	.54	10/18	4/4/4
R3	12/18 (66.7%)	.51	13/18	4/4/3
R4	18/18 (100.0%)	1.00	18/18	5/5/5
R5	10/18 (55.6%)	.38	18/18	4/4/4
R6	13/18 (72.2%)	.61	12/18	4/4/5
R7	15/18 (83.3%)	.76	17/18	5/4/3
R8	12/18 (66.7%)	.55	18/18	4/5/4
R9	16/18 (88.9%)	.84	18/18	5/5/5
Pooled	108/144 (75.0%)	—	124/144	—
Majority label	17/18 (94.4%)	.92	—	—

Table 13: **Item-level Stage 3 judgments for eight quality-eligible reviewers.** Key is the frozen conditional direction; Majority is the modal panel response, with ties assigned to *ambiguous*. Agree and Usable give counts out of eight; C/L/P gives median direction confidence, clarity, and plausibility on 1–5 scales. Item numbers follow the frozen public packet order.

Question	Key	Majority	Agree	C/L/P	Usable
1. Hormuz traffic normal by July 31	More	More	8/8	5/5/4.5	8/8
2. Iran reaches World Cup knockouts	Less	Less	6/8	4.5/4.5/4.5	7/8
3. Israel strikes Yemen by June 30	No effect	No effect	8/8	4.5/4/3.5	6/8
4. Netanyahu exits election by July 31	More	More	6/8	5/5/4.5	8/8
5. Anthropic valuation reaches \$1.1T	Less	Less	7/8	4.5/4.5/5	8/8
6. Oladokun starts Chiefs Week 1	No effect	No effect	6/8	4/4/3.5	7/8
7. S&P 500 reaches 7,700 in June	More	More	4/8	3.5/4/5	7/8
8. Trump unfreezes Iranian assets	Less	Less	5/8	3.5/4/4	7/8
9. May durable-goods orders: 0–2%	No effect	No effect	4/8	3/3/3.5	6/8
10. China Q2 growth: 4.6–4.9%	More	Ambig.	3/8	3/4/3.5	5/8
11. Eisenkot joins Bennett–Lapid alliance	Less	Less	6/8	5/5/5	7/8
12. Russia enters Orikhiv by July 31	No effect	No effect	8/8	5/4/4.5	6/8
13. Romanian PM is independent/technocrat	More	More	6/8	4.5/4.5/4	8/8
14. Russia captures Kostyantynivka	Less	Less	5/8	3.5/5/5	7/8
15. At least four June ChatGPT outages	No effect	No effect	4/8	3/4/4	7/8
16. El-Sayed wins Michigan primary	More	More	8/8	5/5/5	8/8
17. Little is MN-02 Democratic nominee	Less	Less	8/8	5/5/4	7/8
18. OpenAI IPO close: \$1.25–1.5T	No effect	No effect	6/8	4/3.5/3.5	5/8

The response-pattern quality rule excludes any reviewer who selects one direction on at least 17 of 18 balanced items. R2 met this rule by selecting *more likely* on 17 of 18 balanced items, including all

six frozen negative items and five of six frozen null items. Raw responses are retained, and no item is removed or relabeled. Retaining R2 gives 115/162 (71.0%) exact agreement across all nine reviewers.

The task requests no identifying or sensitive information. It was conducted as an internal materials review under the authors’ institutional guidance; no IRB protocol was opened.

Figure 12 records the reviewer-facing workflow. The capture script at `expl_prospective/stage3_materials_annotation/capture_review_screenshots.py` loads the frozen item packet against an empty temporary annotation database. Item order was independently randomized by reviewer.

Orthogonal packets are a useful outcome-relevance control, but are not equivalent to receiving no new text. A literal no-information update condition would separate generic prompt-induced revision from response to topically adjacent content. The very small hosted-system orthogonal revisions (0.5–1.4 pp) do not substitute for that control.

B.5.5 PROMPT, RETRIEVAL, AND RUNTIME CONFOUNDING

Agents shared a substantive task contract, while execution was not byte-identical. Hosted models used Azure Foundry search, while local models used Tavily results in an Ollama workflow; temperatures and provider behavior also differ. Hosted–local gaps therefore combine model capability with retrieval and runtime differences.

B.5.6 SCALE, MISSINGNESS, AND DEPENDENCE

Realized repeat counts are uneven because failures and top-ups differ by provider. Equal-weight market aggregation prevents high-coverage markets from dominating EHC, HFC, and ICS. Movement and sensitivity intervals use 20,000 market-clustered bootstrap draws that retain every packet and repeated initial-run update from each selected market. The mixed-scale Qwen 2.5-7B row described in Section B.3.1 remains an appendix compliance diagnostic and is excluded from the main magnitude comparison.

B.5.7 PREMISE ACCEPTANCE AT STAGE 3

Search is intentionally unavailable during updating, and agents are told to accept the synthetic headline as true. This design measures conditional premise use rather than source verification or epistemic refusal. The human materials review uses the same instruction and therefore assesses comprehension and premise usability. These results do not establish whether a model should trust or verify the synthetic report.

B.6 REPRODUCIBILITY AND ARTIFACT MAP

The paper treats the June 17 model-update report and the August 23 final materials-review summary as the frozen numerical authorities. The relevant artifacts are:

- **Sampling:** `expl_prospective/fetch_markets/select_markets.py` and the June 9 selection/diversity JSON and manifests in `expl_prospective/data/selected_markets/`.
- **Initial elicitation:** `expl_prospective/agent/forecast_agent.py` for hosted runs and `expl_prospective/agent/run_local_forecast.py` for Ollama runs, with records and manifests in `expl_prospective/data/initial_forecasts/`.
- **Interventions:** `expl_prospective/agent/build_counterfactuals.py` and the frozen `expl_prospective/data/counterfactuals/counterfactuals_2026-06-10.jsonl`.
- **Updates:** hosted and local update runners under `expl_prospective/agent/`, with update records in `expl_prospective/data/updated_forecasts/`.
- **Evaluation:** `expl_prospective/agent/evaluate_consistency.py` and `expl_prospective/data/results/consistency_report_2026-06-17.json`;

Time frame: June 10, 2026 is today.
Judge each report from that point in time. Do not use events, outcomes, or information learned after June 10, 2026.

Item 13 of 180 saved

FORECASTING QUESTION

Will China GDP growth in Q2 2026 be between 4.6% and 4.9%?

Information available on June 10, before the synthetic report

Multiple independent forecasts (Goldman Sachs 4.8%, Citi 4.7%, Vanguard 4.7%, Trading Economics 4.7%) for Q2) converge on growth within the 4.6-4.9% band. The 2026 official target was lowered to 4.5-5.0%, making a sub-5.0% print politically acceptable, while the high Q2 2025 base (5.2%) creates mechanical downward pressure. Q1 2026 at 5.0% already shows moderation toward the band. Countervailing risk remains that stimulus could sustain ≥5.0% growth or that tariff shocks could push growth below 4.6%, but the central tendency of expert forecasts strongly favors the 4.6-4.9% range.

Read the full resolution criteria

SYNTHETIC REPORT

This report is fictional research material dated June 10, 2026. Treat June 10, 2026 as today; judge the report as written and do not search for it online.

China expands real-time services data in GDP tally, smoothing quarterly swings

BEIJING, June 10 — China's National Bureau of Statistics said on Wednesday it had broadened the use of tax e-invoice, freight and online retail settlement data in its preliminary quarterly GDP accounting, a technical change officials said would improve coverage of the fast-growing services sector. Two economists briefed by the bureau said the expanded now-casting system would likely reduce large quarter-to-quarter revisions and make year-on-year readings track underlying trend growth more closely. The NBS said the updated compilation process would be used in the Q2 preliminary accounting release due in July.

1 Conditional direction

If this synthetic report were accurate, how would it affect the probability that the market resolves YES?

YES becomes more likely

YES becomes less likely

No material effect on YES

Ambiguous or cannot be determined

2 Quick ratings

Choose one score on each row. For plausibility, ignore whether the fictional report actually occurred.

Confidence in your direction judgment

Not confidentVery confident

12345

Clarity of the report and its relevance

Very unclearVery clear

12345

Plausibility as a real news development

Very implausibleVery plausible

12345

3 Usable hypothetical premise

Is the report coherent and specific enough to use as a hypothetical premise if a forecasting task explicitly tells you to assume it is accurate?

YesUnclearNo

Save and continue

Your answer can be edited later.

Time frame: June 10, 2026 is today.
Judge each report from that point in time. Do not use events, outcomes, or information learned after June 10, 2026.

Item 13 of 180 saved

BEFORE YOU BEGIN

Information and consent

You are invited to review 18 short synthetic news reports created for a forecasting study. Your judgments will be used to characterize the reports, not to evaluate you. The reports are fictional research materials and are not claims about events that actually occurred.

For this task, place yourself on June 10, 2026. Treat that date as today and judge each report using only the question, resolution criteria, the June 10 background summary, and synthetic report on the page. Every item includes the background information available before the synthetic report. Do not use events, outcomes, or information learned after June 10, 2026.

The task should take approximately 15-20 minutes. Do not search the web, use AI assistant, or discuss the items while completing the task. Participation is voluntary, and you may stop at any time. This form does not require your name, email address, demographics, IP address, or sensitive personal information.

☐ I am at least 18 years old, have read the information above, and consent to participate.

Begin review

(a) Information, temporal frame, and consent. (b) A complete item with June 10 context, synthetic report, and five judgments.

Figure 12: **Stage 3 materials-review interface.** Reviewers enter a private code before seeing the information-and-consent screen. Each item repeats the June 10 temporal frame, presents the question and pre-report background, labels the fictional evidence as a synthetic report, and collects conditional direction, confidence, clarity, plausibility, and premise usability. Item order was independently randomized by reviewer. Intended direction and private packet metadata are hidden.

31

expl_prospective/agent/evaluate_clustered_uncertainty.py reconstructs the 20,000-draw market-clustered intervals in expl_prospective/data/results/clustered_movement_uncertainty.json; expl_prospective/agent/audit_orthogonal_abstention.py writes the abstention/magnitude decomposition of Table 9 to expl_prospective/data/results/orthogonal_abstention.json and its table source, under the same resampling scheme.

- **Human materials review:** the frozen protocol, blinded packet, browser application, and tests in expl_prospective/stage3_materials_annotation/; make_final_summary.py reconstructs data/exports/final_summary_current.json, while data/exports/agreement_current.json is the paper-facing agreement export.
- **Figures:** paper/generate_figures.py writes the direction-selective panel; paper/generate_paper_figures.py writes the combined EHC/sensitivity panel.

The model-update tables are defined by the June 17 frozen JSON. The mutable collection directories contain later records, so evaluating their present contents would define a different dataset. Bit-for-bit reconstruction begins from the frozen JSON; a full pipeline reconstruction uses only the manifests and records included in that freeze.

B.7 EXPLORATORY CONTEXT-REVERSAL DEVELOPMENT PILOT

The original packet design leaves a relevance shortcut: a word-only classifier separates directional from orthogonal packets with .989 accuracy. We therefore ran a development pilot in which the news is identical within each family and only its relationship to the forecast outcome changes. These are 20 assistant-authored fictional families, without independent human editing or validation; the pilot develops a measurement instrument rather than providing confirmatory evidence about real-world forecasting. The families span five reused mechanism classes, not 20 distinct mechanisms.

Design and execution. Each family supplies a question, resolution rule, background, prior information, and three relational contexts. Identical news is intended to favor the outcome, oppose it, or have no bearing on it when the connecting relationship is broken. A fourth condition omits the relationship description; it has no prespecified sign or null expectation. Instructions ask models to condition on the supplied fictional premises, with no web retrieval. Condition labels and private annotations are excluded from prompts.

For each context, the model first gives a separate baseline probability. Three independent branches then receive the same baseline and context with either new news, a no-news message, or a verbatim repeat of information already present at baseline. The repeated-information control is not a second presentation of the genuinely new evidence. All branches use the same instructions about novelty, and no branch sees another branch’s answer. This separates contextual changes from news-induced revision.

Qwen2.5-72B-Instruct, Llama-3.1-70B-Instruct, and Qwen3-32B ran from locally cached checkpoints with three repeats, temperature 0.7, top- $p = 1$, and a 128-token output limit; Qwen3 thinking was disabled. Syntax-constrained JSON required a probability in $[0, 1]$ without specifying its value or direction. Each model produced 240 baselines and 720 updates: all 960 records were valid, for 2,880 records in total. Infrastructure retries replaced no valid responses.

Analysis and results. Directional accuracy uses all 120 planned positive/negative trials per model; paired reversal requires both signs to be correct in each of 60 planned pairs. Zero movement, invalid responses, and missing responses count as failures; there is no minimum-movement filter. Broken-link movement uses 60 trials per model. Intervals use 2,000 whole-family bootstrap draws, preserving contexts, branches, and repeats. Thus 60 repeated pairs do not represent 60 independent families.

Table 14: Exploratory context reversal on independently unvalidated materials. Brackets give descriptive 95% family-bootstrap intervals. Sign accuracy and paired reversal include every planned trial; broken-link movement is the mean absolute change in percentage points.

Model	Correct sign (%)	Paired reversal (%)	Broken $ \Delta p $ (pp)
Qwen2.5-72B	62.5 [52.5, 73.3]	30.0 [11.7, 50.0]	13.3 [9.2, 17.4]
Llama-3.1-70B	65.0 [53.3, 77.5]	36.7 [16.7, 56.7]	10.5 [4.8, 17.8]
Qwen3-32B	64.2 [55.8, 73.3]	31.7 [15.0, 48.3]	13.1 [8.8, 17.5]

Table 14 shows only 30.0–36.7% paired reversal. The prespecified broken-link stability criterion requires the signed-mean 90% interval strictly within ± 2 percentage points and the upper endpoint of the 95% interval for mean absolute movement below 2 points. All three models fail: mean absolute movement is 10.5–13.3 points. All 240 no-news branches per model exactly reproduce their priors, making raw and drift-adjusted results identical. Mean absolute movement after repeated information, pooled across contexts, is .375, .021, and .521 points for Qwen2.5-72B, Llama-3.1-70B, and Qwen3-32B, respectively. Under this prompting protocol, ordinary re-elicitation drift does not explain the much larger broken-link response to news.

Exploratory frontier extension. After inspecting the completed local pilot, we added GPT-5.6 Sol (gpt-5.6-sol) on all 20 original families and all four contexts, using only repeat zero. The 80-unit plan was frozen before its API responses were collected; no family, label, or prompt was revised. The Responses API used low reasoning effort, a 2,048-token output limit, strict JSON-schema output, no tools, and no supplied temperature or seed. All 320 calls returned valid probabilities. This is a post-local-results development extension, not a prespecified fourth member of the original local roster. Local models used different decoding and output constraints, so the comparison concerns these deployments rather than an isolated effect of model identity.

Table 15: Matched single-repeat development comparison on the same 20 families. Each model contributes 40 directional trials, 20 reversal pairs, and 20 broken-link trials, with 320/320 valid records. Brackets are descriptive 95% family-bootstrap intervals; inference settings differ between deployments.

Model	Correct sign (%)	Paired reversal (%)	Broken $ \Delta p $ (pp)
Qwen2.5-72B	62.5 [52.5, 72.5]	25.0 [5.0, 45.0]	12.8 [8.8, 16.8]
Llama-3.1-70B	65.0 [52.5, 77.5]	35.0 [15.0, 55.0]	9.8 [3.5, 17.0]
Qwen3-32B	62.5 [50.0, 75.0]	30.0 [10.0, 50.0]	13.2 [8.0, 18.8]
GPT-5.6	92.5 [85.0, 100.0]	85.0 [70.0, 100.0]	0.5 [0.0, 1.5]

GPT-5.6 obtains 37/40 correct directions and 17/20 successful reversal pairs (Table 15). Broken-link absolute movement averages 0.5 points; its signed-mean 90% interval is $[-1.5, 0.0]$ and its upper 95% absolute-movement bound is 1.5, meeting the original 2-point stability criterion on this development set. Nineteen broken-link trials have zero change; one changes by -10 points. All 80 no-news and all 80 repeated-news branches exactly reproduce their own baselines. None of the 40 directional baselines mechanically blocks the expected sign. The three failed directional updates remain in the denominator, including the ambiguous reservoir family. These results support testing relational relevance in a fresh validated cohort, while the local failures show that performance is not uniform across deployments.

Exploratory direction-only companion. After inspecting the probability results, we elicited standalone direction classifications from the same three local checkpoints, retaining all 20 families and four contexts at repeat zero. Each model separately classified 240 scenario–message combinations: new news, no news, and repeated information for each context. No numerical prior was shown or forecast elicited. Constrained JSON selected increase, decrease, unchanged, or unclear, with temperature 0.7, top- $p = 1$, and 128 output tokens; Qwen3 thinking remained disabled. All 720 planned outputs were valid. Scored denominators retain every planned trial; unclear, invalid, and missing outputs fail determinate targets. Masked new-news judgments have no target and remain descriptive. The completed human materials review covers the original 18 evidence packets; its documented scope is unchanged.

Table 16: Exploratory standalone direction judgments on the unchanged development families. Entries are correct/planned counts. Paired reversal requires both signed contexts to be correct. Broken-link news and the 160 no-news/repeated-news controls target unchanged. Each model has 240/240 valid outputs; unclear is valid syntax but fails a determinate target.

Model	Correct sign	Paired reversal	Broken unchanged	Controls unchanged
Qwen2.5-72B	25/40	6/20	6/20	150/160
Llama-3.1-70B	31/40	11/20	7/20	155/160
Qwen3-32B	28/40	9/20	2/20	150/160

Direction-only paired reversal is 6/20, 11/20, and 9/20 for Qwen2.5, Llama, and Qwen3, respectively (Table 16), compared with 5/20, 7/20, and 6/20 in their original repeat-zero probability responses. For both Llama and Qwen3, four of five signed classification-only successes occur where the original numerical prior mechanically blocked the required update: a positive target starting at 1 or a negative target starting at 0. On unblocked trials, Llama scores 27/34 categorically versus 26/34 numerically; Qwen3 scores 24/36 versus 25/36. Qwen2.5 has no blocked endpoints and scores 25/40 in both tasks. The positive net gains for Llama and Qwen3 are therefore concentrated in endpoint cases. Their separate elicitations and single repeat do not isolate the cause of those differences. All three models continue to miss many paired reversals and usually fail to classify broken-link news as unchanged, despite largely correct no-news and repeated-information controls. The original probability outcomes and interval-based stability criterion remain the primary numeric analysis.

Exploratory Qwen3 reasoning follow-up. The same Qwen3-32B checkpoint ran with thinking enabled and disabled on all 20 families and four contexts at repeat zero. Each arm elicited 80 context-specific priors and three separately generated update branches per prior, for 320 planned responses. Updates received only their own arm’s prior. Both modes used unconstrained decoding, identical request seeds, temperature 0.7, top- $p = 1$, a 4,096-token output limit, and an 8,192-token context limit. The original constrained, 128-token non-thinking deployment remains a separate reference. After observing non-thinking outputs, a recorded parsing amendment permitted exactly one optional Markdown fence around otherwise strict final probability JSON. This post-output rule applied symmetrically to both modes; original format compliance is separately reported. Every received raw answer was preserved and reparsed without regeneration. Format recovery generated only previously unattempted update branches whose baselines became parseable. Truncation, unfinished reasoning, and other remaining failures were retained; the parser never extracted a probability from unfinished reasoning. All planned trials remain in scored denominators, with the original sign, paired-reversal, and broken-link stability criteria.

Table 17: Exploratory Qwen3 reasoning follow-up on the same 20 families. Correct-sign and paired-reversal entries retain all planned trials. The two unconstrained arms use the symmetric format amendment; the original constrained arm is a descriptive reference. Broken-link mean absolute movement is observed only, with its missing trial shown explicitly.

Qwen3-32B deployment	Correct sign	Paired reversal	Broken $ \Delta p $ (pp)	Valid/planned
Original constrained, thinking disabled	25/40	6/20	13.25	320/320
Unconstrained, thinking disabled	25/40	6/20	14.05	320/320
Unconstrained, thinking enabled	30/40	12/20	7.63 (19/20 observed)	300/320

All 320 records per arm are present. The disabled arm has 320 valid probabilities after recovery: 103 previously rejected fenced answers were reparsed, and 198 previously blocked update branches were generated for the first time. Before that amendment, its strict-format record counts were 19 valid, 103 parse failures, and 198 blocked branches. The enabled arm retains 300 valid records, five truncated baselines, and their 15 blocked updates; format recovery changes none of these outcomes. Four truncated baselines belong to the radio family (dev_16), and one to the positive carton-transfer context (dev_06). These failures remain in all planned binary denominators.

Thinking enabled obtains 30/40 correct signs and 12/20 paired reversals, compared with 25/40 and 6/20 for matched unconstrained thinking disabled (Table 17). Using 2,000 paired whole-family bootstrap draws, the enabled-minus-disabled contrast is +12.5 percentage points for signs (descriptive

95% interval $[-2.5, 27.5]$) and +30.0 points for paired reversal ($[5.0, 55.0]$). Each paired binary contrast retains all 20 families, including failures. Direction-blocking baseline endpoints occur in 5/40 valid disabled baselines and 1/37 valid enabled baselines; three enabled directional baselines are missing.

Broken-link absolute movement averages 14.05 points for disabled thinking (20/20 observed) and 7.63 for enabled thinking (19/20 observed). On their 19 common complete pairs, the respective means are 14.74 and 7.63 points, giving an enabled-minus-disabled contrast of -7.11 points ($[-16.32, 1.32]$). The missing radio-family measurement is never replaced by zero. The original stability criterion is not met by disabled thinking and is indeterminate for enabled thinking because a planned measurement is missing. Enabled outputs contain nonempty reasoning in all 305 received completions, with 300 closing the reasoning segment and five exhausting the token limit. Mean output length is 966.4 tokens per received enabled completion versus 12.5 for disabled thinking. These results describe the matched deployment change on this development sample; uncertainty, changed baseline judgments, and remaining failures limit conclusions about the mechanism of improvement.

Controlled robustness development on the original families. A separately versioned development run retains all 20 original parent families, including previous successes and failures. Each family has a repaired base, consistent entity renaming, a paraphrase, and a resample with identical visible text but different request IDs and seeds. Explicit transfer destinations repair underspecified routes; signed and broken contexts share the same inventory of positive, negative, and side-channel roles, with actor assignments changed and row order counterbalanced. These revisions do not replace the original results. Repair changes the supplied premises and can change elicited priors, so an old-versus-repaired contrast is not a causal estimate of clarification.

All four versions retain positive, negative, broken, and masked contexts. Each numerical context receives its own baseline and three independent branches: new news, no news, and information already present at baseline. The separate categorical task presents no numerical prior. The three original local checkpoints use constrained JSON, temperature 0.7, top- $p = 1$, and a 128-token output limit; Qwen3 thinking is disabled. Each model has all 1,280 numerical and 960 categorical records. Valid outputs are Qwen2.5-72B: 1280/1,280 numeric and 960/960 categorical; Llama-3.1-70B: 1280/1,280 numeric and 960/960 categorical; Qwen3-32B: 1280/1,280 numeric and 960/960 categorical. Invalid responses remain scored failures, and magnitude estimates use observed measurements with missingness shown explicitly. The independent audit verifies execution, provenance, and scoring; it is not independent human validation of these revised materials.

Table 18 reports every model and version. Zero numerical movement fails a signed target. Paired reversal requires both signed contexts to be correct; categorical unclear fails a determinate target. All intervals are descriptive, using 2,000 whole-parent-family bootstrap draws (seed 20260921), carrying contexts, branches, and versions together. There are 20 sampling families; the four versions do not create 80 independent families.

Table 18: Controlled robustness development. Binary entries are correct/planned counts: 40 signed trials, 20 reversal pairs, and 20 broken-link judgments per version. Broken movement is mean absolute numerical revision in percentage points, with its descriptive 95% family-bootstrap interval.

Model	Version	Numeric sign	Numeric pair	Broken $ \Delta p $ (pp)	Direction pair	Broken unchanged
Qwen2.5-72B	Repaired base	28/40	8/20	15.50 [12.25, 18.25]	9/20	11/20
Qwen2.5-72B	Renamed	27/40	8/20	16.50 [13.50, 19.25]	10/20	10/20
Qwen2.5-72B	Paraphrased	25/40	5/20	16.50 [13.25, 19.25]	7/20	12/20
Qwen2.5-72B	Same-text resample	27/40	7/20	16.25 [13.25, 19.00]	10/20	10/20
Llama-3.1-70B	Repaired base	24/40	8/20	6.46 [2.04, 12.92]	13/20	8/20
Llama-3.1-70B	Renamed	24/40	9/20	11.75 [5.00, 20.00]	14/20	6/20
Llama-3.1-70B	Paraphrased	23/40	6/20	8.70 [3.45, 15.00]	12/20	7/20
Llama-3.1-70B	Same-text resample	25/40	8/20	8.38 [3.25, 15.09]	12/20	5/20
Qwen3-32B	Repaired base	22/40	7/20	10.00 [4.75, 15.75]	8/20	2/20
Qwen3-32B	Renamed	24/40	6/20	18.25 [8.24, 31.25]	9/20	3/20
Qwen3-32B	Paraphrased	23/40	4/20	10.80 [6.55, 15.15]	5/20	2/20
Qwen3-32B	Same-text resample	22/40	5/20	13.50 [6.49, 22.00]	9/20	2/20

For Qwen2.5-72B, Llama-3.1-70B, and Qwen3-32B, respectively, repaired-base numerical paired reversal is 8/20, 8/20, 7/20, and broken-link mean absolute movement is 15.50, 6.46, 10.00 points. The corresponding interval-based stability statuses are not met, not met, not met.

The numerical broken-link stability diagnostic retains the original criterion: the signed-mean 90% interval must lie strictly within ± 2 points and the upper 95% bound on mean absolute movement must be below 2 points, with all planned measurements observed. Detailed statuses are retained in the source summaries; a small point estimate alone does not establish stability.

Editing and stochastic variation. Let Δp_v be a version’s new-news update relative to its own context’s baseline. Table 19 separates mean absolute baseline discrepancy D_0 from mean absolute update discrepancy $D_\Delta = |\Delta p_v - \Delta p_{\text{base}}|$, each averaged across the 80 planned family-context pairs. The within-2-point count describes individual paired observations. The excess column averages $E_\Delta = |\Delta p_v - \Delta p_{\text{base}}| - |\Delta p_{\text{resample}} - \Delta p_{\text{base}}|$ on common observed measurements. The same-text resample is a stochastic reference at the same temperature; its excess is zero by construction. Renaming and paraphrase discrepancies must therefore be read alongside this reference. A single resample does not isolate a causal effect of editing, and an interval including zero does not establish equivalence or invariance.

Table 19: Each edited version or same-text resample compared with repaired base. Numerical discrepancies are in percentage points; brackets give 95% parent-family-bootstrap intervals for excess over resampling. Agreement uses 80 planned categorical pairs, including masked new news. Joint correctness uses 60 determinate pairs, excluding masked new news. Stable wrong answers count as agreement but fail joint correctness. Missing/invalid binary pairs remain failures; magnitudes are observed only.

Model	Vs. repaired base	Baseline D_0	Update D_Δ	Excess E_Δ [95% CI]	$D_\Delta \leq 2$ pp	Direction agree	Both correct
Qwen2.5-72B	Renamed	5.25	5.34	1.28 [−0.81, 3.56]	47/80	72/80	38/60
Qwen2.5-72B	Paraphrased	4.81	5.80	1.74 [−0.69, 4.19]	44/80	70/80	36/60
Qwen2.5-72B	Same-text resample	2.88	4.06	0.00 [0.00, 0.00]	47/80	75/80	39/60
Llama-3.1-70B	Renamed	9.37	7.69	1.48 [−1.95, 5.52]	43/80	72/80	38/60
Llama-3.1-70B	Paraphrased	9.98	7.25	1.04 [−1.90, 4.25]	45/80	69/80	36/60
Llama-3.1-70B	Same-text resample	6.86	6.21	0.00 [0.00, 0.00]	53/80	72/80	37/60
Qwen3-32B	Renamed	11.04	11.19	−0.30 [−4.78, 5.09]	38/80	69/80	26/60
Qwen3-32B	Paraphrased	17.36	14.31	2.83 [−1.63, 7.53]	26/80	62/80	24/60
Qwen3-32B	Same-text resample	11.86	11.49	0.00 [0.00, 0.00]	42/80	72/80	28/60

Controls and lexical checks. No-news and repeated-information controls remain separate (Table 20); the repeated information was already present at baseline and is not a second presentation of genuinely new evidence. Each control has 80 planned contexts per model and version. Their performance does not substitute for directional or broken-link performance.

Table 20: Separate no-news and repeated-information controls. Within-2-point and categorical unchanged entries use all 80 planned contexts. Mean absolute numerical movement is in percentage points; any missing magnitude measurements are shown explicitly. Full descriptive family-bootstrap intervals are retained in the source summaries.

Model	Version	No news: ≤ 2 pp	$ \Delta p $	Unchanged	Repeated: ≤ 2 pp	$ \Delta p $	Unchanged
Qwen2.5-72B	Repaired base	80/80	0.00	80/80	78/80	0.19	80/80
Qwen2.5-72B	Renamed	80/80	0.00	80/80	79/80	0.06	80/80
Qwen2.5-72B	Paraphrased	80/80	0.00	80/80	78/80	0.19	79/80
Qwen2.5-72B	Same-text resample	80/80	0.00	80/80	80/80	0.03	80/80
Llama-3.1-70B	Repaired base	80/80	0.00	80/80	80/80	0.02	79/80
Llama-3.1-70B	Renamed	80/80	0.00	80/80	80/80	0.00	80/80
Llama-3.1-70B	Paraphrased	80/80	0.00	80/80	79/80	0.12	77/80
Llama-3.1-70B	Same-text resample	80/80	0.00	80/80	80/80	0.00	75/80
Qwen3-32B	Repaired base	80/80	0.00	80/80	78/80	0.19	75/80
Qwen3-32B	Renamed	80/80	0.00	80/80	73/80	0.88	74/80
Qwen3-32B	Paraphrased	80/80	0.00	80/80	77/80	0.31	73/80
Qwen3-32B	Same-text resample	80/80	0.00	80/80	75/80	0.75	73/80

Whole-parent-family-held-out news-unigram, full-text unigram, full-text bigram, and character-ngram baselines obtain directional accuracies of 33.3%, 33.3%, 33.3%, 32.5%, and balanced relevance accuracies of 50.0%, 50.0%, 50.0%, 47.8%, respectively. The chance references are 1/3 and 1/2. These measurements concern the tested lexical baselines; they do not establish that all shortcuts are removed. The revised materials and all three checkpoints remain development evidence. No independent fresh-cohort result is asserted here.

Interpretive limits of the original pilot. In the original three-repeat pilot, baseline probabilities at the relevant endpoint mechanically prevent a correct signed update in 0/120 Qwen2.5-72B, 19/120 Llama, and 12/120 Qwen3 trials; these remain failures. The original pilot materials retain ambiguities: the carton-transfer family (dev_06) does not specify where cartons leaving the buffer go, and the reservoir family (dev_07) does not explicitly say that outgoing water bypasses the turbines. The resin-conveyor (dev_10) and silt-pump (dev_18) families likewise leave outlet destinations unspecified. No family was excluded or relabeled after observing responses. The development results do not establish reliable relational relevance judgments on independently validated or real-world materials. Any material revision requires a new development version; a fresh, independently validated and frozen cohort is needed for confirmation.

B.8 FRESH FINITE-RULE CONTEXT REVERSAL: REASONING AND WORDING

The development studies above confound material version with the reasoning setting. We therefore froze a new diagnostic before collecting any target-model responses: 80 new finite-rule instances, ten in each of eight mechanism classes, with the same name, wording, and resampling controls in both Qwen3-32B thinking settings. The earlier qualitative 80-concept blueprint was replaced by these fully specified rules. They are parameterized instances of eight shared classes and one outcome wrapper, not 80 independent mechanisms or a representative natural-news sample. None of the original development outputs enters this cohort.

Explicit mechanisms and residual uncertainty. The classes cover detector calibration, prerequisite timing, component quorums, procedural waivers, proportional brokerage, directed reachability, bounded feedback, and dependence among stored reports. Each of three independently specified channels has private variables and primitive draws, including a binary state X . An audit revises $P(X = 1)$ from $1/4$ to $3/4$ or conversely, without revealing a realized state. Other conditional rules remain fixed. In calibration and report-dependence cases, the stated X probabilities already condition on stored observations; the observation is not applied twice.

The project completes if a final inspection succeeds and either a reserve team succeeds or the ordinary channel fires without the cancellation channel firing. The independent inspection and reserve probabilities are $4/5$ and $1/5$. Writing a, b for the ordinary and cancellation firing probabilities gives

$$P(\text{complete}) = \frac{4}{5}\{\frac{1}{5} + \frac{4}{5}a(1 - b)\}.$$

The archive channel has neither an outcome pathway nor an information link to other channels. Thus identical news can raise, lower, or leave the forecast unchanged solely through actor-to-channel bindings. All channel descriptions appear in every context, with identical word and punctuation inventories. Half the questions ask about noncompletion, and rising/falling audits are balanced. Exact enumeration or Bayes calculations, cross-checked in a second implementation, establish every sign and retain both outcomes after every update. The question probability stays between .16 and .84; no normative forecast is forced to an endpoint. Component probabilities are supplied, but the model's initial forecast is elicited separately in each context and never replaced by the oracle probability.

For example, in the first calibration instance the same audit changes a named channel's $P(X = 1)$ from $1/4$ to $3/4$. Its oracle baseline is .3079 in each context. The new probabilities are .3253, .2773, and .3079 when that actor operates the ordinary, cancellation, or archive channel, respectively. These oracle values are analysis quantities, not supplied model forecasts.

Screening and its limits. No new human validation was obtained. Separate Llama-3.1-8B and Qwen3-14B checkpoints screened all 720 distinct packets (80 instances, three wording versions, three contexts), without intended labels, sibling contexts, or target outputs. Same-text resamples reuse an identical packet. The final screens had poor directional agreement: 224/720 and 182/720 planned judgments matched the rule-derived labels; 44 Llama completions were invalid. These screens do not validate the labels. Every valid endpoint judgment retained both outcomes. Author adjudication repaired two actual text defects in earlier drafts: an over-broad independence statement and reuse of "inspection" for two different events. The final version explicitly distinguishes private variables and a local approval coin from the final inspection. All versions and judgments are retained. No instance was removed because a reviewer answered incorrectly. Label authority is the disclosed formal specification and author cross-check, not model consensus or independent human judgment. Qwen3-14B also shares a model family with the target.

Frozen paired collection. Each instance has base, name-only, paraphrase, and same-text resample variants. Each variant contains three relational contexts and one own-context baseline with three independent new-news, no-news, and already-known-information branches. This gives 3,840 planned records per local reasoning mode. Qwen3-32B uses the same checkpoint, temperature .7, top- $p = 1$, request seeds, unconstrained decoding, 8,192-output-token allowance, and 16,384-token context in both modes; its thinking chat prefix is the intended manipulation. A final probability object may have one Markdown fence in either mode; enabled reasoning must close first. All received invalid/truncated responses and durably recorded interrupted requests are terminal failures, without resampling. This is a comparison of reasoning settings under these deployments, not a general causal decomposition of capability.

Paired reversal requires both signed updates to be strictly correct; zero movement and all invalid or missing responses fail the binary endpoint. Magnitudes use observed values with missingness shown. The prespecified broken-link criterion requires a signed-mean 90% interval strictly inside ± 2 points and an upper 95% mean-absolute bound below 2 points; any missing planned broken measurement makes it indeterminate. Intervals use 2,000 whole-instance bootstrap draws, keeping contexts, wordings, and modes paired. At all-correct or all-zero cells, these empirical intervals collapse; this does not establish population certainty. They are descriptive for the authored corpus; shared mechanism/template dependence limits generalization. Numerical calculation and rounding remain possible sources of error even when the rules do not force endpoint forecasts. The released scoring also reports accuracy in prespecified oracle-effect bins below 1, 1–2, and at least 2 percentage points.

Valid final probabilities: thinking off: 3,840/3,840; thinking on: 3,803/3,840.

Table 21: Fresh local evaluation. Signs are correct updates out of 160 planned signed contexts per wording; pairs are correct reversals out of 80, followed by a 95% interval in percent. Broken movement is mean absolute change with its 95% interval, in percentage points; the final column reports observed broken measurements.

Thinking	Wording	Signs	Pairs [95% CI]	Broken $ \Delta $, pp	Observed
Off	Base	84/160	14/80 [10.0, 26.2]	4.2 [2.9, 5.7]	80/80
Off	Renamed	86/160	14/80 [10.0, 26.2]	2.7 [1.8, 3.6]	80/80
Off	Paraphrased	81/160	12/80 [7.5, 22.5]	2.1 [1.4, 3.1]	80/80
Off	Same-text resample	83/160	11/80 [6.2, 21.3]	3.3 [2.4, 4.4]	80/80
On	Base	157/160	77/80 [91.2, 100.0]	0.2 [0.0, 0.5]	80/80
On	Renamed	159/160	79/80 [96.2, 100.0]	0.0 [0.0, 0.0]	79/80
On	Paraphrased	154/160	74/80 [86.2, 97.5]	0.0 [0.0, 0.0]	79/80
On	Same-text resample	155/160	75/80 [87.5, 98.8]	0.1 [0.0, 0.2]	78/80

The base-wording thinking contrast is 78.8 [68.8, 88.8] percentage points. The leave-one-mechanism-out contrasts range from 75.7 to 82.9 points. Thinking off: no wording variant meets the stability criterion. Thinking on: Base meet the stability criterion.

Table 22: Matched thinking contrasts in paired reversal. The final column subtracts the base-wording contrast; all contexts and modes remain paired by parent instance. Brackets are descriptive 95% intervals.

Wording	On minus off, pp	Contrast minus base, pp
Base	78.8 [68.8, 88.8]	–
Renamed	81.2 [72.5, 88.8]	2.5 [–8.8, 13.8]
Paraphrased	77.5 [67.5, 86.2]	–1.2 [–15.0, 12.5]
Same-text resample	80.0 [71.2, 88.8]	1.2 [–10.0, 12.5]

Table 23: Fresh controls and oracle discrepancy, all in percentage points. No-news and repeated-information columns give mean absolute movement across the three contexts. Baseline and updated error are mean absolute deviations from the exact rule-based probabilities, not calibration against realized real-world outcomes.

Thinking	Wording	No news	Repeated	Baseline error	Updated error
Off	Base	0.00	0.00	28.75	26.89
Off	Renamed	0.00	0.00	27.33	24.52
Off	Paraphrased	0.00	0.00	37.79	35.77
Off	Same-text resample	0.00	0.00	27.84	25.70
On	Base	0.03	0.00	0.13	0.08
On	Renamed	0.01	0.00	0.03	0.04
On	Paraphrased	0.00	0.00	0.05	0.11
On	Same-text resample	0.06	0.05	0.07	0.28

Table 24: Wording discrepancies relative to base, across all three contexts. Excess subtracts the identical-text resampling discrepancy. Joint correctness requires both the base and edited reversal pair to be correct out of 80. Intervals crossing zero do not establish equivalence; agreement alone does not establish correctness.

Thinking	Edit	Discrepancy, pp	Excess over resample, pp	Both pairs correct
Off	Renamed	6.6 [5.6, 7.7]	0.1 [-0.9, 1.2]	5/80
Off	Paraphrased	7.3 [6.1, 8.5]	0.8 [-0.4, 2.1]	3/80
Off	Same-text resample	6.4 [5.5, 7.5]	0.0 [0.0, 0.0]	6/80
On	Renamed	0.2 [0.1, 0.4]	-0.3 [-0.6, -0.1]	76/80
On	Paraphrased	0.3 [0.1, 0.5]	-0.2 [-0.5, 0.0]	71/80
On	Same-text resample	0.5 [0.3, 0.9]	0.0 [0.0, 0.0]	72/80

One frozen frontier evaluation. A prospectively fixed balanced subset contains three instances per mechanism class (24 parents), with base and paraphrase wording. It has 576 planned records, not 48 independent families. GPT-5.6 Sol uses low reasoning, a 2,048-output-token allowance, strict probability JSON, no tools, and zero automatic retries. Batch baselines precede a separate update batch using only their own saved priors. The complete cohort was reserved within the user’s budget before generation; provider token-count receipts and the persistent ledger are archived. No released reservation financed additional target cases. Local rows below use exactly this same subset. Frontier/local contrasts remain descriptive because deployments and reasoning allowances differ.

GPT-5.6 returned 576/576 valid final probabilities.

Table 25: Matched 24-parent subset, with all planned failures retained. Signs are out of 48; paired reversals out of 24. Pair intervals are percentages; broken movement and its interval are percentage points.

Deployment	Wording	Signs	Pairs [95% CI]	Broken $ \Delta $, pp	Observed
Qwen3 off	Base	26/48	4/24 [4.2, 33.3]	2.8 [1.4, 4.3]	24/24
Qwen3 off	Paraphrased	24/48	3/24 [0.0, 29.2]	2.6 [1.2, 4.2]	24/24
Qwen3 on	Base	46/48	22/24 [79.2, 100.0]	0.5 [0.0, 1.6]	24/24
Qwen3 on	Paraphrased	47/48	23/24 [87.5, 100.0]	0.0 [0.0, 0.0]	24/24
GPT-5.6	Base	48/48	24/24 [100.0, 100.0]	0.0 [0.0, 0.0]	24/24
GPT-5.6	Paraphrased	48/48	24/24 [100.0, 100.0]	0.0 [0.0, 0.0]	24/24

Table 26: Base-wording paired reversal by mechanism class. The equal-class macro average equals the overall rate because all classes contain ten instances. These repeated parameterizations do not supply 80 independent mechanism replications.

Thinking	Mechanism class	Base pairs correct
Off	bounded feedback	3/10
Off	broker allocation	3/10
Off	calibration	1/10
Off	directed routes	1/10
Off	evidence dependence	3/10
Off	procedural gates	2/10
Off	quorums	1/10
Off	time windows	0/10
On	bounded feedback	10/10
On	broker allocation	10/10
On	calibration	10/10
On	directed routes	9/10
On	evidence dependence	8/10
On	procedural gates	10/10
On	quorums	10/10
On	time windows	10/10

What a relation-blind predictor can reach. The unit of evidence is the triple: one parent instance and wording, with the evidence text, the prior and the named actor identical across three stated relations. All 320 triples were verified to hold those three constant, to cover the three relations, and to have three distinct correct posteriors. The last property bounds every shortcut: a predictor reading only the evidence, the prior and the actor is constant within a triple, so it matches at most one of three items and never a whole triple. The bound is arithmetic, not empirical, and the predictors below only exhibit it – including an adversary allowed to choose the best constant revision per triple knowing every oracle.

Table 27: Predictors that read only the evidence, the prior and the named actor. Each is constant within a triple, so none can exceed one third of items or win a single triple.

Predictor	Items, direction	Triples, all three
Predict no change	320/960	0/320
Follow the evidence direction	320/960	0/320
Always revise upward	320/960	0/320
Best constant, given every oracle	320/960	0/320
<i>Analytic bound</i>	$\leq 1/3$	0

Exact agreement and the reasoning contrast. Scoring exact posterior agreement rather than revision direction separates systems that a binary criterion cannot: a guess supplies a sign, not a value. Denominators are unconditional, counting refusals, truncations and updates blocked by their own failed baseline as failures.

Table 28: Triple and item accuracy against the exact oracles, with unconditional denominators. Exact agreement is to four decimals, the granularity at which the oracles are stated.

System	Triples	Items, direction	Items, exact	Posterior error, pp
GPT-5.6 Sol	48/48	144/144	144/144	0.000
Claude Opus 5	17/48	78/144	54/144	0.013
Qwen3-32B thinking off	29/320	519/960	1/960	28.220
Qwen3-32B thinking on	297/320	937/960	811/960	0.129

Table 29: Every planned call, classified by attribution. A refusal or unparseable answer is the system’s; a truncation at the output budget or a request killed mid-batch is ours; a record blocked by its own failed baseline is neither.

System	Planned	Valid	Model	Instrument	Cascade
GPT-5.6 Sol	576	576	0	0	0
Claude Opus 5	576	339	84	0	153
Qwen3-32B thinking off	3840	3840	0	0	0
Qwen3-32B thinking on	3840	3803	5	5	27

Interpretive boundary and audit. These data address the original relevance shortcut under explicit finite rules: news-only, full-word, and word-bigram classifiers all achieve 1/3 direction accuracy and .50 balanced relevance accuracy in five-fold evaluation holding out entire parent instances and all their variants. This does not eliminate every possible structural heuristic. The new diagnostic does not retroactively validate the original packet relevance comparison, establish natural-news coverage, or replace the limitations of model screening. Raw completions, context-specific priors, effective chat hashes, token counts, failure denominators, and primary statistics were checked before these tables were generated. The artifact root is `expl_prospective/context_reversal/results/fresh_evaluation_v1`; the immutable design is under `expl_prospective/context_reversal/data/fresh_evaluation_v1/frozen_v1`.

B.9 EXPLORATORY PROBE: IS THE SATURATION A PROPERTY OF THE ITEMS?

The frozen cohort separates reasoning from no reasoning and little else at the top: two frontier deployments answer every item they leave intact to four decimals. That ceiling admits two readings. The items may be too easy, in which case harder ones would restore discrimination, or the design family may itself be saturating, in which case they would not. We ran a development probe to tell these apart before committing to a second frozen cohort. It is unscreened, unfrozen, and carries no human validation; nothing here supports a claim about model capability, and it is reported because its outcome changed what we built.

What was varied. The relation in this cohort is an actor-to-channel binding that the prompt states outright. We held the arithmetic fixed — the same mechanisms, the same outcome wrapper, the same exact oracles — and varied only how far that binding sits from the surface: stated directly; bound through permits; bound through permits and offices; and that, plus inert relations naming each actor beside a channel it does not operate. The inert relations take the cyclic successor of an actor’s own channel, which is a derangement, so no actor is named beside the channel it runs and each channel still appears exactly once. Every rung keeps the properties the triple analysis depends on: within a triple the evidence text, the prior, and the named actor are identical, the three correct posteriors are distinct, and the word and punctuation inventories are identical across the three contexts, so the added layers give a bag-of-words reader nothing. Sixteen parent instances per rung, two per mechanism class.

Table 30: Development calibration of binding depth. Qwen3-32B with thinking enabled, 384/384 records and no failures. Triples require all three relations correct. Error is mean absolute distance from the exact posterior. Sixteen triples per rung does not support inference about small differences; the relevant observation is the absence of a large one.

Binding	Triples	Signs	Exact	Error (pp)
Direct statement	16/16	48/48	44/48	0.011
One step, via permits	15/16	46/48	44/48	0.007
Two steps, via permits and offices	16/16	48/48	45/48	0.050
Two steps, plus inert relations	14/16	46/48	41/48	0.320

The probe came back negative. Table 30 shows the direct rung reproducing the frozen cohort, which uses that form and which this deployment wins at 297/320, so the instrument measures what it should. Moving the binding two inferential steps from the surface and adding the inert relations

costs two triples of sixteen, which at this sample size is noise. Mean posterior error rises from 0.011 to 0.320 pp, a thirtyfold increase that remains far below any threshold an evaluation could use.

We therefore did not build a second frozen cohort on this axis. The probe also suggests why the ceiling is difficult to move from inside this design. Exact oracles require the scenario to be fully specified, and full specification is what makes the task tractable for a deployment that can already enumerate; indirection then lengthens what must be read without adding anything that must be reasoned about. A discriminating axis would have to introduce something that survives full specification, or abandon exact oracles — and the oracles are what license the bound of Table 27, which is the strongest claim this experiment makes. We report the ceiling rather than weaken the bound to move it. Materials, runner, and scorer are under `expl_prospective/context_reversal/ladder`.

C EXPERIMENT 2: CONTEXTUAL RELATIONSHIP SELECTION AND REVISION IN COIN CITY

Experiment 2 tests whether context selects which response relationship a model applies and whether direct target observations revise that selection. The six primary deployments and symbol-control comparisons reuse the same 250 seeded episodes, so prompt-arm differences are paired rather than comparisons between independently sampled tests.

C.1 DESIGN AND DATA-GENERATING PROCESS

C.1.1 DESIGN SUMMARY

The design identifier is `coin_city_stable_relationship_claude_n250_v4`. It contains 250 independent episodes generated from seeds 930000–930249, five target-city prefixes $k \in \{0, 1, 2, 3, 4\}$, and three semantic-selection arms. Cue-inversion and arbitrary-symbol modules reuse those episode identities and numerical inputs; each has a hashed task manifest (Section C.3). Every model received the same episode records, queries, and arm-specific prompt text.

C.1.2 BALANCED LATENT REGIMES

Each episode crosses two binary factors: whether City A or City B is the strongly responsive reference city, and whether target City C is strongly or weakly responsive. The closest-balanced allocation contains 63 A-strong/C-strong, 62 A-strong/C-weak, 62 B-strong/C-strong, and 63 B-strong/C-weak episodes. Within an episode, each city receives an independently sampled stable slope

$$g_{ec} \sim \begin{cases} \mathcal{N}(0.90, 0.03^2), & \text{strong regime,} \\ \mathcal{N}(0.25, 0.03^2), & \text{weak regime.} \end{cases}$$

Thus the cue identifies a response regime but not the realized coefficient.

C.1.3 OBSERVED CASES AND FORECAST QUERY

Every city has four reference cases arranged as two matched pairs. Rows in a pair share a starting poll and news magnitude but have opposite signs; the first sign is random. For reference row r ,

$$p_{ecr}^{\text{end}} = \text{round}_{0.1} (p_{ecr}^{\text{start}} + g_{ec}x_{ecr} + \epsilon_{ecr}), \quad \epsilon_{ecr} \sim \mathcal{N}(0, 5^2).$$

Starting polls are uniform on $[35, 65]$, and reference-news magnitudes are drawn from $\{6, 7, 8, 9, 10\}$. The query shock is drawn from $\{-10, \dots, -4, 4, \dots, 10\}$. Its scoring target is the latent conditional expectation $p_0 + g_C x$, with no additional query-time observation noise.

The prefix k exposes the first k City C rows in their stored order. Consequently, $k = 2$ completes the first matched pair and $k = 4$ completes both; $k = 1$ and $k = 3$ deliberately expose partial pairs. City A and B always contribute all four rows in the ABC arms.

Table 31: Coin City design and evaluation dimensions. Final analytic record counts retain one record per task cell; raw transport-level calls are discussed in Section C.9.

Quantity	Value
Independent episodes	250
Factorial cells	63, 62, 62, 63 across the four A/B-by-C regimes
Target-city prefixes	$k \in \{0, 1, 2, 3, 4\}$
Primary prompt arms	target only; ABC/no C context; ABC/correct C context
Intervention modules	ABC/misleading C context; ABC/episode-randomized arbitrary labels
Primary deployment roster	Claude Opus 4.8; GPT-5.6; DeepSeek-V4-Pro; Kimi K3; Gemini 3.6 Flash; Claude Opus 5
Symbol-comparison roster	Qwen3 4B/8B/14B/32B; Qwen2.5 7B/14B/32B/72B; Llama 3.1 8B/70B; DeepSeek-V4-Pro; Kimi K3; Claude Opus 4.8/5; GPT-5.6 Sol
Primary task cells / final records	22,500 / 22,500 (22,497 parseable)
Misleading task cells / final records	7,500 / 7,500 (all parseable)
Symbol-comparison task cells / response records	56,250 / 56,250 (15 models)
Strong / weak slope means	0.90 / 0.25 poll points per news point
Slope standard deviation	0.03 within each regime
Reference-case noise	5 poll points
Reference cases per city	4 (two matched pairs)

C.1.4 PREFLIGHT VALIDATION AND TASK INTEGRITY

The GPT-5.6, DeepSeek-V4-Pro, Kimi K3, Gemini 3.6 Flash, and Claude Opus 5 replications reuse the exact task files and prompt hashes created for the original Claude Opus 4.8 evaluation; no episode was regenerated. Before model calls, the manifest hashed the engine, task files, answer and scoring keys, validator, runner, renderer, and launcher. Automated validation confirmed the intended dimensions and factorial balance, prompt contracts and hashes, correct cue-to-regime mapping, visible matched pairs, and benchmark reconstruction. No model response or performance statistic was used by this preflight validator.

C.2 PROMPT ARMS AND INFORMATION AVAILABLE

C.2.1 SHARED TASK INSTRUCTIONS

All prompts state that rows are separate polling cases rather than a time series, positive news favors the candidate, poll change is end minus start, and within-city responsiveness is stable even though individual polls are noisy. Rows sharing a pair label are described as having the same starting poll and equal-magnitude, opposite-sign news. Every displayed row includes the starting poll, news value, change, and ending poll.

The required answer is a JSON object with exactly a one-sentence `rationale` and a numeric `predicted_poll` between 0 and 100. The prompt asks for no step-by-step derivation.

C.2.2 FIVE INFORMATION CONDITIONS

Target only. The model sees the first k City C cases and is explicitly told that no additional background information is available. This arm measures recovery from target-city evidence alone.

ABC, no City C context. The model also sees four cases from each reference city, each introduced by a background sentence. The strongly responsive reference city is described verbatim as receiving “national news coverage, where campaign stories receive sustained attention and tend to produce larger immediate poll movements”, and the weakly responsive city as receiving “local news coverage, where campaign stories compete with local issues and tend to produce smaller immediate poll movements”. The direction of the coverage-to-response mapping is therefore stated in words, though no numerical value is. City C receives no background description. This arm tests numeric transfer without a cue that aligns C to either reference regime.

ABC plus correct City C context. The numeric rows and query are identical to the preceding arm. The sole added information is that City C receives national or local coverage, correctly aligned with its latent strong or weak regime. This is the primary cue-transfer condition.

ABC plus misleading City C context. This intervention arm is identical to the correct-context prompt except that its City C sentence names the opposite coverage regime. Every number, case row, query, and other instruction is unchanged. This arm tests whether the model follows the semantic cue when it conflicts with the truth.

ABC plus an arbitrary City C label. Every national/local description is replaced by one of two arbitrary labels, KIV or ZOR. The prompt states only that cities sharing a label share a typical response regime and that the labels have no inherent meaning or ordering. Their higher/lower-response mapping reverses on alternating episodes; City A and B’s numerical cases are therefore the only evidence from which the episode-local mapping can be induced before applying City C’s label.

No prompt contains a fitted coefficient, an estimator, the strong or weak numerical means, the case-noise scale, a rule for pooling cities, or any analyst benchmark. The correct-versus-no-context contrast therefore isolates the one City C background sentence while holding all displayed numeric evidence fixed; the correct-versus-misleading comparison isolates the sentence’s regime assignment. The symbol-versus-no-context contrast asks whether the same assignment can be induced without lexical knowledge of national and local coverage.

Because the reference sentences state the direction of the mapping, the semantic arms are not by themselves a test of whether that direction can be induced. A policy that reads no displayed number and assigns slope 1.0 to the nationally covered regime and 0.2 to the locally covered one attains $k = 0$ forecast MAE 0.523 and implied-slope correlation .996, better than any deployment and better than the analogical-regression benchmark of Table 33; 42 of 108 prior pairs on the grid $\{0.4, \dots, 1.5\} \times \{0.0, \dots, 0.8\}$ beat every deployment’s $k = 0$ MAE. The semantic arms therefore measure how a stated regime assignment is converted into a numerical response. Two separate controls carry the induction claim: the arbitrary-symbol arm, whose labels have no stated direction and whose mapping reverses across episodes, and the reference-selection analysis of Section C.6, which shows that the forecast tracks the named reference city’s own sampling error rather than any fixed prior.

C.3 MODEL EXECUTION AND RESPONSE PROCESSING

C.3.1 DEPLOYMENT CONFIGURATIONS

Table 32 reports the effective API settings used for production. Temperature refers to the value sent to the provider: GPT-5.6 and Claude Opus 5 omit it even though the local GPT run manifest records a default of zero. No deployment used tools or retrieval.

Table 32: Production deployment configuration and primary-arm parse counts.

Deployment	Interface	Reasoning / sampling	Output ceiling	Parsed
Claude Opus 4.8	Anthropic Messages via Azure	temperature 0	512	3,750
GPT-5.6	OpenAI Responses	low reasoning; temperature omitted	4,096	3,749
DeepSeek-V4-Pro	OpenAI Chat Completions	temperature 0	4,096	3,750
Kimi K3	OpenAI Chat Completions	low reasoning; temperature 0	16,384	3,750
Gemini 3.6 Flash	Gemini OpenAI-compatible Chat	low reasoning; temperature 0	16,384	3,750
Claude Opus 5	Foundry Anthropic Messages	default extended thinking; temperature omitted	4,096	3,748

Claude Opus 5 returns extended thinking before its answer; the runner therefore selects its first text block. Gemini is the sole fast-tier deployment, so the six systems should be read as a replication set rather than a tier-matched leaderboard. Authentication material was supplied at runtime and was not written to source, manifests, logs, or response files.

C.3.2 DEPLOYMENT ATTEMPTS, RETRIES, AND PILOTS

Every primary model–prompt pair received one application-level attempt. A failed or answerless task was not re-asked. HTTP-client retries were disabled for GPT, DeepSeek, Kimi, and Opus 5. The original Opus 4.8 client allowed up to two transport retries. Gemini’s runner allowed adaptive pre-model transport retries but recorded none in the completed shards. These transport policies do not alter the one-answer-per-task analytic rule.

Before production, 119 Kimi, 15 Gemini, and 17 Opus 5 task-shaped calls were used only to select configurations that reliably returned the required JSON schema. Forecast accuracy was not inspected. The calls are archived separately and are excluded from every denominator and analysis; each production namespace began empty. Production tasks were sharded only for throughput.

C.3.3 RESPONSE PARSING AND RECORD SELECTION

The parser scans the raw reply for JSON objects, retains the last object with an exact `predicted_poll` key whose value converts to a number in $[0, 100]$, truncates the retained rationale to 500 characters, and stores at most 8,000 raw characters. When duplicate task records exist, the analysis keeps the latest parseable response, allowing a valid response to supersede an earlier error. A single offline normalization removed whitespace from JSON keys in one otherwise valid Opus 4.8 reply; this recovered one primary record without another model call. Other record-level corrections are documented in Section C.9.

C.3.4 MISLEADING-CONTEXT DESIGN

The cue-inversion module begins from the frozen semantic-selection tasks. Before any inversion call, its builder verifies that every prompt equals its correct-context counterpart under exactly one sentence substitution and that all numeric content is unchanged. The cue is inverted for all 1,250 episode-prefix tasks: 625 strong-city tasks are described as weak and 625 weak-city tasks as strong. Each deployment uses the same client and production configuration as its correct-context arm.

C.3.5 ARBITRARY-SYMBOL DESIGN

The symbol builder generates 1,250 matched prompts from the same frozen episodes and verifies that their complete sequence of numeric tokens is identical to the correct-context arm. The higher-response label is KIV in 125 episodes and ZOR in 125; the target regime-by-label cells contain 62–63 episodes each. A prompt validator rejects national/local and strong/weak semantic descriptions. At $k = 0$, City C supplies no numerical cases, so its correct regime can be recovered only by decoding the label from City A and B’s examples.

For DeepSeek-V4-Pro, we collected all 1,250 symbol tasks and contemporaneously re-collected all 2,500 no-context and semantic-context tasks on the same reachable Azure resource, using one deployment identifier and configuration. The released 15-model comparison combines that triplet with Qwen3-4B, Qwen3-8B, Qwen3-14B, Qwen3-32B, Qwen2.5-7B, Qwen2.5-14B, Qwen2.5-32B, Qwen2.5-72B, Llama 3.1-8B, Llama 3.1-70B, Kimi K3, Claude Opus 4.8, Claude Opus 5, and GPT-5.6 Sol.

The eight Qwen and two Llama checkpoints received all 3,750 prompts through vLLM in bfloat16, with temperature-zero decoding, a 3,072-token model context, and a 512-token output ceiling. Qwen uses the no-thinking chat template; Llama uses its native instruction template. Qwen2.5-14B uses two-way and Qwen3-32B and Qwen2.5-32B four-way tensor parallelism, while Qwen2.5-72B and Llama 3.1-70B use eight-way tensor parallelism.

For Kimi K3, Claude Opus 5, and GPT-5.6 Sol, the comparison reuses each deployment’s preserved no-context and semantic-context responses and adds only the 1,250 symbol prompts under the same model-specific production protocol. Claude Opus 4.8’s symbol arm instead ran on a second Azure resource serving the same deployment and used a 1,024-token rather than 512-token ceiling; all 1,250 responses returned valid JSON. We therefore treat its arm contrast as a cross-run robustness result and do not use it to rank providers.

Every runner verifies the frozen task hash before generation and writes a timestamped manifest. Each released model has 1,250 unique response records in every arm; this count does not imply that

every prediction parsed. No primary response file or frozen task is overwritten. Original outputs are preserved, and the final tables use matched complete cases after the no-call arithmetic reparse described in Section C.9.

C.4 ESTIMANDS, BENCHMARKS, AND STATISTICAL INFERENCE

C.4.1 FORECAST AND LATENT-RESPONSE ESTIMANDS

Forecast error is measured against the latent conditional expectation rather than a newly sampled noisy ending poll:

$$\text{MAE} = \frac{1}{n} \sum_{e=1}^n |\hat{p}_e - (p_{0e} + g_{Ce}x_e)|.$$

For every nonzero query shock, a forecast also identifies an implied slope $\hat{g}_e = (\hat{p}_e - p_{0e})/x_e$. We report its mean absolute error from g_{Ce} and its Pearson correlation with g_{Ce} across episodes. Correlation is needed because a nearly constant prior-mean slope can attain moderate MAE while failing to distinguish strong from weak cities.

Within each deployment and prefix, the three primary-arm metrics use the common intersection of episodes with parseable responses in all three arms. Arm contrasts are paired by episode and prefix. The misleading-arm table uses its final unique parseable record for every task. The arbitrary-symbol analysis uses each model’s corresponding common intersection across its three matched arms.

C.4.2 ANALYST-ONLY BENCHMARKS

The ABC OLS benchmark fits poll change on news through the origin using every displayed A, B, and available C row. The target-only OLS benchmark uses only displayed City C rows and is undefined at $k = 0$. Both are reconstructed from the visible data by the analyst; neither the estimator nor its output appears in a model prompt.

The label-conditioned analogical regression implements the minimal structured inductive policy defined in Methods. It fits separate through-origin slopes to City A and City B, uses City C’s coverage label to select the reference assigned to the same response regime, and predicts with that slope. At $k > 0$, it pools the selected reference rows with the visible City C rows before refitting. In the misleading arm, the inverted label selects the opposite reference. The implementation uses only displayed rows; neither the estimator nor its output appears in a model prompt.

Table 33: Zero-shot behavior of the analogical policy and the six model deployments. At $k = 0$, City C has no numerical cases. Model columns give the range across deployments; forecast MAE is in poll points.

Cue	Policy MAE	Policy ρ	Model MAE range	Model ρ range
Correct	1.842	.656	1.596–2.474	[.52, .70]
Inverted	4.561	–.667	4.204–4.361	[–.68, –.50]

C.4.3 UNCERTAINTY AND COMPARISON FAMILIES

Primary arm contrasts use percentile intervals from 2,000 bootstrap resamples of the paired episode-level absolute-error difference. Deterministic, model-specific seeds make these intervals reproducible. The $k = 0$ pairwise model table uses 4,000 paired bootstrap resamples with seed 20260812 and the common episode intersection for each model pair. Negative differences favor the first-listed condition or model.

The episodes are the resampling units. Intervals are pointwise descriptions of this fixed seeded synthetic evaluation: they are not simultaneous bands across prefixes, and no multiple-comparison correction is applied. Accordingly, the number of pairwise intervals excluding zero is not used as an estimand; each interval supports only its displayed cell.

C.5 RESPONSE-INFERENCE DIAGNOSTICS

The tables below retain a common arm and prefix ordering. Within a deployment and prefix, all three primary arms use the same complete-case episode set; n is repeated so that every exception is visible. Boldface identifies the lowest forecast MAE or slope MAE, and the highest slope correlation, within that deployment and prefix.

C.5.1 PER-DEPLOYMENT PRIMARY-ARM RECOVERY

Table 34: Claude Opus 4.8 forecast and latent-response diagnostics for every prompt arm. Forecast MAE is in poll points; slope MAE and ρ compare the implied response with the true City C slope.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	2.321	.332	.019
	ABC, no C context	250	2.651	.379	-.063
	ABC + C context	250	1.817	.265	.700
1	Target only	250	3.730	.549	.371
	ABC, no C context	250	2.438	.358	.322
	ABC + C context	250	2.147	.316	.584
2	Target only	250	2.668	.403	.383
	ABC, no C context	250	2.041	.296	.492
	ABC + C context	250	2.079	.308	.491
3	Target only	250	2.485	.358	.449
	ABC, no C context	250	2.128	.305	.513
	ABC + C context	250	1.977	.291	.627
4	Target only	250	2.209	.324	.498
	ABC, no C context	250	1.929	.281	.542
	ABC + C context	250	1.760	.257	.658

Table 35: GPT-5.6 forecast and latent-response diagnostics on the same frozen prompts. Metrics use episodes parseable in all three arms at each prefix.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	3.596	.511	.054
	ABC, no C context	250	2.518	.352	-.041
	ABC + C context	250	1.596	.224	.678
1	Target only	249	2.993	.439	.394
	ABC, no C context	249	2.768	.399	.302
	ABC + C context	249	1.586	.225	.669
2	Target only	250	2.560	.387	.459
	ABC, no C context	250	2.050	.299	.534
	ABC + C context	250	1.854	.268	.606
3	Target only	250	2.240	.328	.528
	ABC, no C context	250	1.967	.286	.584
	ABC + C context	250	1.773	.256	.600
4	Target only	250	2.153	.318	.577
	ABC, no C context	250	1.746	.253	.653
	ABC + C context	250	1.594	.228	.664

Table 36: DeepSeek-V4-Pro forecast and latent-response diagnostics on the same frozen prompts. All 250 episodes parsed in every arm and prefix.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	2.447	.353	-.027
	ABC, no C context	250	2.841	.401	-.027
	ABC + C context	250	2.474	.344	.523
1	Target only	250	3.585	.525	.334
	ABC, no C context	250	3.130	.458	.302
	ABC + C context	250	2.456	.347	.578
2	Target only	250	3.380	.501	.319
	ABC, no C context	250	3.423	.502	.365
	ABC + C context	250	2.687	.391	.492
3	Target only	250	2.908	.424	.454
	ABC, no C context	250	3.224	.463	.403
	ABC + C context	250	2.797	.403	.504
4	Target only	250	2.812	.415	.430
	ABC, no C context	250	3.040	.459	.366
	ABC + C context	250	2.809	.409	.470

Table 37: Kimi K3 forecast and latent-response diagnostics on the same frozen prompts. All 250 episodes parsed in every arm and prefix.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	3.352	.476	-.016
	ABC, no C context	250	2.548	.357	-.021
	ABC + C context	250	1.886	.266	.593
1	Target only	250	3.362	.494	.365
	ABC, no C context	250	3.370	.490	.345
	ABC + C context	250	2.382	.341	.579
2	Target only	250	3.194	.491	.387
	ABC, no C context	250	3.114	.487	.352
	ABC + C context	250	2.306	.349	.455
3	Target only	250	2.816	.412	.430
	ABC, no C context	250	2.636	.407	.410
	ABC + C context	250	2.005	.292	.609
4	Target only	250	2.378	.357	.444
	ABC, no C context	250	1.957	.284	.622
	ABC + C context	250	1.840	.263	.601

Table 38: Gemini 3.6 Flash forecast and latent-response diagnostics on the same frozen prompts. All 250 episodes parsed in every arm and prefix.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	2.534	.370	.002
	ABC, no C context	250	2.552	.356	.078
	ABC + C context	250	2.174	.302	.559
1	Target only	250	3.733	.550	.403
	ABC, no C context	250	3.717	.548	.361
	ABC + C context	250	3.300	.500	.292
2	Target only	250	3.564	.558	.348
	ABC, no C context	250	3.703	.579	.335
	ABC + C context	250	3.600	.569	.330
3	Target only	250	3.370	.513	.314
	ABC, no C context	250	3.092	.460	.397
	ABC + C context	250	3.047	.443	.471
4	Target only	250	3.308	.498	.389
	ABC, no C context	250	2.911	.447	.433
	ABC + C context	250	2.793	.419	.447

Table 39: Claude Opus 5 forecast and latent-response diagnostics on the same frozen prompts. Two responses exhausted the output budget on extended thinking and returned no answer, one in each ABC arm.

k	Arm	n	Forecast MAE	Slope MAE	$\rho(\hat{g}, g_C)$
0	Target only	250	2.359	.333	-.018
	ABC, no C context	250	2.619	.371	-.007
	ABC + C context	250	1.764	.247	.617
1	Target only	249	3.549	.520	.391
	ABC, no C context	249	2.303	.338	.370
	ABC + C context	249	1.741	.254	.601
2	Target only	250	3.804	.596	.309
	ABC, no C context	250	2.680	.412	.389
	ABC + C context	250	2.060	.314	.524
3	Target only	250	3.022	.461	.379
	ABC, no C context	250	2.363	.356	.459
	ABC + C context	250	2.075	.309	.550
4	Target only	249	2.700	.418	.443
	ABC, no C context	249	2.369	.366	.459
	ABC + C context	249	2.201	.335	.499

C.5.2 MISLEADING-CUE RESULTS

Table 40: Misleading-cue arm: forecast MAE and latent-slope correlation at each City C prefix. The prompt is byte-identical to the correct-context arm except that the City C sentence names the opposite response regime. Negative ρ means the implied response is anti-correlated with the truth, i.e. the deployment has adopted the stated regime. Each of the six deployments contributes all 1,250 prompts.

Model	$k=0$		$k=1$		$k=2$		$k=3$		$k=4$	
	MAE	ρ	MAE	ρ	MAE	ρ	MAE	ρ	MAE	ρ
Claude Opus 4.8	4.20	-.66	3.88	-.51	3.04	.09	2.61	.22	2.31	.34
GPT-5.6	4.23	-.68	4.15	-.63	2.78	.17	2.71	.25	2.10	.46
DeepSeek-V4-Pro	4.36	-.62	4.46	-.48	3.61	.07	3.39	.14	3.04	.32
Kimi K3	4.31	-.62	4.12	-.25	3.03	.15	2.72	.14	2.18	.44
Gemini 3.6 Flash	4.29	-.50	4.40	-.06	3.93	.24	3.52	.25	2.93	.36
Claude Opus 5	4.25	-.62	3.15	-.21	2.59	.27	2.41	.41	2.36	.47

C.5.3 ARBITRARY-SYMBOL INDUCTION RESULTS

Table 41: **Fifteen-system episode-randomized arbitrary-symbol comparison at $k = 0$.** Panel A reports forecast calibration; Panel B reports response-regime discrimination. All arms share the frozen numeric task content; Section C.3 documents model-specific execution and the Opus 4.8 cross-resource deviation. MAE is in poll points; ρ correlates implied and true latent slopes; Acc. classifies the response regime at slope 0.575. Each Δ is symbol minus no context, with a paired 95% episode-bootstrap interval. With no City C examples, its label can be decoded only from City A/B examples.

Panel A: Forecast calibration

Model	n	No context	Semantic	Symbol	Δ MAE [95% CI]		
Qwen3-4B	250	2.896	2.365	2.892	−.005	[−.088, .074]	
Qwen2.5-7B	250	2.911	2.669	3.248	+.338	[.027, .682]	
Qwen3-8B	250	2.985	2.859	2.973	−.011	[−.239, .207]	
Qwen2.5-14B	250	2.606	2.548	2.711	+.105	[−.178, .388]	
Qwen3-14B	250	2.632	2.241	2.812	+.180	[−.188, .548]	
Qwen3-32B	250	2.482	2.069	2.914	+.432	[.119, .752]	
Qwen2.5-32B	250	2.679	2.404	3.208	+.529	[.169, .924]	
Qwen2.5-72B	250	2.474	2.383	3.398	+.924	[.582, 1.310]	
Llama 3.1 8B	242	3.113	3.386	3.605	+.492	[.029, .980]	
Llama 3.1 70B	250	3.483	3.761	3.931	+.448	[.030, .886]	
DeepSeek V4-Pro	250	2.622	2.588	2.729	+.107	[−.222, .452]	
Kimi K3	250	2.548	1.886	2.051	−.497	[−.765, −.211]	
Claude Opus 4.8	250	2.651	1.817	2.067	−.584	[−.858, −.314]	
Claude Opus 5	250	2.619	1.764	2.017	−.603	[−.869, −.336]	
GPT-5.6 Sol	250	2.518	1.596	1.744	−.774	[−1.011, −.533]	

Panel B: Response-regime discrimination

Model	n	$\rho(\hat{g}, g_C)$			Regime accuracy		
		No context	Symbol	$\Delta\rho$ [95% CI]	No context	Symbol	Δ Acc. [95% CI]
Qwen3-4B	250	-.011	.085	+.095 [−.060, .230]	.496	.512	+.016 [−.008, .044]
Qwen2.5-7B	250	-.034	.230	+.264 [.112, .406]	.500	.572	+.072 [.020, .128]
Qwen3-8B	250	-.014	.009	+.023 [−.128, .171]	.492	.516	+.024 [−.012, .060]
Qwen2.5-14B	250	-.029	.208	+.236 [.107, .356]	.492	.580	+.088 [.020, .152]
Qwen3-14B	250	-.057	.342	+.398 [.253, .541]	.496	.660	+.164 [.088, .240]
Qwen3-32B	250	.075	.377	+.302 [.177, .433]	.524	.656	+.132 [.060, .200]
Qwen2.5-32B	250	-.015	.358	+.372 [.226, .502]	.500	.648	+.148 [.088, .208]
Qwen2.5-72B	250	.027	.408	+.381 [.250, .508]	.516	.632	+.116 [.056, .176]
Llama 3.1 8B	242	-.042	.003	+.045 [−.123, .250]	.467	.512	+.045 [−.029, .120]
Llama 3.1 70B	250	-.039	.385	+.424 [.314, .543]	.500	.592	+.092 [.036, .152]
DeepSeek V4-Pro	250	.018	.454	+.436 [.323, .558]	.520	.692	+.172 [.104, .240]
Kimi K3	250	-.021	.612	+.633 [.520, .742]	.468	.804	+.336 [.272, .400]
Claude Opus 4.8	250	-.063	.560	+.623 [.492, .744]	.448	.760	+.312 [.240, .380]
Claude Opus 5	250	-.007	.583	+.591 [.485, .697]	.496	.780	+.284 [.220, .348]
GPT-5.6 Sol	250	-.041	.655	+.696 [.597, .807]	.464	.836	+.372 [.312, .432]

Table 42: **Matched open-weight scaling across target-city observations.** Symbol-minus-no-context effects are reported for Qwen2.5 32B/72B and Llama 3.1 8B/70B at each City C prefix. Negative ΔMAE favors the induced symbol mapping; positive $\Delta\rho$ and $\Delta\text{accuracy}$ indicate better regime discrimination.

Model	k	n	ΔMAE [95% CI]	$\Delta\rho$ [95% CI]	$\Delta\text{Acc.}$ [95% CI]
Qwen2.5-32B	0	250	+.529 [.169, .924]	+.372 [.226, .502]	+.148 [.088, .208]
	1	250	+.124 [−.188, .441]	+.096 [.010, .190]	+.068 [.016, .124]
	2	250	+.184 [−.044, .421]	−.010 [−.059, .037]	−.012 [−.048, .024]
	3	250	+.462 [.222, .710]	−.033 [−.096, .024]	−.004 [−.044, .036]
	4	250	+.255 [−.001, .524]	−.010 [−.059, .033]	−.016 [−.048, .016]
Qwen2.5-72B	0	250	+.924 [.582, 1.310]	+.381 [.250, .508]	+.116 [.056, .176]
	1	250	+.189 [−.148, .520]	+.166 [.075, .265]	+.076 [.024, .128]
	2	250	−.024 [−.286, .239]	+.070 [.017, .127]	+.016 [−.028, .060]
	3	250	+.252 [−.050, .549]	+.019 [−.051, .092]	+.024 [−.024, .068]
	4	250	+.019 [−.280, .322]	+.011 [−.067, .088]	+.020 [−.024, .064]
Llama 3.1 8B	0	242	+.492 [.029, .980]	+.045 [−.123, .250]	+.045 [−.029, .120]
	1	240	−2.318 [−3.256, −1.373]	+.120 [−.022, .263]	−.050 [−.129, .029]
	2	238	−2.454 [−3.433, −1.529]	+.106 [−.038, .260]	−.017 [−.097, .055]
	3	244	−.867 [−1.739, .085]	−.088 [−.238, .061]	−.078 [−.156, .000]
	4	246	−1.396 [−2.321, −.489]	+.100 [−.067, .245]	−.049 [−.122, .024]
Llama 3.1 70B	0	250	+.448 [.030, .886]	+.424 [.314, .543]	+.092 [.036, .152]
	1	250	+.194 [−.245, .658]	+.077 [−.026, .190]	−.016 [−.076, .040]
	2	250	−.200 [−.548, .149]	+.039 [−.026, .106]	−.016 [−.068, .036]
	3	250	−.345 [−.621, −.075]	+.037 [−.015, .087]	+.020 [−.016, .056]
	4	250	−.079 [−.306, .158]	+.034 [−.006, .082]	+.012 [−.020, .048]

At $k = 0$, Qwen3 provides a within-family contrast. Symbol-minus-no-context $\Delta\rho$ is .095 [−.060, .230] at 4B and .023 [−.128, .171] at 8B, but .398 [.253, .541] at 14B and .302 [.177, .433] at 32B; the regime-accuracy contrasts follow the same pattern. The 4B and 8B intervals span zero, whereas the 14B and 32B intervals do not. The 32B point estimate sits below the 14B one, so this is a boundary among the tested Qwen3 checkpoints rather than a monotone or universal parameter threshold. The smaller checkpoints are not wholesale task failures: their semantic-context slopes correlate .370 and .286 with truth, versus −.011 and −.014 without context.

All four tested Qwen2.5 checkpoints discriminate under arbitrary labels ($\Delta\rho = .264/.236/.372/.381$ at 7B/14B/32B/72B), reinforcing the family qualification. Their symbol-minus-no-context MAE changes are +.338/+.105/+.529/+.924: mapping discrimination does not ensure calibrated numerical use. Direct City C examples make the large-Qwen symbol advantage redundant; by $k = 3-4$, the 32B and 72B $\Delta\rho$ and $\Delta\text{accuracy}$ intervals span zero.

Llama 3.1 provides a second within-family contrast: $\Delta\rho = .045 [−.123, .250]$ at 8B and .424 [.314, .543] at 70B. The 70B regime-accuracy contrast is .092 [.036, .152], although its MAE rises by .448 [.030, .886]. In Qwen3 the mapping-recovery interval excludes zero only from 14B upward, and in Llama only at 70B. The positive Qwen2.5-7B and Qwen2.5-14B estimates, despite Qwen3-8B’s interval spanning zero, indicate that the tested transition depends on model family rather than parameter count. Mapping recovery and numerical calibration remain separate outcomes.

DeepSeek improves discrimination ($\Delta\rho = .436 [.323, .558]$) without a detectable zero-shot MAE change. Kimi, both Opus deployments, and GPT combine mapping recovery with lower forecast error: their $\Delta\rho$ values are .633/.623/.591/.696, while MAE falls by .497/.584/.603/.774 points. Cross-provider magnitudes remain descriptive because runtimes, reasoning settings, and output ceilings differ.

C.5.4 PAIRED CUE EFFECTS AND INTERPRETATION

Table 43: Paired forecast-MAE difference between the correct-context arm and the context-blind ABC OLS benchmark, which the analyst fits through the origin on every displayed A, B and available City C row. Negative values favour the deployment. Boldface marks intervals excluding zero.

Model	$k=0$	$k=1$	$k=2$	$k=3$	$k=4$
Claude Opus 4.8	-.73	-.18	-.05	-.01	-.08
GPT-5.6	-.95	-.74	-.28	-.21	-.24
DeepSeek-V4-Pro	-.07	.13	.56	.81	.97
Kimi K3	-.66	.06	.18	.02	.01
Gemini 3.6 Flash	-.37	.98	1.47	1.06	.96
Claude Opus 5	-.78	-.57	-.07	.09	.36

At $k = 0$ every deployment matches or beats the context-blind benchmark, and five of the six intervals exclude zero: with no City C cases, the only way to separate City C from the pooled A/B fit is to use its description. The advantage does not persist once City C supplies its own rows. By $k = 2$ DeepSeek-V4-Pro and Gemini 3.6 Flash are significantly worse than the benchmark, Claude Opus 5 becomes so at $k = 4$, and the remaining deployments are indistinguishable from it. The cue effect this experiment measures is therefore a zero-shot phenomenon, and a plain through-origin estimator remains competitive once direct target evidence accumulates.

In the zero-shot matched-reference comparison, the target-city description changes the influence of the A/B cases on forecasts. All six $k = 0$ bootstrap intervals exclude zero. A label-conditioned benchmark that selects a reference, estimates its slope, and transfers that estimate reproduces the aggregate pattern; the comparison does not identify the model’s internal algorithm. More target rows do not monotonically reduce error: DeepSeek and Opus 5 drift by 0.34 and 0.44 points from $k = 0$ to $k = 4$, and Gemini’s context advantage is significant only at $k \in \{0, 1\}$. Short, noisy prefixes can temporarily mislead either arm.

C.6 DOES THE FORECAST TRACK THE CUE-NAMED REFERENCE?

Two accounts explain the correct-versus-no-context contrast equally well. Under a lexical-prior account the model maps the City C coverage sentence straight to a remembered response magnitude and never reads City A or City B’s numbers. Under a reference-selection account it uses the cue to choose a reference city, estimates that city’s slope from its four displayed cases, and transfers the estimate. Both predict a regime-level correlation, because the cue and the latent regime coincide.

The design separates them. A through-origin fit to four noisy reference cases has a within-regime standard deviation of .332 against a generator standard deviation of .03, so the matched reference’s fitted slope correlates only .656 with the true City C slope where the regime indicator correlates .996. That sampling error is the discriminating signal: a forecaster that is deterministic given the cue cannot reproduce it, whereas one that transfers a fitted reference slope inherits it. For every deployment, arm and prefix we therefore report

$$r(\hat{g}_e, \hat{\beta}_e^{\text{ref}} \mid R_e),$$

the partial correlation between the forecast-implied slope and a reference city’s through-origin OLS slope, controlling for the response regime R_e that the City C label names. Arms with no City C label are scored against the true regime and report the regime-matched and regime-mismatched reference instead. Intervals are 2,000-draw episode bootstraps with deterministic per-cell seeds, and the analysis makes no model call.




Figure 13: **The cue selects one of the two displayed references.** Partial correlation between each deployment’s forecast-implied City C slope and a reference city’s own fitted slope at $k = 0$, controlling for the response regime the City C sentence names; bars are 95% episode-bootstrap intervals. Orange is the reference that sentence points to, gray the other. With no reference cases displayed neither loads; with both displayed and no cue they load equally, the signature of pooling; the correct cue concentrates loading on the reference it names; and an inverted cue moves that loading onto the reference it falsely names. A prior attached to the words *national* and *local* predicts no loading in any panel. Per-deployment values and the arbitrary-symbol version are in Appendix C.6.

Table 44: Partial correlation between the forecast-implied City C slope and each reference city’s fitted slope at $k = 0$, controlling for the regime the City C label names. “Named” is the reference the label points to, or the regime-matched reference in the two arms without a label. Boldface marks named-reference intervals excluding zero.

Model	Target only		ABC, no label		ABC + correct		ABC + inverted	
	Named	Other	Named	Other	Named	Other	Named	Other
Claude Opus 4.8	-.08	-.02	.39	.53	.62	.07	.71	.04
GPT-5.6	.15	-.03	.67	.68	.86	.04	.90	.03
DeepSeek V4-Pro	.00	.01	.40	.42	.51	.05	.61	.06
Kimi K3	-.01	.04	.55	.60	.71	.10	.64	.06
Gemini 3.6 Flash	-.09	-.07	.42	.44	.69	.09	.63	.05
Claude Opus 5	-.07	-.00	.52	.53	.78	.07	.75	.05

The four arms make four different predictions and all four hold. With no reference cases displayed at all, the target-only arm loads on neither reference: eleven of its twelve intervals span zero. With both references displayed but no City C label, every deployment loads on the two references almost equally, the pooling signature of a model that cannot tell which reference applies. Adding the correct label concentrates the loading on the named reference and removes it from the other, every “other” interval spanning zero. Inverting the label moves the loading wholesale onto the reference that the false cue names, leaving the truth-matched reference near zero.

The forecast therefore tracks the particular reference city the label names, including that city’s sampling error, whether or not the label is correct. No prior attached to the words *national* and *local* reproduces this. It also explains why implied-slope correlations plateau near .66 rather than near .996: faithful transfer of a noisy reference estimate is capped at that reference’s own .656 correlation with the truth. GPT-5.6 and Claude Opus 4.8 slightly exceed the cap at $k = 0$, which is what blending a reference estimate with a regime prior produces.

The estimator only detects an inefficiency. The signal this analysis reads is noise the forecast did not have to import. Each city draws its slope independently, so a reference city’s deviation from its regime mean is uncorrelated with City C’s: simulating the generator gives $\text{corr}(\beta^{\text{ref}} - \mu, g_C - \mu) = -.0000$ (Section C.7.4). Within a regime the reference’s fitted slope is therefore pure sampling error about a mean the regime already supplies, and the response-recovery-optimal policy is to read the cue, ignore the reference numbers, and predict the regime mean—which attains $\rho = .996$ against the .656 a faithful transfer is capped at. A forecaster that did this would show no loading on either reference and would be indistinguishable, under this estimator, from the lexical-prior account the section argues against. The estimator thus identifies reference selection only in forecasters that fail to shrink, and has least power on those closest to the shrinking policy: the two deployments above the transfer cap are

the two with the highest slope correlation, one of which also has the lowest $k = 0$ forecast error of the six. We read a positive loading as evidence for reference selection, and a null under this estimator as uninformative rather than as evidence for a prior.

Table 45: The same partial correlation in the arbitrary-symbol arm at $k = 0$, where KIV and ZOR carry no stated direction and their mapping reverses across episodes. The no-label columns repeat each model’s pooling control. Boldface marks symbol named-reference intervals excluding zero.

Model	n	ABC, no label		ABC + symbol	
		Named	Other	Named	Other
Qwen3-4B	250	.05	−.09	.00	−.03
Qwen2.5-7B	250	.11	.15	.28	.10
Qwen3-8B	250	.12	.06	.09	.11
Qwen2.5-14B	250	.31	.20	.32	.03
Qwen3-14B	250	.26	.20	.46	.06
Qwen3-32B	250	.20	.19	.40	.03
Qwen2.5-32B	250	.22	.15	.38	.02
Qwen2.5-72B	250	.25	.18	.53	.04
Llama 3.1 8B	249	−.02	.07	−.00	−.08
Llama 3.1 70B	250	.28	.19	.54	.00
DeepSeek V4-Pro	250	.40	.42	.51	.07
Kimi K3	250	.55	.60	.86	.01
Claude Opus 4.8	250	.39	.53	.71	−.05
Claude Opus 5	250	.52	.53	.83	.02
GPT-5.6 Sol	250	.67	.68	.98	.02

Applying the estimator to the arbitrary-symbol arm removes the lexical account entirely. Twelve of the fifteen released systems have a named-reference interval excluding zero, and the three that do not—Qwen3-4B, Qwen3-8B and Llama 3.1-8B—are exactly the three whose symbol-minus-no-context $\Delta\rho$ intervals span zero in Table 41. The two statistics are computed from different quantities and agree on all fifteen systems.

C.7 LINEAR-PROBE ANALYSIS OF CUE-CONDITIONED REPRESENTATIONS

C.7.1 BEHAVIORAL SCREENING CRITERIA

We screen the same held-out prompts with Qwen3-8B, Qwen3-14B, dense Qwen3-32B, and Qwen2.5-72B-Instruct. Each checkpoint received 16 test episodes crossed with $k \in \{0, 4\}$ and correct, absent, or inverted City C context (96 prompts), with 96/96 parseable forecasts.

Table 46: **Behavioral screening results for the linear-probe analysis.** Entries are paired mean differences with 95% episode-bootstrap intervals ($n = 16$). Evidence is no-context $k = 4$ minus $k = 0$ aligned slope; selection is correct minus absent context at $k = 0$; reversal is inverted minus correct context at $k = 0$; forecast penalty is inverted minus correct poll MAE at $k = 0$. Positive evidence and selection, negative reversal, and a positive forecast penalty are the predicted directions.

Model	Evidence	Selection	Reversal	Forecast penalty
Qwen3-8B	+ .119 [−.094, +.374]	+ .078 [−.012, +.219]	− .140 [−.315, +.007]	+ .812 [−.063, +2.023]
Qwen3-14B	+ .394 [+ .152, +.734]	+ .261 [+ .134, +.400]	− .380 [−.597, −.172]	+ 1.662 [+ .252, +3.155]
Qwen3-32B	+ .315 [+ .017, +.646]	+ .234 [+ .111, +.407]	− .390 [−.554, −.221]	+ 1.965 [+ 1.013, +3.083]
Qwen2.5-72B	+ .246 [+ .001, +.529]	+ .096 [−.077, +.259]	− .433 [−.674, −.220]	+ 2.206 [+ .957, +3.620]

The 14B and 32B intervals exclude zero in all four predicted directions. For 72B, the evidence, reversal, and forecast-penalty intervals exclude zero, while the correct-versus-absent selection interval spans zero. All four 8B intervals span zero. We use 32B for the dense Qwen3 scaling comparison and 72B for a cross-generation comparison; neither is interpreted as another point on a monotonic size curve. Qwen3-4B and Llama-3.1-8B use the identical task file, held-out split, and model, layer, target, and endpoint selection rules.

C.7.2 ACTIVATION AND PROBE PROTOCOL

We selected 64 episodes without replacement, balanced over which reference city was strongly responsive and whether City C’s latent regime was strong or weak. Forty-eight episodes form a development set and 16 form a sealed test set. Each episode is rendered at $k = 0$ and $k = 4$ under correct, absent, and inverted context, giving 384 activation records. Within an episode-prefix pair the three prompts have identical numeric content. The only intervention is the City C relationship description.

For the final input token, we save the embedding output and every transformer-block output. The resulting tensor dimensions are [384, 37, 2560] for Qwen3-4B; [384, 37, 4096] for Qwen3-8B; [384, 33, 4096] for Llama-3.1-8B; [384, 41, 5120] for Qwen3-14B; [384, 65, 5120] for Qwen3-32B; and [384, 81, 8192] for Qwen2.5-72B-Instruct. We fit separate dual-ridge probes at each layer and k , using only correct-context development episodes. Four cell-balanced development folds select the ridge strength from $\{10^{-4}, 3 \times 10^{-3}, 10^{-1}, 3, 100\}$ and then select one reported layer. The untouched test episodes are evaluated in all three context arms ($n = 16$ per arm). A 10,000-repeat held-out label permutation test keeps the development-selected layer, ridge strength, and fitted probe fixed. Thus neither the test labels nor test-arm scores select the reported layer.

We probe three increasingly strict targets: the binary strong/weak regime; the full realized slope g_C ; and the within-regime residual $g_C - .90$ for strong episodes or $g_C - .25$ for weak episodes. The first two are recoverable from the displayed evidence and the third is not, by construction: Section C.7.4 derives what a Bayes-optimal reader of the same prompt could score on each, and the residual is included as a specificity control rather than as a target a model could hit. Mean-pooled embedding and final-layer states are additional pooling controls.

Table 47: **Six-model correct-context probe summary.** Each entry gives $k = 0/k = 4$ on the 16 sealed episodes. Layers and ridge strengths were selected on development episodes only. Residual p is the one-sided 10,000-repeat held-out label-permutation value.

Model	Regime AUC	Full-slope R^2	Residual-slope R^2	Residual p
Qwen3-4B	1.000/1.000	.960/.885	-.389/-.227	.927/.614
Qwen3-8B	.844/1.000	.713/.653	-.359/-.374	.502/.735
Llama-3.1-8B	1.000/.984	.913/.731	-.126/-.207	.380/.572
Qwen3-14B	1.000/.984	.898/.627	-.531/-.128	.315/.450
Qwen3-32B	1.000/1.000	.967/.811	-.269/.003	.477/.221
Qwen2.5-72B	1.000/1.000	.975/.869	-.451/-.053	.675/.447

Table 48: **Held-out layer-wise open-weight probe.** AUC is reported for the regime target and R^2 for both slope targets. Correct, none, and wrong score the same development-selected probe against the true City C target under the three context arms; wrong/cue instead scores the inverted arm against the relationship named by its cue. p is the one-sided held-out label-permutation value on the correct-context test episodes.

Model	k	Target	Layer	Dev CV	Correct	None	Wrong/true	Wrong/cue	p
Qwen3-4B	0	Regime (AUC)	1	1.000	1.000	.562	.000	1.000	< .001
		Full slope (R^2)	24	.936	.960	-.517	-2.614	.956	< .001
		Residual slope (R^2)	15	.054	-.389	-.296	-.258	–	.927
	4	Regime (AUC)	19	1.000	1.000	.609	.094	.906	< .001
		Full slope (R^2)	20	.805	.885	-.065	-1.147	.492	< .001
		Residual slope (R^2)	14	.004	-.227	.089	-.183	–	.614
	Qwen3-8B	0	Regime (AUC)	19	1.000	.844	.453	.078	.009
			Full slope (R^2)	23	.870	.713	-.697	-2.344	.771 < .001
			Residual slope (R^2)	1	.063	-.359	-.673	-.378	– .502
		4	Regime (AUC)	23	.925	1.000	.578	.000	1.000 < .001
			Full slope (R^2)	23	.462	.653	-.512	-1.550	.693 < .001
			Residual slope (R^2)	33	.030	-.374	-1.003	-.029	– .735
	Llama-3.1-8B	0	Regime (AUC)	1	1.000	1.000	.359	.000	1.000 < .001
			Full slope (R^2)	1	.907	.913	-18.953	-2.629	.972 < .001
			Residual slope (R^2)	12	.018	-.126	-.458	-.227	– .380
		4	Regime (AUC)	1	1.000	.984	.438	.000	1.000 < .001
			Full slope (R^2)	1	.751	.731	-15.450	-2.082	.808 < .001
			Residual slope (R^2)	10	.006	-.207	-.148	-.150	– .572
Qwen3-14B	0	Regime (AUC)	2	1.000	1.000	.422	.000	1.000	< .001
		Full slope (R^2)	26	.874	.898	-.522	-2.423	.943	< .001
		Residual slope (R^2)	34	.159	-.531	-9.255	-.096	–	.315
	4	Regime (AUC)	26	.880	.984	.688	.109	.891	< .001
		Full slope (R^2)	26	.408	.627	-.411	-1.191	.409	< .001
		Residual slope (R^2)	1	-.013	-.128	-.605	-.159	–	.450
Qwen3-32B	0	Regime (AUC)	27	1.000	1.000	.359	.000	1.000	< .001
		Full slope (R^2)	47	.930	.967	-.027	-2.570	.958	< .001
		Residual slope (R^2)	1	.062	-.269	-.232	-.259	–	.477
	4	Regime (AUC)	44	1.000	1.000	.562	.000	1.000	< .001
		Full slope (R^2)	47	.797	.811	-.160	-1.694	.666	< .001
		Residual slope (R^2)	1	.043	.003	.062	-.000	–	.221
Qwen2.5-72B	0	Regime (AUC)	53	1.000	1.000	.469	.000	1.000	< .001
		Full slope (R^2)	58	.955	.975	-1.050	-2.958	.969	< .001
		Residual slope (R^2)	32	.061	-.451	-1.192	-.401	–	.675
	4	Regime (AUC)	58	.995	1.000	.922	.062	.938	< .001
		Full slope (R^2)	58	.678	.869	-.468	-1.072	.493	< .001
		Residual slope (R^2)	1	-.017	-.053	-.063	-.053	–	.447

C.7.3 REGIME AND FULL-SLOPE DECODABILITY

Across the six checkpoints and two evidence depths, correct-context regime AUC is .844–1.000 and full-slope R^2 is .627–.975 (Table 47). This information is not a context-invariant encoding of the true coefficient: removing the City C description drives selected-layer full-slope R^2 below zero in all 12 model- k cells. Under inversion, true-target regime AUC falls to .000–.109 while cue-target AUC is .891–1.000; the full-slope probe likewise fits the cue-assigned relationship better than the true one. The Llama replication shows that this pattern is not specific to the Qwen checkpoints. Across both




Figure 14: **Semantic-arm slope decoding follows cue assignment.** Each line is a sealed model- k cell across the six open checkpoints of the 16-episode roster; purple is the median. Probes are fit on correct-context development episodes. Under inversion, held-out R^2 collapses against the true slope and recovers against the cue-assigned slope. The input-embedding control below shows this signal is available before any computation, which is why the arbitrary-label replication rather than this figure carries the representational claim.

model families, the linearly available signal therefore tracks the response family that context assigns to City C.

C.7.4 WHAT EACH PROBE TARGET CAN BE SCORED AGAINST

A probe is scored against a latent generator quantity, so a null is only a fact about a model if the displayed evidence determines that quantity. We therefore derive, for each target and evidence depth, the score a Bayes-optimal reader of the same prompt would obtain. This is an analytic property of Equation 1 and its row construction, computed in closed form and re-derived by simulating the generator’s own row logic; the two agree to within .005 and no model output enters; the derivation is `exp2_v2/biased_news/analysis/compute_bayes_ceilings.py`. Table 49 gives the ceilings. The same module reports forecast-MAE floors: 2.28 poll points at $k = 0$ with no cue, 0.17 for an oracle told the regime means, and 1.78 for a reader who must instead estimate the named regime’s slope from its four displayed rows.

The same simulation records one quantity the reference-selection analysis of Section C.6 depends on. Cities draw their slopes independently, so a reference city’s deviation from its regime mean carries no information about City C’s: $\text{corr}(\hat{\beta}^{\text{ref}} - \mu, g_C - \mu) = -.0000$ over 400,000 simulated episodes, against a reference-fit standard deviation of .318 within regime. That is why a loading on the reference’s own sampling error identifies reference selection, and equally why a forecaster that shrank optimally would show none.

Table 49: Maximum score a Bayes-optimal reader of the same prompt can obtain on each probe target, under the frozen data-generating process. The regime and the full slope are recoverable; the within-regime residual is not, because it is generated with standard deviation .03 while a through-origin fit to City C’s displayed rows carries a standard deviation of .318. At $k = 0$ City C displays no rows at all and the residual ceiling is exactly zero.

Probe target		$k = 0$	$k = 4$
Regime, correct semantic cue	AUC	1.000	1.000
Regime, arbitrary label	AUC	0.979	—
Full slope g_C	R^2	0.992	0.992
Residual slope $g_C - \mu$	R^2	0.000	0.009




Figure 15: **Development layer trajectories for the six-model probe.** Curves show four-fold development CV rather than test performance; stars mark the development-selected layer. Layer depth is normalized because the checkpoints have 32–80 transformer blocks. The regime and full-slope signals are broad, whereas residual-slope CV is weak and does not survive held-out evaluation.

C.7.5 SPECIFICITY: WITHIN-REGIME RESIDUALS

The residual target is the specificity control, and it behaves as a target with no available signal should. The 12 correct-context residual-slope R^2 values range from $-.531$ to $.003$, none of their permutation tests rejects the null ($p = .221$ – $.927$), and 33 of 36 arm-level residual scores are nonpositive. Against a ceiling of $.000$ at $k = 0$ and $.009$ at $k = 4$, that is the result the design permits: no reader of these prompts, and no representation computed from them, could score materially above zero. We therefore read the residual null as evidence that the probe does not manufacture structure where none is available, and not as a limit on what these checkpoints encode. A probe that returned a positive residual score here would be the finding, and it would be a warning about the probe.

The two recoverable targets do carry information, and the input-embedding control is what bounds their interpretation. Mean-pooled input embeddings for every Qwen checkpoint already give perfect regime AUC and full-slope $R^2 = .988-.993$ under correct context—at the .992 ceiling—because the explicit description names a regime and the regime means account for essentially all slope variance. Llama input embeddings likewise give perfect regime AUC and full-slope $R^2 = .756/.827$. The layer-wise semantic result is therefore evidence that the cue-selected response family is linearly represented, and reaches nothing further about the continuous g_C : the part of g_C beyond the family is not in the prompt to begin with. The original behavioral inversion screen connects this representation to output sensitivity in the three larger screened checkpoints: inverted context raises $k = 0$ forecast MAE by 1.66, 1.96, and 2.21 points for 14B, 32B, and 72B, respectively.

C.7.6 POWER-MATCHED AND CROSS-MODEL ARBITRARY-SYMBOL REPLICATION

The six-checkpoint roster above seals 16 test episodes per arm. Its exact label-permutation null spans ROC AUC $[.203, .797]$, so it can separate a score from chance only above about .75. That resolution suffices for the semantic arm, whose scores sit at ceiling, but it cannot evaluate a weaker signal at all. We therefore evaluate Qwen3-14B at the design’s balanced maximum—248 episodes, 160 development and 88 sealed test, 1,488 activation records—once with the semantic cue and once with the episode-randomized KIV/ZOR labels. Anchor, targets, fold structure, ridge grid, layer-selection rule, pooling controls and the 10,000-repeat held-out permutation test are unchanged; only the number of episodes differs. The 88-episode null reaches AUC .603. The enlarged development set is disjoint from every episode used by the 16-episode analysis, providing an independent estimate rather than a re-slice.

Table 50: **Power-matched Qwen3-14B probe, 88 sealed episodes per arm.** Correct, none and wrong score the development-selected probe against the true City C target under the three cue arms; wrong/cue scores the inverted arm against the relationship its cue names. “Layer” is the development-selected block of 40 and “Embed.” is the mean-pooled input-embedding control under correct context. p is the one-sided held-out label-permutation value.

Run	Target	Correct	None	Wrong	Wrong/cue	Layer	Embed.	p	Behav. $\Delta\rho$
Qwen3-14B, semantic	$k=0$ regime AUC	1.000	.429	.000	1.000	1	1.000	<.001	—
	$k=4$ regime AUC	1.000	.560	.000	1.000	1	1.000	<.001	
	$k=0$ slope R^2	.942	-.300	-2.806	.961	27	.989	<.001	
	$k=4$ slope R^2	.910	-24.039	-2.387	.906	1	.989	<.001	
Qwen3-14B, symbol	$k=0$ regime AUC	.678	.467	.321	.679	26	.555	.0016	.398
	$k=4$ regime AUC	.713	.689	.747	.253	40	.582	<.001	
	$k=0$ slope R^2	.016	-.094	-.130	.065	26	.003	.1162	
	$k=4$ slope R^2	.030	-.237	.126	-.855	40	-.032	<.001	
Qwen3-8B, symbol	$k=0$ regime AUC	.566	.535	.567	.433	10	.534	.1471	.023
	$k=4$ regime AUC	.610	.545	.593	.407	19	.621	.0401	
	$k=0$ slope R^2	-.001	.005	.004	-.007	25	.003	.4883	
	$k=4$ slope R^2	-.151	-.219	-.188	-.504	20	-.018	.0314	
Qwen3-4B, symbol	$k=0$ regime AUC	.520	.419	.329	.671	23	.498	.3756	.095
	$k=4$ regime AUC	.685	.689	.637	.363	22	.585	.0017	
	$k=0$ slope R^2	-.009	-.065	-.079	.023	23	.004	.2899	
	$k=4$ slope R^2	-.021	-.265	-.031	-.687	23	.004	.0014	
Qwen2.5-32B, symbol	$k=0$ regime AUC	.778	.571	.200	.800	50	.550	<.001	.372
	$k=4$ regime AUC	.797	.831	.662	.338	46	.589	<.001	
	$k=0$ slope R^2	.217	-.256	-1.033	.254	50	.004	<.001	
	$k=4$ slope R^2	.263	-.563	-.044	-.774	46	-.029	<.001	
Qwen2.5-72B, symbol	$k=0$ regime AUC	.662	.492	.210	.790	63	.556	.0026	.381
	$k=4$ regime AUC	.793	.745	.651	.349	61	.575	<.001	
	$k=0$ slope R^2	-.005	-.490	-1.034	.187	62	.004	.0072	
	$k=4$ slope R^2	.222	-.624	-.129	-.870	61	-.047	<.001	
Llama 3.1-8B, symbol	$k=0$ regime AUC	.423	.447	.452	.548	11	.539	.8925	.045
	$k=4$ regime AUC	.543	.604	.549	.451	31	.604	.2470	
	$k=0$ slope R^2	-.097	-.089	-.086	-.001	11	.003	.8607	
	$k=4$ slope R^2	-.275	-.091	-.156	-.358	31	.009	.4529	
Llama 3.1-70B, symbol	$k=0$ regime AUC	.728	.468	.292	.708	75	.495	<.001	.424
	$k=4$ regime AUC	.810	.740	.721	.279	42	.606	<.001	
	$k=0$ slope R^2	.076	-8.200	-.961	.082	75	.003	<.001	
	$k=4$ slope R^2	.222	-3.224	.087	-1.007	55	.007	<.001	

The two arms separate sharply, and the input-embedding control is what separates them. Under the semantic cue the development-selected layer is the first transformer block for three of the four conditions, and mean-pooled input embeddings alone already reach regime AUC 1.000 and full-slope $R^2 = .989$: the decodable quantity is the cue token itself. Under arbitrary labels the same control collapses to AUC .555 and $R^2 = .003$, the selected layers move to blocks 26 and 40 of 40, and held-out regime AUC is .678 at $k = 0$ ($p = .0016$) and .713 at $k = 4$ ($p < .001$). The induced label-to-regime mapping is therefore linearly represented, but only after computation, far more weakly than a regime the model was told outright, and below the .979 AUC a Bayes-optimal reader of the same two reference cities could reach. Unlike the residual target, the full slope is recoverable here—its ceiling is .992—so its symbol-arm score of $R^2 = .016$ at $k = 0$, which does not reject its null, is a genuine gap rather than a property of the design.

Representation and behavior agree closely once the cue must be induced. On the same 88 sealed episodes the model’s own forecasts give regime AUC .702, .518 and .331 under correct, absent and inverted labels, against probe AUCs of .678, .467 and .321; inversion raises $k = 0$ poll MAE from 2.49 to 4.28. With four City C cases the inverted-arm representation reverses: true-target regime AUC rises from .321 to .747, and the full-slope probe fits the cue-named relationship at $R^2 = -.855$ while fitting the true one at $+.126$. Direct target evidence overrides the induced label in the representation as it does in the forecast.

Six further checkpoints received the identical symbol protocol at the same 88 sealed episodes: Qwen2.5-32B and Qwen2.5-72B, which change model generation; Qwen3-8B and Qwen3-4B, which

hold family, tokenizer and architecture fixed and change only scale; and Llama 3.1-8B and Llama 3.1-70B, which do the same within a second vendor and architecture. Held-out regime AUC at $k = 0$ is .778 for Qwen2.5-32B ($p < .001$, selected layer 50 of 64), .728 for Llama 3.1-70B ($p < .001$, layer 75 of 81), .662 for Qwen2.5-72B ($p = .003$, layer 63 of 80), .566 for Qwen3-8B ($p = .15$), .520 for Qwen3-4B ($p = .38$) and .423 for Llama 3.1-8B ($p = .89$). No null checkpoint rejects in any $k = 0$ condition—regime, slope or residual slope. Mean-pooled input embeddings stay between .495 and .556 in every symbol run, so no symbol arm carries the mapping before computation.

Decodability tracks behavioral recovery across all seven checkpoints, four positive and three null. The four whose forecasts recover the mapping are the four in which it decodes (Llama 3.1-70B, $\Delta\rho = .424$; Qwen3-14B, .398; Qwen2.5-72B, .381; Qwen2.5-32B, .372), and the three whose forecasts do not are the three in which it does not (Qwen3-4B, .095; Llama 3.1-8B, .045; Qwen3-8B, .023; all three intervals spanning zero). The Llama 3.1 contrast carries the weight, because it holds family, tokenizer and architecture fixed, varies only scale, and is the one within-family contrast whose behavioral classification survives the harder generator population of Section C.8: 8B is null on both populations and 70B positive on both, while the probe runs .423 to .728 across them. Because the null and the positive share an architecture and a tokenizer, the split is not an artifact of either.

The Qwen3 sequence—.520, .566, .678 from 4B to 14B—is consistent with this but does not add independent weight, because the population that would test it is the one on which its behavioral boundary moves. That movement is smaller than the $\Delta\rho$ column alone suggests. Symbol-arm regime accuracy is .512, .516, .660, .656 originally against .528, .554, .589, .606 on the harder population. The larger pair remains above the smaller pair, but 14B and 32B exchange order and the between-group separation shrinks. The $\Delta\rho$ contrast then subtracts a no-context baseline whose standard error is about .064, and baselines of $-.036$ for Qwen3-8B and $+.025$ for Qwen3-14B are enough to put one interval above zero and the other .004 below it, with the two intervals overlapping by 78%. We therefore read the harder population as preserving the broad size-group contrast without locating the significance threshold at a particular checkpoint. Section C.8 should be read alongside the seven-for-seven agreement, which is evidence that the probe reads what the forecasts express on the episodes both were measured on and not evidence for a scale law. The positives span three model generations and two vendors, and every positive shows the same shape: the decoded regime follows the cue at $k = 0$ and the truth at $k = 4$. Qwen3-4B and Qwen3-8B do reject at $k = 4$ (.685, $p = .002$ and .610, $p = .04$), which is expected rather than contrary: at $k = 4$ the displayed target cases exhibit the regime whatever the label says, and the mean-pooled input embeddings rise to .585 and .621 accordingly, in Qwen3-8B’s case above the probe itself. The induction claim is a $k = 0$ claim throughout. The sealed test is 88 episodes per model, well above the 16 the roster began with but not a large sample, and the permutation nulls are computed rather than assumed.

At 16 episodes, the same arbitrary-symbol procedure returns $k = 0$ regime AUC .312 with $p = .90$. Its [.203, .797] null is too broad to resolve a signal of the magnitude observed at 88 episodes.

C.7.7 CAUSAL TEST: PATCHING THE CITY C LABEL REPRESENTATION

Decodability is not mediation. To test whether the representation the probe reads is the one the forecast uses, we intervene on it directly. Under arbitrary labels, an episode’s correct and swapped prompts are positionally aligned through the City C label and differ at exactly two label-token positions. We replace the residual stream at those two positions with the states the opposite arm produces, in both directions, and regenerate the forecast greedily. The endpoint is the episode-level mean of the regime-aligned correct-into-swapped shift and the sign-reversed swapped-into-correct shift, tested one-sided against an episode sign-flip null.

The analysis uses three patch sites and two controls. The primary site is a three-layer window centred on the development-selected relational layer; patching at the embedding output only, and at every causally effective layer, are the comparison sites; the unpatched arm and a self-patch that writes each state back over itself are the controls. Layer, window and ridge were selected on development episodes with the sealed test unread.

Table 51: **Causal effect of patching the two City C symbol subtokens, Qwen3-14B.** Shift is the regime-aligned symmetric patch mean; positive means the forecast moved toward the relationship the donor label names. The 16- and 88-episode analyses select their windows independently on their own development sets. The null interval is the 88-episode sign-flip null.

Patch site	k	16 sealed, window 33–35			88 sealed, window 19–21			Null 95%
		Shift	p	n	Shift	p	n	
Selected window	$k=0$	-.033	.75	14	.333	<.0001	82	$[-.177, .178]$
	$k=4$	-.007	.75	15	-.035	.8256	87	$[-.073, .073]$
Embedding only	$k=0$	.236	.15	14	.319	.0002	84	$[-.182, .182]$
	$k=4$	.115	.23	15	-.017	.6419	87	$[-.087, .087]$
All layers	$k=0$	.236	.15	14	.319	.0002	84	$[-.182, .182]$
	$k=4$	.115	.23	15	-.017	.6419	87	$[-.087, .087]$

At $k = 0$ the intervention moves the forecast: patching the selected window shifts the regime-aligned prediction by .333 against a null interval of $[-.177, .178]$ ($p < .0001$, 82 complete episodes), and patching the embedding alone gives .319. All eight stability checks pass, including 347/347 exact self-patch matches.

Two readings of the table matter more than the significance. First, the embedding-only and all-layer rows are identical, and necessarily so: the two prefixes differ only at the patched label positions, so overwriting their embeddings and letting the model recompute reproduces what overwriting every later state does. They are one intervention reported at two depths, not two converging results. What distinguishes the mid-window row is that it reaches the same magnitude while touching only three layers late in the network.

Second, the effect is confined to $k = 0$. With four City C cases every patch site returns a null ($p = .83$ and $p = .64$), so the induced label no longer controls the forecast once direct target evidence is available. That is the revision claim of Section C.5 in causal form, and it matches the representational result: at $k = 4$ the inverted-arm probe recovers the true regime rather than the cue-assigned one.

At 16 episodes, the primary patch estimate is $-.033$ at $p = .75$ and the sign-flip null spans $[-.387, .387]$. Its 64-episode development set selects layers 33–35, whereas the 160-episode development set selects 19–21 at a higher cross-validation score (.903 against .806).

Permutation implementation. The test enumerates all 2^n sign assignments for 16 episodes; exhaustive enumeration is infeasible for 88 ($2^{88} \approx 3 \times 10^{26}$). Enumeration is retained for $n \leq 20$; above it the same null is sampled with 200,000 draws under a fixed seed, giving a Monte Carlo standard error of .0005 at $p = .05$. The estimand, endpoint, and direction are identical under both implementations.

Scope. This is one checkpoint under arbitrary labels. It shows that the label representation Qwen3-14B computes is causally responsible for its zero-shot regime assignment, and that direct evidence displaces it. It does not establish the same for the hosted deployments, the semantic cue, or other architectures.

C.7.8 REPRODUCIBILITY AND SCOPE

Design scope. The cue is binary and the regimes well separated, so these results establish selection between two displayed references rather than graded sensitivity to a continuous cue. The per-episode evidence is too small to test exact coefficient recovery: the claim is which reference the forecast follows, not that the model reproduces its slope. Causal patching covers one of seven open checkpoints, and the hosted deployments are tested only behaviorally, since their internals are unavailable.

The checkpoints are pinned to commits `cdbee75f17c01a7cc42f958dc650907174af0554` (Qwen3-4B), `b968826d9c46dd6066d109eabc6255188de91218` (Qwen3-8B), `0e9e39f249a16976918f6564b8830bc894c89659` (Llama-3.1-8B), `40c069824f4251a91eefaf281ebe4c544efd3e18` (Qwen3-14B), `9216db5781bf21249d130ec9da846c4624c16137` (Qwen3-32B), and

495f39366efef23836d0cfae4fbe635880d2be31 (Qwen2.5-72B). The activation hashes are e191385963d639aa5b8574a9abb91b3519aed934ac491f863655631b26dbbc (Qwen3-4B) and 7f65aeec7a91e404cc60068618fedf25c679abde5d48c2b8477f736670fbc54f (Llama-3.1-8B); all six hashes and tensor dimensions are enforced by the fail-closed generator. Frozen tasks, activation tensors, extraction manifests, generated behavior, selected predictions, complete layer curves, and probe results are under exp2_v2/biased_news/data/coin_city_stable_relationship_claude_n250_v4/mechanistic_probe/. The analysis and asset generator are exp2_v2/biased_news/mechanistic_probe/analyze_probe.py and paper/generate_coin_city_mechanistic_comparison.py. The causal run is mechanistic_probe/qwen3_14b_symbol_relational_v2/, frozen by freeze_symbol_relational_protocol.py and executed by extract_symbol_label_states.py, patch_symbol_label_activations.py and analyze_symbol_relational_probe.py; every count in that pipeline is derived from the frozen task manifest and checked by tests/test_symbol_relational_scaling.py. paper/generate_coin_city_causal_patch_table.py fails closed on the run’s study name, sealed-test size and stability gate. The power-matched runs are mechanistic_probe/qwen3_14b_probe_v2/, mechanistic_probe/qwen3_14b_symbol_probe_v2/, mechanistic_probe/qwen3_8b_symbol_probe_v2/, mechanistic_probe/qwen3_4b_symbol_probe_v2/, mechanistic_probe/llama3_1_70b_symbol_probe_v2/, mechanistic_probe/qwen2_5_72b_symbol_probe_v2/, mechanistic_probe/qwen2_5_32b_symbol_probe_v2/ (checkpoint 5ed1c97bbab6ce5cda5812749b4c0bdf79b18dd) and mechanistic_probe/llama3_1_8b_symbol_probe_v2/ (checkpoint 0e9e39f249a16976918f6564b8830bc894c89659), built by exp2_v2/biased_news/mechanistic_probe/build_probe_tasks.py and build_arbitrary_symbol_probe_tasks.py with 62 and 31 episodes per factorial cell; paper/generate_coin_city_probe_power_table.py regenerates their table and validates the sealed-test size and task hash. Each run’s activation tensor is regenerable from its frozen task file and pinned checkpoint; the extraction manifests record the tensor hashes rather than the tensors. Cross-vendor probe evidence here is the Llama 3.1-8B null rather than a positive. The diagnostic covers six checkpoints spanning Qwen and Llama at 16 sealed episodes per arm, plus the matched 88-episode semantic and arbitrary-symbol Qwen3-14B pair. For the semantic cue it is correlational and does not establish that the decoded direction mediates the forecast; under arbitrary labels Section C.7.7 provides a causal intervention for one checkpoint. Both complement the paired behavioral interventions above.

C.8 ROBUSTNESS ACROSS GENERATOR POPULATIONS AND DEPLOYMENTS

C.8.1 GENERATOR-POPULATION ROBUSTNESS

Fresh seeds 940000–940249 target strong/weak means of .75/.40 with slope SD .10, jointly narrowing the regime gap and widening within-regime dispersion. The realized target slopes are $.752 \pm .100$ and $.404 \pm .105$. Draws were neither truncated nor reordered: two reference pairs reverse their nominal latent ordering and remain in the analysis. Case noise, factorial balance, prompts, prefix set, local checkpoint settings, and episode-randomized symbol mapping are unchanged.

All ten open checkpoints were run on the three matched arms, producing 37,500 response records; 37,355 contain a numeric forecast, and unresolved formats remain missing. Table 52 compares the $k = 0$ symbol-minus-no-context contrast across generator populations. Bold intervals exclude zero; negative ΔMAE would favor symbol context. Because this population changes separation and dispersion together, it tests the joint perturbation rather than identifying either parameter’s effect separately.

Table 52: Arbitrary-symbol recovery on the original and harder generator populations. Intervals are 95% episode-bootstrap intervals with 5,000 draws; n is the harder-population complete-case count.

Model	n	Original $\Delta\rho$ [95% CI]	Harder $\Delta\rho$ [95% CI]	Symbol acc.	Harder Δ MAE [95% CI]
Qwen3-4B	250	+0.095 [−.060, .230]	+0.050 [−.079, .197]	.528	+0.052 [−.056, .161]
Qwen3-8B	249	+0.023 [−.128, .171]	+0.231 [.074, .370]	.554	−.113 [−.269, .036]
Qwen3-14B	246	+0.398 [.253, .541]	+0.155 [−.004, .303]	.589	+0.361 [.076, .644]
Qwen3-32B	249	+0.302 [.177, .433]	+0.193 [.063, .325]	.606	+0.848 [.539, 1.168]
Qwen2.5-7B	250	+0.264 [.112, .406]	+0.149 [.009, .279]	.540	+0.303 [.000, .627]
Qwen2.5-14B	250	+0.236 [.107, .356]	+0.133 [.005, .254]	.584	+0.365 [.115, .609]
Qwen2.5-32B	250	+0.372 [.226, .502]	+0.247 [.109, .383]	.616	+0.882 [.590, 1.186]
Qwen2.5-72B	250	+0.381 [.250, .508]	+0.336 [.204, .463]	.604	+1.012 [.717, 1.316]
Llama 3.1-8B	229	+0.045 [−.123, .250]	−0.077 [−.216, .079]	.524	+0.597 [.229, .952]
Llama 3.1-70B	249	+0.424 [.314, .543]	+0.242 [.119, .369]	.558	+0.856 [.465, 1.242]

Which Qwen3 checkpoints clear significance changes, while the larger pair retains higher symbol-arm regime accuracy than the smaller pair. Qwen3-8B moves from an original null to $\Delta\rho = .231$ [.074, .370] while Qwen3-14B falls to .155 [−.004, .303]; Qwen3-4B remains null and Qwen3-32B remains positive. Regime accuracy in the symbol arm, which subtracts no baseline, is .512, .516, .660, .656 originally against .528, .554, .589, .606 here. The 14B and 32B checkpoints exchange order, and the smaller models improve slightly, but the gap between 8B and 14B shrinks from .144 to .035. The $\Delta\rho$ contrast then subtracts a no-context baseline carrying a standard error of about .064, and baselines of −.036 for Qwen3-8B against +.025 for Qwen3-14B suffice to place one interval above zero and the other .004 below it, with the two overlapping by 78%. We read this as a detection threshold that is no longer locatable to a checkpoint on the compressed population; it does not support a fixed scale threshold.

All four Qwen2.5 intervals remain positive, from .149 [.009, .279] at 7B to .336 [.204, .463] at 72B, and both Llama checkpoints keep their classification, 8B null in each population and 70B positive in each. Seven of ten intervals therefore exclude zero in each population and eight of ten checkpoints keep their classification, the Qwen3 8B/14B pair being the two that do not. Mapping recovery survives the substantially less separable population; the checkpoint at which it becomes detectable does not, and neither does a scale reading across families, since Qwen2.5-14B sits below Qwen2.5-7B in both. Part of the shrinkage is mechanical: the regime explains less of the slope here, capping ρ at .87 against .996. The observed $\Delta\rho$ falls further than that alone predicts. Numerical use stays distinct: Qwen2.5-72B’s discrimination improves while symbol context raises forecast MAE by 1.012 [.717, 1.316].

C.8.2 ONE INDEPENDENT HOSTED REPEAT

GPT-5.6 was replayed once on the 250 $k = 0$ prompts in each no-context, semantic-context, and arbitrary-symbol arm. The deployment, low reasoning effort, provider-native sampling, prompt hashes, parser, and one-attempt policy were unchanged; all 750 repeat responses parsed.

Table 53: Independent GPT-5.6 $k = 0$ repeat. Paired intervals resample episodes; mean $|\Delta\hat{y}|$ is the per-task absolute change in forecast.

Arm	Forecast MAE			$\rho(\hat{g}, g_C)$			Mean $ \Delta\hat{y} $
	Original	Repeat	Repeat − original [95% CI]	Original	Repeat	Repeat − original [95% CI]	
No C context	2.518	2.451	−.068 [−.130, −.012]	−.041	−.013	+0.029 [−.008, .066]	.166
Semantic context	1.596	1.613	+0.016 [−.051, .082]	.678	.678	+0.000 [−.032, .037]	.193
Arbitrary symbol	1.744	1.708	−.036 [−.104, .023]	.655	.648	−.006 [−.016, .002]	.167

Per-task mean absolute forecast changes are .166, .193, and .167 points across the three arms. The arbitrary-symbol-minus-no-context correlation effect is .696 originally and .661 on repeat; their difference is −.035 [−.075, .003]. The corresponding forecast-MAE benefit is −.774 originally and −.743 on repeat, with difference +.031 [−.046, .104]. This one repeat supports stability of the paired conclusions for this named snapshot, not a general estimate of provider or model variance.

C.9 MISSINGNESS, SCOPE, AND ARTIFACT MAP

C.9.1 MISSINGNESS AND RECORD RECOVERY

Claude Opus 5 has two unparseable primary task cells: one correct-context response at $k = 1$ and one no-context response at $k = 4$. Both exhausted the output budget during extended thinking and returned no answer. GPT-5.6 has one no-context failure at $k = 1$, where the provider returned a model-side HTTP 500 before any response. None was rerun. Complete-case pairing therefore gives $n = 249$ for the corresponding deployment-prefix rows and $n = 250$ elsewhere. These three failures constitute 0.013% of the 22,500 primary task cells.

Claude Opus 4.8, DeepSeek-V4-Pro, Kimi K3, and Gemini 3.6 Flash have 3,750 parseable primary responses each after the one offline Claude key-whitespace normalization described above. That normalization changed neither the forecast value nor the raw model answer and made no additional API call.

All six misleading-context analytic files contain exactly 1,250 unique parseable tasks. Concurrent writes in the Kimi collection produced 11 duplicate records and 14 malformed intermediate lines. Construction of the analytic file retained the earliest response for each task and identified 11 absent task IDs; those 11 tasks were collected under the unchanged Kimi configuration. Thus 7,500 is the number of final unique analytic records for the misleading arm rather than the number of intermediate records or provider-level attempts.

The expanded arbitrary-symbol evaluation contains fifteen response-record triplets, or 56,250 records. Within every model, each arm contains 1,250 unique task IDs matched to the frozen task-file prompt hash. Original outputs are preserved. When a model placed an explicit arithmetic expression rather than a number in `predicted_poll`, a no-call reparse accepts only numeric constants, parentheses, unary signs, and the four arithmetic operators, and requires the result to lie in the primary parser’s 0–100 range. For Llama 3.1-8B, this reparse leaves 1,234/1,241/1,232 parseable no-context, semantic-context, and symbol records, respectively. Its matched complete-case n is 242 at $k = 0$ and ranges from 238 to 246 across the five prefixes. Qwen3-32B retains one unresolved symbol response, leaving 1,249 parseable symbol records; every other model parses all 1,250 in every arm, and every model except Llama 3.1-8B has $n = 250$ at $k = 0$. The final table reports these complete-case denominators, and unresolved outputs remain missing rather than being re-asked.

C.9.2 INTERPRETIVE SCOPE

The aggregate behavior is consistent with context-conditioned response-regime assignment and integration of direct City C evidence in this controlled synthetic setting. A label-conditioned analogical-regression benchmark—map the label to a reference, estimate that reference’s slope, transfer it to City C, and revise with City C’s own cases—is sufficient to reproduce the pattern but does not identify the model’s internal algorithm. In the misleading arm, reversing only the City C label reverses forecast assignment while all numeric content remains fixed.

The primary binary semantic cue is aligned by construction with the higher- or lower-response family and has a fixed meaning across episodes. The arbitrary-symbol control removes that meaning by reversing the KIV/ZOR mapping across episodes; it therefore tests whether the semantic-arm result can arise solely from a pre-existing national-greater-than-local prior. The expanded 15-model comparison shows detectable symbol-based discrimination for Qwen3-14B and Qwen3-32B, all four tested Qwen2.5 checkpoints, Llama 3.1-70B, and all five hosted systems. In this original population, intervals span zero for Qwen3-4B, Qwen3-8B, and Llama 3.1-8B. Within Qwen3 the mapping is recovered only from 14B upward, and within Llama only at 70B, providing within-family contrasts on the original population; Qwen2.5-7B’s positive result alongside Qwen3-8B’s interval spanning zero, and Qwen3-32B’s smaller point estimate than Qwen3-14B’s, show that the transition is neither monotone in nor determined by parameter count alone. Four quantitative analyses bear on whether a model induces the arbitrary label-to-regime mapping, and they are computed from different quantities: the symbol-minus-no-context correlation contrast on forecasts; the partial correlation between the implied slope and the cue-named reference’s own fitted slope (Section C.6); held-out linear decoding of hidden states; and causal patching of those states (Section C.7.7). Where more than one was run on a checkpoint they agree, including Llama 3.1-8B, where every analysis run is null. Agreement among measures that could have disagreed is the reason to prefer the induction account over a lexical or

pooling one. It does not extend the claim to the hosted deployments, whose internals are unavailable, and only one checkpoint carries the causal test. A fifth, descriptive audit asks what models themselves say in their one-sentence rationales.

Deterministic lexical rules were developed on a hash-selected 50-episode subset and then applied without revision to the remaining 200 episodes. Across all three comparison arms, 56,057 of 56,250 response records contain a rationale. At $k = 0$ in the arbitrary-symbol arm, 2,992 of 3,000 held-out records contain one; the coder identifies a selected City A/B reference in 1,832, and every codable statement names the reference that actually shares City C’s episode-local symbol. In the inverted semantic arm, all 1,200 held-out records contain a rationale: all 1,089 codable reference statements name the reference selected by the false cue, and all 1,186 codable national/local statements name the false regime. These outputs are behavioral self-reports, not privileged evidence of computation; in particular, the coder abstains whenever a statement does not make an unambiguous reference or regime claim. The conditional pattern also appears in the three checkpoints with null behavioral symbol contrasts, so it establishes declared cue interpretation rather than successful numerical use.

Because discrimination and forecast MAE can move in opposite directions, mapping induction is not equivalent to calibrated numerical use. Ambiguous, continuous, and adversarial cues remain untested. All prompts also tell models that within-city response is stable and explain the matched-pair structure.

The query target is a latent conditional mean with no newly sampled observation noise. This isolates response recovery but is not equivalent to calibration on a future noisy poll. The six primary deployments differ in provider, reasoning mode, output budget, and service tier. The symbol module mixes local vLLM checkpoints with hosted calls and includes the documented cross-resource Opus 4.8 arm. Cross-model comparisons are therefore descriptive rather than configuration-matched capability rankings. The cue-inversion module is a matched stress test over the same task identities and numeric inputs.

The primary and harder designs are two reproducible seeded populations in the same synthetic world. The second jointly narrows the regime gap and widens slope dispersion, but two settings do not characterize generator-parameter uncertainty or isolate those changes. Bootstrap intervals resample episodes within a population and are pointwise; they do not incorporate provider drift or correct for the many model, arm, and prefix comparisons. The evaluated populations also do not characterize factorial variation in generator parameters, ambiguous cues, or calibration against noisy held-out outcomes.

C.9.3 REPRODUCIBILITY AND ARTIFACT MAP

The frozen data root is `exp2_v2/biased_news/data/coin_city_stable_relationship_claude_n250_v4/`, and the experiment-code root is `exp2_v2/biased_news/`. Paths below are relative to those roots unless a repository-level path is shown.

Design and preflight. Under the data root, `design/manifest.json` and `design/validation.json` freeze and validate the primary design; `design/tasks_abc_wrong_context.manifest.json` records the misleading arm, and `design/tasks_abc_symbol_context.manifest.json` records the balanced, episode-randomized symbol control.

Generator and task builders. Under the code root, the generator is `engine/coin_city_stable_relationship_claude_n250.py`; the builders are `eval/build_coin_city_stable_relationship_claude_n250.py` and `eval/build_coin_city_wrong_context_tasks.py`. The symbol implementation is `engine/coin_city_symbol_context_arm.py`, with task builder `eval/build_coin_city_symbol_context_tasks.py`.

Execution and parsing. The manifest-locked Opus 4.8 wrapper is `eval/run_coin_city_stable_relationship_claude_n250.py`; shared append-only execution and parsing are in `eval/run_frozen_task_shard.py`. Model-specific runners and `eval/run_coin_city_wrong_context.py` preserve configurations and response records under the data root. The symbol runner is `eval/run_coin_city_symbol_context.py`; its contemporaneous comparison arms are under `responses/symbol_control_matched_20260813/`. The open-model runner is `eval/run_coin`

city_symbol_context_open_model.py; the module-specific Qwen and Llama responses, hosted symbol arms, launchers, and run manifests are under local_results/symbol_context_model_comparison_20260824/ relative to the code root.

Primary analysis. analysis/analyze_coin_city_stable_relationship_multimodel.py under the code root produces analysis/multimodel_results.json under the data root, the authority for the primary-arm metrics and paired contrasts.

Arbitrary-symbol analysis. analysis/analyze_coin_city_symbol_context.py supplies the common estimator; local_results/symbol_context_model_comparison_20260824/build_final_comparison.py writes the paper authority analysis/symbol_context_model_comparison_20260824.json. Focused design checks are in tests/test_coin_city_symbol_context.py; paper/generate_coin_city_symbol_context_tables.py generates both appendix tables from the final comparison.

Reference-selection analysis. analysis/analyze_coin_city_reference_selection.py under the code root recomputes the partial correlations of Section C.6 from the frozen responses and writes analysis/reference_selection_results.json under the data root; it makes no model call. Checks are in tests/test_coin_city_reference_selection.py, including one that requires the partial correlation and the symbol-arm $\Delta\rho$ contrast to name the same systems. paper/generate_coin_city_reference_selection_tables.py writes both appendix tables.

Probe-target ceilings. analysis/compute_bayes_ceilings.py under the code root derives the Bayes-optimal ceiling for each probe target in Table 49 and the forecast-MAE floors drawn in Figure 5, importing the generator constants from the engine module so they cannot drift, and writes analysis/bayes_ceilings.json under the data root. It reads no response file and makes no model call, and fails closed if its closed form and its simulation of the generator’s own row logic disagree. Checks are in tests/test_bayes_ceilings.py.

Misleading arm and paper outputs. The repository-level script paper/generate_coin_city_split_figures.py recomputes misleading-arm metrics from the final response files and scoring key. paper/generate_coin_city_appendix_tables.py generates the primary tables and pairwise contrasts.

Generator and deployment robustness. engine/coin_city_robustness_population.py and eval/build_coin_city_robustness_population.py build the fresh-seed population; its design, responses, scheduler ledgers, and analysis are under data/coin_city_population_075_040_sd010_v1/. analysis/analyze_coin_city_rationales.py writes analysis/rationale_audit_v1.json under the primary data root. eval/run_coin_city_gpt56_k0_repeat.py and analysis/analyze_coin_city_gpt56_k0_repeat.py isolate and score the hosted repeat. paper/generate_coin_city_robustness_tables.py renders both appendix tables from the saved analysis artifacts.

D EXPERIMENT 3: CONTROLLED RELATIONSHIP TRANSFER

As the controlled transfer step, Experiment 3 asks whether outcome-based training learns to select and apply an episode-matched relationship beyond its training setting. The primary Qwen3-8B endpoint jointly changes the domain and response mechanism; Qwen3-4B and Llama-3.1-8B provide scale and architecture comparisons. The original three-seed design and results are retained below as an execution record. Section D.1.6 reports the later eight-seed Qwen extension used for the main-text conclusions; the Llama and structureless comparators retain their original seed counts.

D.1 DESIGN AND EVALUATION

D.1.1 DATA-GENERATING PROCESSES

Training uses only Structure A, the direct relationship from Experiment 2. On the displayed scale of domain d ,

$$y_{t+1} = b_d + g_j x_t + \epsilon_t, \quad \epsilon_t \sim \mathcal{N}(0, \sigma_d^2). \quad (7)$$

A pulse acts in the next period and has no residual effect under zero subsequent input. The unseen evaluation mechanism, Structure B, is

$$m_{t+1} = \rho m_t + g_j x_t, \quad (8)$$

$$y_{t+1} - b_d = \phi(y_t - b_d) + \lambda m_{t+1} + \epsilon_t. \quad (9)$$

This is a linear-Gaussian latent-state process in the sense of [Kalman \(1960\)](#): it introduces a persistent mediator and outcome carry-over, so a single pulse has both immediate and delayed effects. Model prompts disclose neither equation, parameter values, internal variable, nor structure label.

Coin City uses $b = 50$, range $[0, 100]$, driver pool $\{-12, -8, -4, 0, 4, 8, 12\}$, and $\sigma = 1.6$. Held-out Coin Harbor uses $b = 180$, range $[40, 320]$, driver pool $\{-24, -16, -8, 0, 8, 16, 24\}$, and $\sigma = 4.5$, together with new terminology, units, and qualitative descriptions. Evaluation gains follow $g \sim \text{Uniform}(.22, .92)$. Structure B independently samples $\rho \sim \text{Uniform}(.48, .74)$, $\phi \sim \text{Uniform}(.18, .42)$, and $\lambda \sim \text{Uniform}(.68, .94)$. A reference and its matched target share these parameters but have independent inputs and noise.

Table 54: **Exact Coin City–Coin Harbor domain contrast.** Both surfaces instantiate the same Structure A and B simulators and task template; Coin Harbor changes the semantic setting, time unit, numeric scale, interventions, noise, and pathway descriptions.

Field	Coin City (train and test)	Coin Harbor (test only)
Driver $\rightarrow$ outcome	Net campaign news $\rightarrow$ candidate support	Net shipping bulletin $\rightarrow$ cargo-flow index
Time unit	Week	Tide cycle
Resting level / support	50 / $[0, 100]$	180 / $[40, 320]$
Observed-input pool	$\{-12, -8, -4, 0, 4, 8, 12\}$	$\{-24, -16, -8, 0, 8, 16, 24\}$
Evaluation shocks	$\{-12, -6, 0, 6, 12\}$	$\{-24, -12, 0, 12, 24\}$
Observation-noise SD	1.6 support points	4.5 cargo-index points
Structure A description	Synchronized citywide alerts; reactions occur mostly in the same week	Synchronized fleetwide radio; responses occur mostly in the same tide cycle
Structure B description	Neighborhood conversations and local organizations; reactions build gradually and persist	Dockside crews and harbor groups; responses build gradually and persist

D.1.2 TRAINING EPISODES AND MATCHED CONTROL


Each training prompt contains two 14-case Coin City reference systems, both governed by Equation 7. One draws $g \in [.18, .46]$ and the other $g \in [.68, .98]$; their order is randomized. The target shares exactly one reference gain, selected in balanced fashion, and shows $k \in \{0, 2, 4, 8\}$ independently generated cases. Its national/local background description identifies the matching weak or strong response as in Experiment 2. There are 1,200 rows at each k , or 4,800 per training run.

The episode-matched and population-prior datasets have byte-identical prompts. Episode-matched reward uses the target’s noise-free forecasts. For the control, targets are deranged within each k stratum with a fixed permutation: no row keeps its own target, while the complete multiset of target vectors is exactly preserved. Consequently the control can learn the output format and population distribution but cannot improve reward by using that row’s references or target history. The Qwen3-4B structureless diagnostic additionally scrambles displayed drivers and rewards a flat intervention response. It isolates task-format learning and is excluded from the episode-matched-versus-prior estimand.

Table 55: **One training episode, and what each policy is rewarded against.** Panel A is the prompt for episode `train:1:k2`, abridged from the rendered text. Panel B is the prompt-key pairing used during training, drawn for four episodes at one evidence depth: P_i is episode i ’s prompt and y_i^* its ten-answer key. Episode-matched and population-prior training use byte-identical prompts and the same multiset of keys; only the pairing differs. The control shifts cyclically across all 1,200 episodes at a depth, so its key is a real key drawn from the same distribution but never this prompt’s, and no reward can be earned by reading this row’s references. Base is untrained and receives no supervision.

A. Shared prompt. Coin City; resting level 50.0, bounded to $[0, 100]$; no equation or coefficient is shown.
 REFERENCE 1 "...immediate national alerts ...react **strongly**" 14 wk: $-8.0 \rightarrow 43.8, -12.0 \rightarrow 40.9, \dots, 0.0 \rightarrow 51.6$
 REFERENCE 2 "...local civic briefings ...react **weakly**" 14 wk: $+8.0 \rightarrow 52.6, +4.0 \rightarrow 51.3, \dots, -12.0 \rightarrow 46.9$
 TARGET ($k=2$) "...immediate national alerts ...react **strongly**" $-8.0 \rightarrow 42.3, +4.0 \rightarrow 53.9$
 QUERY ten counterfactuals A-J: shock $\in \{-12, -6, 0, 6, 12\} \times \text{horizon} \in \{1, 3\}$; return {"forecasts": $[A, \dots, J]$ }

B. How prompts are paired with answer keys during training.



D.1.3 EVALUATION INTERFACE AND QUALITATIVE HINT

Every evaluation prompt contains a complete 14-period Structure A reference, a complete 14-period Structure B reference, and a target governed by one of them. For Coin City the two reference descriptions are:

Residents receive campaign developments through synchronized citywide alerts. Reactions occur mostly during the same week as the news.
 Campaign developments usually spread through neighborhood conversations and local organizations. Reactions may build gradually and persist after the original news has faded.

Coin Harbor uses the matched fleetwide-radio versus dockside-network descriptions. The target receives the correct description, no description, or the other pathway’s description. The three prompt variants reuse the exact reference and target trajectories, scenario grid, and gold forecasts; the manifest records a common numeric hash for every cue triplet.

Each task requests ten counterfactuals: shocks $\{-2u_d, -u_d, 0, u_d, 2u_d\}$ at horizons one and three, with $u_d = 6$ in Coin City and $u_d = 12$ in Coin Harbor. The shock occurs in the next period and all later inputs are zero. The response vector consists of the eight nonzero-shock minus zero-shock contrasts. The full endpoint contains

$$2 \text{ domains} \times 2 \text{ structures} \times 3 \text{ hints} \times 4 \text{ } k \text{ values} \times 30 \text{ episodes} = 1,440$$

paired prompts. Training and evaluation seed bands begin at 81,000,000 and 91,000,000, respectively.

Exact task template and prompt variants. The following is the task string emitted by the prompt generator before the model-specific chat wrapper is applied. Bracketed uppercase fields below denote row-specific substitutions; they are not shown to the model. Line wrapping is typographic, while the blank-line-separated sections and all fixed wording are part of the prompt.

```
Forecasting task in [DOMAIN].

The resting level of [OUTCOME] is [BASELINE]. Values are bounded to [LOWER,
UPPER].

The reference systems are complete examples. The target system may share a
stable response pattern with one reference. Infer from the information
shown; no response equation or coefficient is provided.

REFERENCE SYSTEM 1
Background: [REFERENCE 1 CONTEXT]
[REFERENCE 1 TABLE: period, driver, and 14 observed outcomes]

REFERENCE SYSTEM 2
```

```

Background: [REFERENCE 2 CONTEXT]
[REFERENCE 2 TABLE: period, driver, and 14 observed outcomes]

TARGET SYSTEM
[TARGET BACKGROUND, WHEN PRESENT]
[TARGET TABLE WITH  $k$  OBSERVATIONS, OR THE EXACT NO-OBSERVATION SENTENCE]

COUNTERFACTUAL FORECASTS
For each scenario, apply the stated [DRIVER] in the next [PERIOD] and then
set it to 0 in later periods. Forecast [OUTCOME] at the stated horizon.

A: shock=[VALUE]; horizon=1 ... E: shock=[VALUE]; horizon=1
F: shock=[VALUE]; horizon=3 ... J: shock=[VALUE]; horizon=3

Return only strict JSON with exactly ten finite numbers in scenario order:
{"forecasts": [A,B,C,D,E,F,G,H,I,J]}

```

The renderer realizes the variants as follows. Training always includes the correct strong/weak background and uses the same input string for episode-matched and population-prior training; only the hidden reward target differs. At evaluation, correct inserts the target pathway’s description, none omits the background line entirely, and misleading inserts the other pathway’s description. Every cue triplet otherwise has identical tables, interventions, and targets. If $k = 0$, the target block contains the exact sentence *No target-system observations are available yet.*; otherwise it contains the same table format as the references with the first k rows. The Qwen3-4B structureless diagnostic retains this shell but swaps every displayed $-12.0 \leftrightarrow 12.0$ and $-8.0 \leftrightarrow 8.0$ token through a collision-safe substitution. Untrained bases receive the same held-out prompts as trained models. Thus cue, reward-arm, structureless, and base comparisons do not introduce an unreported instruction change.

D.1.4 ESTIMANDS, ROSTER, AND VALIDATION CRITERIA

The original factorial crosses Qwen3-4B, Qwen3-8B, and Llama-3.1-8B, two reward arms, and three training seeds, for 18 runs. Three Qwen3-4B structureless runs and three untrained endpoints provide diagnostic comparisons. The primary estimand is population-prior minus episode-matched response MAE on Structure B in Coin Harbor with the correct hint, averaging the four k values equally. Positive values favor episode-matched training. The same estimator defines factorial contrasts for Structure A/Coin Harbor (domain transfer), Structure B/Coin City (structural transfer), and the in-distribution cell. Paired absent-minus-correct and misleading-minus-correct contrasts measure hint use. Their paired $k = 8$ minus $k = 0$ difference tests whether observations override a misleading prior; negative values indicate a shrinking misleading-hint penalty.

Every reported round-300 endpoint uses five draws per prompt at temperature .7; temperature-zero checkpoints serve only as online training monitors.

Each episode’s score is the mean of its five per-draw response MAEs, so “averaging draws” averages errors and not forecasts. The hierarchical bootstrap resamples the three training seeds at the top level and paired episodes within each selected seed. Parse failures receive the full observable-range error penalty rather than being discarded. That penalty is the clip width—280 points in Coin Harbor, roughly fifty times a typical error—so it carries large leverage: a single unparsed draw of 150 raises the Qwen3-8B base joint-cell MAE from 5.26 to 5.72, which is most of that arm’s reported gap to matched training. Both trained arms parse completely in that cell, so the primary estimand is unaffected; the base comparison is not. The launch ledger contains all 24 endpoints and 172,800 stochastic rows; 172,732 parse (99.96%), and every endpoint has at least 99.82% coverage.

Original three-seed uncertainty analysis. The original hierarchical interval resamples three top-level clusters, and a nonparametric cluster bootstrap at $G = 3$ is known to understate between-cluster uncertainty (Cameron et al., 2008). We therefore report a deliberately conservative alternative alongside it: a Student- t interval on the three seed-level means with two degrees of freedom, which uses the same paired episode differences but charges the full small-sample penalty for estimating a variance from three numbers. Under that alternative none of the twelve transfer-cell intervals excludes zero, the primary Qwen3-8B contrast becoming .251 [−.259, .762] and the Qwen3-4B comparator .319 [−.120, .758]. Nine of the twelve intervals that exclude zero under the hierarchical bootstrap therefore do not survive the t alternative, and we do not treat any original three-seed interval as robust to the choice of small-cluster procedure.

In that original roster, the sign pattern was consistent without estimating a variance. Across the four transfer cells and three training seeds there are twelve seed-level contrasts per model. All twelve favor episode-matched training for Qwen3-8B and all twelve for Qwen3-4B, so all 24 Qwen seed-by-cell contrasts point the same way; these reuse training seeds across cells and are not independent replications. The corresponding count for Llama-3.1-8B is three of twelve. This sign pattern motivated the later precision extension; its interpretation is superseded by the eight-seed results in Section D.1.6.

Before training, the dataset passed 17 fail-closed checks: exact row counts and factorial balance; byte-identical episode-matched/population-prior prompts and held-out datasets; exact within- k target marginals with no fixed reward links; absence of Structure B from training prompts and metadata; cue-triplet numeric identity; train/evaluation identifier disjointness; hidden-equation leak scans; strict gold round trips; and oracle separation from the blind pathway mixture. Mean A-B response separation is 2.65 points in Coin City and 4.80 in Coin Harbor. Across both tokenizers, the longest prompt is 1,071 tokens against the configured 2,304-token limit.

We estimate population-prior-minus-episode-matched intervals in all four transfer cells, paired intervals for both cue penalties at every k , and the paired $k = 8$ minus $k = 0$ change in misleading-minus-correct error to evidence override. The protocol archive records versioned input, analysis, and output hashes. Full optimization and compute parameters are in Section F.

D.1.5 ORIGINAL THREE-SEED SCALE AND ARCHITECTURE COMPARATORS

This subsection preserves the original estimates. Its Qwen contrasts and cross-size synthesis precede the eight-seed extension and do not establish the current transfer claim; see Section D.1.6.

Qwen3-4B scale replication. The Qwen3-4B comparison includes three seeds in each of the episode-matched, population-prior, and structureless arms. Each is evaluated on the same 1,440 tasks at round 300 with five stochastic draws per task; the untrained base uses the same endpoint. Parse rates are 100% for every trained arm and 99.97% for the stochastic base.

Joint-cell MAE is 6.15, 6.76, 4.81, and 4.49 for base, structureless, population-prior, and episode-matched training, respectively. The prior-minus-episode-matched difference is .319 [.134, .489]; episode-matched training has the lowest point estimate in all four cells, and all four original intervals exclude zero. The original Qwen3-8B estimate is likewise positive, .251 [.058, .470], giving the same interval-excluding-zero result at both tested Qwen sizes.

Qwen3-family summary. In the hardest joint-transfer cell, Qwen3-8B response MAE is 5.72 for the untrained base, 4.71 after population-prior training, and 4.46 after episode-matched training. The prior-minus-matched contrasts are .319 [.134, .489], .251 [.058, .470], and .120 [−.027, .261] at 4B, 8B, and 14B. An equal-weight synthesis across the Qwen3 family is .230 [.128, .333], a 4.9% MAE reduction. All three size-level estimates, eight of nine size-seed estimates, and 11 of 12 size-cell estimates agree in direction. The original estimates suggested an advantage that attenuated with scale; the eight-seed extension does not sustain that cross-size transfer conclusion. At three seeds and 8B, about one fifth of the improvement over the untrained base comes from preserving which evidence belongs to which outcome rather than learning the format or common answers.

Averaged over evidence depths, omitting the hint raises episode-matched response MAE by .151 [.079, .235], and a misleading hint raises it by .313 [.228, .415]. The misleading-hint penalty is smaller at $k = 8$ than at $k = 0$, but the evidence-override contrast −.116 [−.309, .100] includes zero.

Llama-3.1-8B architecture comparator. In the joint-transfer cell the prior-minus-episode-matched training contrast is positive, .420 [−.115, .958], but its interval spans zero. The three seed-level effects are +.762, +.622, and −.125, so this result does not establish a robust cross-architecture training advantage.

The joint cell is also the most favorable of the four Llama cells, and reporting it alone would misrepresent the factorial, so we give all four. The contrasts are −.351 [−.680, −.130] in the in-distribution cell, −.167 [−.862, .306] under structure-only shift, −.581 [−.868, −.243] under domain-only shift, and +.420 [−.115, .958] under joint shift. In the two Structure A cells the population-prior arm is therefore better than the episode-matched arm by an interval that excludes zero, and only three of the twelve Llama seed $\times$ cell contrasts favor episode-matched training. On this architecture the factorial runs against the hypothesis in the cells whose response mechanism

was seen during training, and recovers a positive point estimate only where both shifts apply. We report this as a failure of the episode-matching effect to replicate on Llama-3.1-8B rather than as a directionally-consistent-but-underpowered result. The later Qwen extension also weakens the original cross-size transfer interpretation.

Llama nevertheless uses the qualitative pathway cue at the stochastic endpoint: absent-minus-correct MAE is .219 [.123, .313] and misleading-minus-correct is .210 [.092, .320]. Its misleading-hint penalty shrinks from $k = 0$ to $k = 8$ by $-.621 [-.957, -.323]$. All three cue-related intervals exclude zero. These contrasts indicate sensitivity to the supplied mechanism cue and reduced misleading-cue effects as observations accumulate.

Additional Llama-3.1-70B scale check. To distinguish the Llama-8B result from a general failure of the architecture, we launched a Llama-3.1-70B extension using the same training and held-out data. We report a fixed exploratory snapshot: seed 42 at round 100 of 300, temperature zero, one draw on each of the same 1,440 held-out tasks. The intervals below bootstrap paired tasks conditional on this single training seed; they do not measure between-seed uncertainty and are excluded from the primary three-seed factorial estimates.

The capacity result is already pronounced on the response mechanism seen in training. In the correct-hint Coin City/Structure A cell, response MAE falls from 4.114 for the untrained 70B checkpoint to .410 under population-prior training and .450 under episode-matched training at round 100. For comparison, the completed round-300 Llama-8B temperature-zero endpoints average 4.159 and 4.328 in the corresponding arms. Thus, after only one third of the planned schedule, 70B has roughly one tenth the direct-response error of 8B; both 70B arms parse every output and classify Structure A correctly on every task in this cell.

This snapshot does not yet establish the treatment-specific transfer effect. Across the in-distribution, domain-only, structure-only, and joint-transfer cells, the round-100 prior-minus-episode-matched contrasts are respectively $-.040 [-.098, .020]$, $+.132 [-6.166, 6.820]$, $-.068 [-.156, .018]$, and $+1.778 [-.143, 5.406]$. Positive values favor episode-matched training. The large joint-cell point estimate is driven by one additional unparsed population-prior response: among tasks parsed by both arms, the contrast is only $+.039 [-.209, .276]$, and Structure B classification is near zero in both arms. The in-distribution contrast also changes from $+.291 [.209, .371]$ at round 50 to $-.040$ at round 100. We therefore draw the narrow conclusion from this snapshot that scale lets Llama execute the seen numerical relationship far more accurately, while episode-conditioned structural transfer remains unresolved. This additional snapshot is distinct from the primary three-seed, five-draw round-300 comparison.

Table 56: **Stochastic Coin City appendix comparators in the joint-transfer cell, original three-seed roster.** The updated eight-seed Qwen results appear in Table 60. Entries are response MAE under the correct hint after averaging five draws per task; lower is better. Positive prior-minus-episode-matched contrasts favor episode-matched training. Intervals use the paired seed-first bootstrap with a fixed cell-specific bootstrap seed.

Model	Base	Prior	Episode-matched	Δ [95% CI]
Qwen3-4B	6.15	4.81	4.49	0.32 [0.13,0.49]
Llama-3.1-8B	9.43	4.69	4.27	0.42 [-0.12,0.96]

Three further outputs of the same three-seed analysis were generated but never reported. They are included here for completeness, on the original roster and with the same estimator; none of them bears on the eight-seed transfer claim of Section D.1.6.

Prespecified endpoint result. In the prespecified joint domain-plus-structure transfer cell, population-prior minus episode-matched response-MAE estimates (positive favors episode-matched reward) under five-draw stochastic decoding were Qwen3-4B 0.32 [0.13,0.49]; Qwen3-8B 0.25 [0.06,0.47]; Llama-3.1-8B 0.42 [-0.12,0.96]. The table reports the prespecified seed-first, paired-episode intervals. In the same joint-transfer cell, the prespecified stochastic cue contrasts were

Qwen3-4B: absent-minus-correct 0.15 [0.08,0.24], misleading-minus-correct 0.31 [0.23,0.41], and the $k = 8$ minus $k = 0$ change in misleading penalty -0.12 [-0.31,0.10]; Qwen3-8B: absent-minus-correct 0.12 [0.01,0.23], misleading-minus-correct 0.31 [0.21,0.41], and the $k = 8$ minus $k = 0$ change in misleading penalty -0.01 [-0.28,0.25]; Llama-3.1-8B: absent-minus-correct 0.22 [0.12,0.31], misleading-minus-correct 0.21 [0.09,0.32], and the $k = 8$ minus $k = 0$ change in misleading penalty -0.62 [-0.96,-0.32]. Negative values for the final contrast mean that target observations reduced the misleading-hint penalty.

Table 57: **Qwen3-4B across all four transfer cells, original three-seed roster.** Table 56 reports the joint cell alone; the other three cells and the structureless diagnostic appear only here. Entries are response MAE under the correct hint, averaged over five draws per task, lower is better.

Test cell	Base	Structureless	Prior	Episode-matched	Prior — episode-matched [95% interval]
City / A (ID)	3.82	2.56	0.79	0.58	0.21 [0.11, 0.32]
Harbor / A (domain)	8.53	5.23	1.91	1.68	0.23 [0.04, 0.38]
City / B (structure)	2.66	3.59	2.53	2.34	0.19 [0.02, 0.35]
Harbor / B (joint)	6.15	6.76	4.81	4.49	0.32 [0.13, 0.49]

Table 58: **Hint-condition contrasts by evidence depth, original three-seed roster.** A correct hint names the generating structure, an absent hint omits it, and a misleading hint names the wrong one. Positive contrasts are error increases relative to the correct hint. The Qwen3-8B rows are the original three-seed estimates; its eight-seed transfer result is separate and appears in Table 60.

Model	k	Correct	Absent	Misleading	Absent — correct [95% CI]	Mislead. — correct [95% CI]
Qwen3-4B	0	4.49	4.67	4.89	0.18 [0.01,0.37]	0.40 [0.20,0.61]
Qwen3-4B	2	4.44	4.53	4.69	0.09 [-0.06,0.27]	0.24 [0.12,0.38]
Qwen3-4B	4	4.53	4.68	4.85	0.16 [0.04,0.29]	0.33 [0.16,0.51]
Qwen3-4B	8	4.51	4.69	4.79	0.18 [0.03,0.33]	0.28 [0.15,0.46]
Qwen3-8B	0	4.40	4.58	4.73	0.18 [-0.06,0.44]	0.33 [0.08,0.58]
Qwen3-8B	2	4.60	4.68	4.83	0.07 [-0.10,0.25]	0.22 [0.02,0.43]
Qwen3-8B	4	4.47	4.74	4.85	0.26 [0.04,0.52]	0.37 [0.19,0.57]
Qwen3-8B	8	4.35	4.30	4.66	-0.05 [-0.21,0.10]	0.31 [0.16,0.47]
Llama-3.1-8B	0	3.82	4.19	4.35	0.37 [0.17,0.60]	0.52 [0.29,0.78]
Llama-3.1-8B	2	4.06	4.24	4.36	0.18 [0.06,0.32]	0.30 [0.11,0.52]
Llama-3.1-8B	4	4.65	4.85	4.76	0.19 [0.03,0.36]	0.11 [-0.11,0.32]
Llama-3.1-8B	8	4.53	4.66	4.44	0.13 [-0.15,0.42]	-0.10 [-0.33,0.10]

Table 59: **Per-cell diagnostics by model, domain and structure, original three-seed roster.** Δ_{matched} is population-prior minus episode-matched, positive favoring episode-matched training. The structureless column is a diagnostic arm that scrambles the displayed drivers and is available for Qwen3-4B only.

Model	Domain	Structure	Base	Pop. prior	Episode-matched	Δ_{matched} [95% CI]	Structureless
Qwen3-4B	City	A	3.82	0.79	0.58	0.21 [0.11,0.32]	2.56
Qwen3-4B	City	B	2.66	2.53	2.34	0.19 [0.02,0.35]	3.59
Qwen3-4B	Harbor	A	8.53	1.91	1.68	0.23 [0.04,0.38]	5.23
Qwen3-4B	Harbor	B	6.15	4.81	4.49	0.32 [0.13,0.49]	6.76
Qwen3-8B	City	A	3.96	0.53	0.45	0.08 [0.00,0.19]	–
Qwen3-8B	City	B	2.30	2.36	2.33	0.03 [-0.01,0.07]	–
Qwen3-8B	Harbor	A	7.30	2.03	1.79	0.24 [0.03,0.55]	–
Qwen3-8B	Harbor	B	5.72	4.71	4.46	0.25 [0.06,0.47]	–
Llama-3.1-8B	City	A	5.84	0.73	1.08	-0.35 [-0.68,-0.13]	–
Llama-3.1-8B	City	B	4.54	2.31	2.48	-0.17 [-0.86,0.31]	–
Llama-3.1-8B	Harbor	A	11.16	1.69	2.28	-0.58 [-0.87,-0.24]	–
Llama-3.1-8B	Harbor	B	9.43	4.69	4.27	0.42 [-0.12,0.96]	–

D.1.6 LATER EIGHT-SEED PRECISION EXTENSION

Qwen3-8B and Qwen3-4B were extended to seeds 42–49 in both reward arms, holding fixed the 300-round endpoint, five-draw decoder, data, and estimands. The added seed rosters and four estimators were recorded before their added endpoints existed. The seed count was chosen after seeing the original three-seed results; this is a later precision extension, not an eight-seed preregistration of the original study. The Qwen3-4B extension followed the disclosed Qwen3-8B interim look at seven seeds, whose joint-cell estimate was .108 [−.035, .265] while a crashed seed-48 run was repeated. Eight seeds remained the fixed endpoint. The amendment and interim records are in `exp3_training_transfer/five_seed_extension/`.

Table 60: **Eight-seed prior-minus-episode-matched response MAE.** Positive values favor episode matching; brackets give hierarchical 95% intervals. All cells use the correct cue and average evidence depths equally.

Cell	Qwen3-8B	Qwen3-4B
In distribution	.060 [.018, .112]	.067 [−.056, .194]
Domain shift	.110 [−.026, .269]	.087 [−.091, .255]
Structure shift	−.003 [−.056, .040]	.052 [−.075, .179]
Joint shift	.126 [−.002, .265]	.136 [−.032, .298]

The Qwen3-8B joint estimate falls from .251 [.058, .470] to .126 [−.002, .265]; all four Qwen3-4B intervals now span zero. The sole cell supported by all four procedures is Qwen3-8B in distribution: .060 [.018, .112], Student-*t* interval [.005, .115], wild-cluster $p = .0089$, and 8/8 positive seeds (one-sided sign-test $p = .0039$). The hierarchical bootstrap uses 5,000 repetitions; the Student-*t* interval has seven degrees of freedom; the wild cluster bootstrap uses 9,999 Webb-weight repetitions; the sign test is exact. Seed-level contrasts and all four procedures for every cell are retained in `qwen3_8b_step300_stochastic_eight_seed.json` and `qwen3_4b_step300_stochastic_eight_seed.json` under the Coin City reports root. These artifacts supply `tables/exp3_rebuilt_data.tex`.

The wrapper deviation record separately documents changed checkpoint-saving flags in the Qwen3-8B added-seed block and the crossover reruns of seeds 42–44. Their mean joint-cell contrast is .255, close to the original .251, although individual seed contrasts change. This supports robustness of the mean at three seeds, not exact equivalence of the wrappers. Those diagnostic runs are excluded from the endpoint roster. Qwen3-4B’s added runs restored the parent flags before producing endpoints. The original Llama-8B and Qwen3-4B structureless comparisons remain three-seed analyses.

D.1.7 STRUCTURE-SELECTION EXPERIMENT

Protocol `coin_city_structure_selection_v1`. The parent Coin City design cannot test whether cue-driven selection transfers on functional form, because its training set contains only the direct structure. This dataset makes structure selectable. Each episode displays one direct and one mediated reference whose horizon-1 responses are matched exactly—the direct gain is scaled by the mediated system’s λ , since a mediated horizon-1 response is $\lambda g s$ against a direct $g s$ —so persistence at horizon 3 is the only discriminating feature and no model can select on response magnitude instead. Mediated parameters are resampled every episode, so a memorised template does not solve the task.

Solvability was verified before any data was generated, the check the parent design lacked: a one-feature classifier on the coefficient of the lagged outcome, applied to the fourteen reference rows, separates the structures with $d = 4.47$ in Coin City and 3.10 in Coin Harbor and classifies at 99.0% and 97.2%. Choosing correctly is worth 2.18 response-MAE points.

Training is Coin City only, semantic labels only, correct cue only, balanced 2,400/2,400 across target structures. Held out entirely: Coin Harbor, and an arbitrary-label variant in which each structure is denoted by a symbol whose mapping is redrawn every episode, so it can be recovered only from the reference trajectories. Twenty-eight fail-closed checks gate the dataset, including that the dimension under test varies in training and that the arbitrary symbols never appear in a training prompt. Evaluation is fully crossed over two domains, two label kinds, two target structures, three cue conditions and four evidence depths, 120 tasks per cell.

The primary measure is the implied persistence ratio, mean $|\text{predicted horizon-3 response}|$ over mean $|\text{predicted horizon-1 response}|$, computed as a ratio of means so it stays finite when individual horizon-1 predictions are near zero. Its true value is 0 for direct targets and approximately .65 for mediated. The cue-following index is the difference in that ratio between episodes whose cue names a mediated structure and those whose cue names a direct one.

The untrained reference was evaluated on the same 2,880-task set at 99.91% parse. Its index is .020 and -.005 on Coin City under semantic and arbitrary labels and .068 and .002 on Coin Harbor, with implied persistence between 1.05 and 1.21 in every condition: the untrained model predicts full persistence irrespective of cue or truth.

Results are a three-seed pilot of the episode-matched arm, reported as a pilot. No claim about the episode-matched minus population-prior contrast is made at three seeds; that contrast requires the eight-seed treatment applied to the parent roster. A mid-training interim look at the greedy monitors was taken at the author’s direction and is recorded in `INTERIM_LOOK_DISCLOSURE.md`; it agreed with the endpoint.

D.1.8 COMPONENT DIAGNOSIS WITH PAIRED CUE FLIPS

The preceding RL pilot motivates a fixed, exploratory decomposition, frozen before these inference outputs. We sample 96 fresh latent episodes, 24 in each of the two domain $\times$ two label cells. Each episode is presented with both cue assignments under four interfaces, giving 768 prompts per checkpoint: (i) identify the named reference without forecasting; (ii) forecast with exact unit-shock response coefficients at horizons 1 and 3 supplied for both references; (iii) forecast with the selected reference and its coefficients supplied; and (iv) the original forecasting task. The coefficients are latent simulator summaries, not the requested ten forecasts; this assistance tests whether a supplied relationship is usable, not whether exact coefficients are recoverable from fourteen noisy observations. We use $k = 0$ target observations so numerical evidence is identical across the original-task cue flips. This diagnostic does not test integration with target observations at $k > 0$.

The fixed roster comprises Qwen3-8B base, its matched and fixed-prior RL checkpoints at seeds 42–44, and an untrained Qwen3-32B feasibility comparison. All 6,144 greedy generations parse. Selection-only accuracy is 100% for every checkpoint in every domain/label cell. Table 61 reports the frozen scores, including every assistance condition. At 8B, supplying patterns does not reduce average response error in any domain/label cell for the base or either trained arm. This locates a failure beyond identifying the reference named by the cue; it does not uniquely separate numerical application from sensitivity to the assisted prompt format.

Domain	Labels	Model	Name (%)	Original	Both	Selected
City	semantic	8B base	100	2.105	2.393	2.413
City	semantic	8B matched RL	100	1.053	2.213	1.440
City	semantic	8B prior RL	100	0.990	3.657	1.224
City	semantic	32B base	100	4.110	1.710	1.423
City	arbitrary	8B base	100	2.117	2.710	2.961
City	arbitrary	8B matched RL	100	1.130	3.793	1.567
City	arbitrary	8B prior RL	100	1.084	4.156	1.342
City	arbitrary	32B base	100	2.671	1.993	1.604
Harbor	semantic	8B base	100	4.967	7.841	7.872
Harbor	semantic	8B matched RL	100	3.318	9.657	8.633
Harbor	semantic	8B prior RL	100	2.903	7.938	7.197
Harbor	semantic	32B base	100	7.543	5.288	4.730
Harbor	arbitrary	8B base	100	5.353	12.155	8.826
Harbor	arbitrary	8B matched RL	100	3.418	13.837	9.303
Harbor	arbitrary	8B prior RL	100	3.015	10.971	8.224
Harbor	arbitrary	32B base	100	7.012	5.242	5.200

Table 61: **Reference naming succeeds while forecasting remains difficult.** Name is selection-only accuracy; Original, Both, and Selected are response MAE for the three forecast interfaces. Trained rows average all three seeds, not selected endpoints. Each checkpoint contributes 24 paired episodes per cell. Domain scales differ, so compare interfaces within a domain.

Forecast structure accuracy assigns each predicted eight-coordinate response vector to the closer of the episode’s direct and mediated simulator templates in MAE; ties count as incorrect. The paired cue-change calibration is

$$\alpha = \frac{1}{N} \sum_{i=1}^N \frac{(\hat{r}_{i,M} - \hat{r}_{i,D})^\top (r_{i,M} - r_{i,D})}{\|r_{i,M} - r_{i,D}\|_2^2},$$

where r subtracts the zero-shock forecast separately within each horizon, and M, D index the mediated and direct cue assignments. Zero denotes invariance to the cue; one denotes the simulator’s change along this direction, but need not imply a correct full response vector. This differs from the preceding persistence-ratio index, whose selecting reference is approximately .65. Matched RL’s mean α is .000, .017, .136, and .008 for City semantic, City arbitrary, Harbor semantic, and Harbor arbitrary, respectively. The corresponding fixed-prior means are $-.008$, .027, .085, and .035; these three-seed descriptions are not a new confirmatory comparison between training arms.

Stronger-model feasibility and output semantics. The frozen scorer clips forecasts to the valid domain range before computing responses. Inspection revealed that Qwen3-32B often returns deviations from rest instead of absolute levels; clipping those values distorts the response vector, especially in Harbor, whose lower bound is 40. We therefore report a separately labeled *post hoc* raw-output audit (Table 62), without replacing the frozen scores or counting offset-corrected answers as correct forecasts. With the selected pattern supplied in City/semantic, raw response MAE is .384 versus 4.794 on the original interface. Complete relative vectors are accurate to .01 in 52.1% of cases, but 95.8% omit the resting level. This shows partial conditional numerical ability, not a solution of the original task. The 8B assistance failures remain predominantly response-pattern errors after a common-offset correction.

D.1.9 SUPERVISED LEARNABILITY AND ITS TRANSFER BOUNDARY

We train two fixed forecast-SFT recipes from the same Qwen3-8B base on the original 4,800 balanced Coin City/semantic prompts, with simulator-correct absolute forecasts as targets. Both use one epoch, effective batch 16, 300 optimizer updates, completion-only token cross entropy, LoRA rank 32 and alpha 64 on all seven projection modules, AdamW (.9, .95), no weight decay, constant learning rate, and gradient norm cap 1.0. Rates are 10^{-6} (the RL rate) and 2×10^{-5} (an optimization-feasibility condition). Neither Harbor nor arbitrary labels appear in training. Final endpoints alone are evaluated, on the unchanged 768-prompt component diagnostic.

Domain	Labels	Original	Both	Selected	Exact (%)	Omit rest (%)
City	semantic	4.794	0.883	0.384	52.1	95.8
City	arbitrary	2.721	1.225	0.525	41.7	91.7
Harbor	semantic	10.754	2.910	3.187	18.8	85.4
Harbor	arbitrary	7.652	1.500	2.494	35.4	85.4

Table 62: **Post hoc Qwen3-32B audit of incorrectly offset answers.** The first three numeric columns are unclipped response MAE. Exact means all eight relative response coordinates are within .01 of truth; Omit rest means both zero-shock predictions are near zero. These last two columns concern the selected-pattern interface only. They do not measure correct absolute forecasts.

This matches distinct prompts and prompt presentations, not computation or optimizer updates: RL samples eight completions per prompt and uses about 2,400 updates. The higher-rate SFT recipe also changes learning rate. Both supervised outcomes are reported in Table 63; the seed-42 RL checkpoint is the primary diagnostic comparison. The other two matched RL seeds yield City/semantic structure accuracies 52.1% and 50.0%, with response MAE 1.029 and 1.053.

Domain	Labels	Recipe	Structure (%)	Resp. MAE	α	95% interval
City	semantic	Base	45.8	2.105	-0.211	[-0.467, 0.014]
City	semantic	RL	50.0	1.077	0.005	[-0.074, 0.078]
City	semantic	SFT 10^{-6}	58.3	1.253	0.166	[-0.235, 0.549]
City	semantic	SFT 2×10^{-5}	97.9	0.658	1.413	[1.205, 1.620]
City	arbitrary	Base	54.2	2.117	0.344	[0.162, 0.529]
City	arbitrary	RL	50.0	1.222	0.012	[-0.068, 0.102]
City	arbitrary	SFT 10^{-6}	43.8	1.727	0.001	[-0.424, 0.489]
City	arbitrary	SFT 2×10^{-5}	39.6	1.274	-0.211	[-0.516, 0.107]
Harbor	semantic	Base	43.8	4.967	0.170	[-0.157, 0.499]
Harbor	semantic	RL	50.0	3.445	0.180	[0.050, 0.325]
Harbor	semantic	SFT 10^{-6}	54.2	5.317	0.058	[-0.447, 0.538]
Harbor	semantic	SFT 2×10^{-5}	45.8	12.192	-0.932	[-1.528, -0.375]
Harbor	arbitrary	Base	50.0	5.353	-0.094	[-0.301, 0.109]
Harbor	arbitrary	RL	52.1	3.619	0.026	[-0.115, 0.190]
Harbor	arbitrary	SFT 10^{-6}	47.9	5.965	0.120	[-0.132, 0.381]
Harbor	arbitrary	SFT 2×10^{-5}	50.0	10.655	0.004	[-0.282, 0.255]

Table 63: **A supervised recipe learns in-domain structure selection but fails to transfer.** All trained recipes in this table use seed 42; the successful recipe’s three seed replications are in Table 64. Structure accuracy is derived from forecasts, not reference naming. All outputs parse. Intervals resample the 24 paired episodes within each cell (2,000 draws); they do not capture training-seed uncertainty. Both fixed SFT learning rates are retained.

At the higher rate, City/semantic structure accuracy is 47/48 (97.9%), response MAE is .658, and forecast-level MAE is .527. Calibration $\alpha = 1.413$ indicates overshoot despite accurate choice of response form. The lower-rate recipe reaches only 58.3% structure accuracy. Transferred from City/semantic training, higher-rate SFT does not recover structure selection under arbitrary labels or in Harbor, and its Harbor response errors increase substantially. Original-task SFT response scores are unchanged when clipping is removed, and an additive-offset audit does not remove the transfer failure. Reference naming remains perfect. Thus the successful recipe establishes in-domain learnability, while a semantic cue-to-pattern association remains a possible explanation. It does not establish transferable induction from reference trajectories or an isolated effect of SFT versus RL.

Seed replication of the successful recipe. The 2×10^{-5} result above was obtained at one seed. Before evaluating them, we froze three replications at seeds 43–45 that change the training seed and nothing else—identical data, optimizer, schedule, LoRA configuration and 300-update budget—and did not extend the 1×10^{-6} arm, which had not learned the task. All three replicate the in-distribution result at 100.0% City/semantic structure accuracy, against 97.9% at seed 42, and lower response MAE (.277–.314 versus .658). The transfer cells remain at chance at every seed: City/arbitrary spans 39.6–50.0%, Harbor/semantic 45.8–52.1% and Harbor/arbitrary 47.9–50.0%. All outputs parse and

reference naming is 100% in all sixteen cells of all four checkpoints. In-distribution learnability under this recipe is therefore seed-robust, and so is its failure to transfer; the caveat that remains is the comparison to RL, which still differs in update budget and learning rate.

Seed	City/sem.	City/arb.	Harbor/sem.	Harbor/arb.	Resp. MAE	Name (%)
42	97.9	39.6	45.8	50.0	0.658	100
43	100.0	50.0	50.0	50.0	0.277	100
44	100.0	50.0	50.0	50.0	0.301	100
45	100.0	45.8	52.1	47.9	0.314	100

Table 64: **The supervised in-distribution result is seed-robust, and so is its transfer failure.** The 2×10^{-5} recipe at four training seeds, evaluated on the unchanged 768-prompt component diagnostic. Only the seed differs. Response MAE and Name are the City/semantic cell.

The boundary is arbitrary labels, not surface domain. Transfer failure alone cannot say whether the untrained cells are unreachable by this recipe or only unreached from City/semantic training. We therefore trained the identical recipe natively in each of the three cells held out above—City/arbitrary, Harbor/semantic and Harbor/arbitrary—at three seeds each, changing only the training cell. Each model is evaluated on its own cell’s original task. Results are in Table 65.

The two conditions separate cleanly. Harbor/semantic reaches 100.0% structure accuracy at every seed, matching the 97.9% obtained in City/semantic, so the second surface domain is fully learnable by this recipe and the earlier Harbor failure was a transfer failure rather than a ceiling. Both arbitrary-label cells remain at chance under native supervision—41.7% in City and 53.5% in Harbor, with seed ranges of [37.5, 43.8] and [52.1, 54.2] that exclude the semantic result—at the same recipe, budget and seed count. Reference naming stays at 100.0% in all three cells, so the failure is not one of identifying the named reference. The boundary this recipe encounters is therefore the arbitrary label mapping specifically, and it survives native training; surface domain shift is not a barrier to it.

Domain	Labels	Structure (%)	Seed range	Resp. MAE	Name (%)
City	arbitrary	41.7	[37.5, 43.8]	1.101	100.0
Harbor	semantic	100.0	[100.0, 100.0]	0.902	100.0
Harbor	arbitrary	53.5	[52.1, 54.2]	1.932	100.0

Table 65: **Native training separates the label condition from the surface domain.** The higher-rate SFT recipe trained in each cell and evaluated on that cell’s original task, averaged over seeds 42–44. Seed range is the minimum and maximum structure accuracy across those seeds. Name is selection-only reference identification. Harbor is learnable under semantic labels; neither arbitrary-label cell exceeds chance.

Adaptive decision disclosure. The initially specified numerical launch gate was not met. Before any SFT training, a documented outcome-informed amendment used the stronger model’s conditional numerical ability and absolute-level errors to justify running the two already prepared, unchanged recipes. This is an exploratory learnability pilot, not a passed preregistered gate or a confirmatory method comparison. There was no SFT endpoint selection. Separate reference-identification supervision was unnecessary given the perfect selection-only results.

D.1.10 MECHANISM-FAMILY EVIDENCE USE UNDER COHERENT GAIN CHANGES

The mechanism-family accuracy comparison uses a fixed population-prior training control. It is distinct from the parent Coin City shuffled-pair control and does not preserve the matched arm’s exact target-vector distribution. To test whether the accuracy gain depends on displayed episode-specific relationships, we freeze a fresh gain-intervention set before inference: 16 pairs in each of 12 mechanisms, comprising 128 pairs from eight seen mechanisms and 64 from four held-out compositions or parameter-extrapolation mechanisms. Both mechanism-disclosure conditions use the same episodes, with $k = 9$ target observations.

Within each pair the target gain changes from the 20th to the 80th percentile of that world’s native uniform support. Inputs, exact noise draws, calibration trajectories, baseline, and mechanism are fixed; affected target observations, terminal states, and future outcomes are regenerated coherently. The disclosed description provides the same gain range on both sides, never the actual gain. This tests a gain change, not a change of mechanism identity. Native supports are retained, including the narrow, noisy extrapolation world.

The primary metric is response-change MAE: subtract the zero-shock forecast within each horizon, take the high-minus-low change of the resulting eight responses, and compare with the simulator’s change. This removes purely additive level tracking. Greedy inference evaluates Qwen3-4B-Instruct-2507 base and the final matched and shuffled-target adapters at all five training seeds, 42–46. All pairs parse. The mean contrasts and bootstrap intervals in Table 66 resample paired training seeds and world-stratified episodes (2,000 draws). The comparator is the distribution-matched shuffled-target arm of Appendix D.1.11, not the fixed population prior, so the contrast isolates episode-to-target pairing at a preserved answer distribution. Individual seed results are retained in the released analysis.

Disclosure	Mechanisms	Base	Matched	Shuffled	No change	Slope M	Slope S	Shuffled–matched	95% interval
Disclosed	Seen	3.405	1.250	1.512	1.591	0.323	0.084	+0.262	[+0.212, +0.319]
Disclosed	Held out	2.629	0.898	0.874	0.826	0.122	0.049	-0.024	[-0.079, +0.027]
Undisclosed	Seen	3.691	1.115	1.488	1.591	0.461	0.107	+0.373	[+0.318, +0.426]
Undisclosed	Held out	3.664	0.878	0.862	0.826	0.238	0.060	-0.016	[-0.061, +0.027]

Table 66: Better gain tracking on seen mechanisms does not establish transferred evidence use. Entries are response-change MAE (lower is better); trained entries average the five training seeds 42–46. No change is an invariant predictor. Slope M and Slope S are matched and shuffled tracking slopes (ideal 1, zero for no response to changed evidence). Positive shuffled-minus-matched contrasts favor matched training. Both held-out intervals include zero, and matched held-out errors exceed the no-change baseline.

Matched training improves on the shuffled-target control for seen mechanisms by .262 [.212, .319] (disclosed) and .373 [.318, .426] (undisclosed), but held-out contrasts are −.024 [−.079, .027] and −.016 [−.061, .027], both spanning zero. The tracking slopes separate the same way: on seen mechanisms matched training reaches .323 and .461 against .084 and .107 for the shuffle, whereas on held-out mechanisms matched tracking falls to .122 and .238 and its errors exceed the .826 no-change baseline in both conditions. At five seeds, therefore, the evidence-use benefit is established on seen mechanisms and absent on held-out ones. The previously reported ordinary forecast-accuracy advantage remains a narrower result than transferred evidence use: these interventions neither show a complete absence of sensitivity nor establish robust tracking of held-out episode-specific gains.

D.1.11 DISTRIBUTION-MATCHED SHUFFLED-TARGET CONTROL

The accuracy comparison above uses a fixed population-prior control, which does not preserve the matched arm’s exact target-vector distribution. We therefore ran a separately frozen shuffled-target control. It holds the prompts, the training budget and the exact joint target-vector multiset fixed within compatible world/evidence-depth groups, changing only which compatible episode receives which target vector, so the two arms differ in the episode-to-target pairing and in nothing else. Qwen3-4B-Instruct-2507 trains for 300 rounds in each of 5 seeds per disclosure condition. The matched arm is rerun on the same A5000 hardware and actor-memory allocation, so both arms share hardware; endpoints use a frozen fresh held-out set of 720 tasks per seed under greedy and five-draw stochastic decoding. The pre-registered primary estimand is the stochastic shuffled-minus-matched response MAE. All 40 endpoint/decode cells completed and every response parsed.

Matched training retains an advantage on the primary endpoint in both disclosure conditions. With the mechanism inferred it is .094 (5/5 seeds positive, Student-*t* [.041, .146], hierarchical [.032, .188], wild cluster $p = .0138$, exact sign test $p = .0312$), the only cell that clears every estimator. With the mechanism described the point estimate is larger, .155 [.046, .260], but rests on 4/5 seeds, one of which reverses sign, and does not clear the cluster-robust test ($p = .0794$). Greedy decoding is directionally consistent in both conditions—.120 described and .053 inferred—but neither cell clears the cluster-robust test.

Disclosure	Decode	Shuffled—matched	95% Student- <i>t</i>	95% hierarchical	Webb <i>p</i>	Seeds > 0
Disclosed	Greedy	+0.120	[-0.027, +0.267]	[+0.027, +0.210]	0.094	4/5
Disclosed	Stochastic	+0.155	[+0.005, +0.305]	[+0.046, +0.260]	0.079	4/5
Undisclosed	Greedy	+0.053	[-0.022, +0.128]	[+0.000, +0.099]	0.149	4/5
Undisclosed	Stochastic	+0.094	[+0.041, +0.146]	[+0.032, +0.188]	0.014	5/5

Table 67: **Matched training retains a smaller advantage over a distribution-matched shuffled-target control.** Entries are shuffled-minus-matched response MAE at 5 training seeds; positive favors matched training. The hierarchical interval resamples seeds and trajectories; Webb *p* is a two-sided wild cluster bootstrap. The stochastic inferred-mechanism cell is the pre-registered primary endpoint and the only cell clearing every estimator.

Disclosure	Decode	Shuffled—matched	95% Student- <i>t</i>	95% hierarchical	Webb <i>p</i>	Seeds > 0
Disclosed	Greedy	+0.048	[-0.014, +0.110]	[+0.011, +0.092]	0.088	4/5
Disclosed	Stochastic	+0.076	[-0.018, +0.170]	[+0.017, +0.173]	0.070	4/5
Undisclosed	Greedy	-0.019	[-0.057, +0.018]	[-0.046, +0.007]	0.206	1/5
Undisclosed	Stochastic	+0.018	[-0.093, +0.129]	[-0.040, +0.142]	0.881	1/5

Table 68: **The matched advantage does not appear in level error.** Shuffled-minus-matched level MAE, same runs and estimators as Table 67. The inferred-mechanism cells show no matched benefit. This contrast was not pre-registered and is descriptive.

The benefit is concentrated in response error, which is what the manipulation predicts. Because donors are drawn from the same compatibility group, the shuffle changes an episode’s target *level* almost not at all while changing its response *shape* almost entirely: averaged over all 4,800 episodes, the recipient-donor level difference is .0086–.0093 against a shape difference of .892–.908, a ratio of 96–104 across both disclosure conditions and all five seeds, with a within-group level standard deviation of .0076. Both arms therefore receive essentially correct level information and only the matched arm receives the correct episode-specific relationship, so a level-error advantage is not expected. Its absence when the mechanism is inferred (.018 [-.040, .142], $p = .8812$ stochastic; -.019 [-.046, .007] greedy) is a manipulation check rather than a negative result, and it argues against the benefit being generic improvement, which would raise both metrics.

We report this as supporting rather than decisive, for three reasons. The level contrast was not pre-registered. Level MAE is computed on the raw forecast vector, so it partly contains the response component and is roughly twice its magnitude, diluting any response-driven gain. And the pattern is not uniform: the described condition does show positive level contrasts (.048 greedy, .076 stochastic). The asymmetry above is a property of the frozen builder, not of any prediction; `measure_level_shape_asymmetry.py` recomputes it from the training targets alone and records it in `level_shape_asymmetry.json`. None of this speaks to held-out mechanisms, where gain tracking remains unestablished (App. D.1.10). These estimates use the same-hardware comparator; a supplementary historical-hardware comparison reusing the earlier A6000 matched seeds 45–46 is retained in the released results and supports no claim here.

Control status and reproducibility. Only completed diagnostics enter the tables and claims here. Their artifacts, manifests, raw outputs, and the outcome-informed SFT amendment reside under `exp3_training_transfer/overnight_diagnostics_20260921/`, in `structure/`, `learnability/`, `evidence_use/`, and `shuffled_control/`. `paper/generate_exp3_diagnostic_tables.py` renders the four diagnostic tables from completed summaries and records source hashes in `paper/tables/exp3_diagnostics_sources.json`; `paper/generate_exp3_shuffled_control_tables.py` renders the two control tables and records hashes, the design freeze and the analysis-code hash in `paper/tables/exp3_shuffled_control_sources.json`. Both refuse to emit a partial roster.

E EXPERIMENT 4: HISTORICAL-MARKET TRAINING AND HELD-OUT EVALUATION

Experiment 4 asks whether base-relative gains extend from the controlled setting to events and normalized question families unseen during training. It uses proper-score training, a point-in-time cutoff, and family-disjoint held-out data.

E.1 SAMPLE CONSTRUCTION AND LEAKAGE CONTROLS

Source archive and sample. The source export² has 1,788,703 rows and SHA-256 5697f767...f638e7d (the full digest is stored in the manifest). The builder scans the entire file and fails on checksum drift. The calendar splits contain 1,736 train, 512 development, and 1,024 test markets, with YES rates .311, .275, and .301. Train contains 991 distinct events and 1,456 normalized question templates; development contains 292 and 453; test contains 493 and 866. The selected archive rows have no official series identifiers, so event and template components supply the operative family groups.

Eligibility and point-in-time censoring. We admit literal Yes/No markets resolved by an extreme final token price. A boundary-aware domain rule selects political, governmental, macroeconomic, and geopolitical questions while explicitly rejecting culture, awards, sports, esports, and crypto categories. This rule is fixed before model training. One task is built 30 days before the earliest archived terminal timestamp. At least four distinct UTC dates must have genuine pre-cutoff YES candles, the latest candle must be no more than seven days stale, and its price must remain in [.02, .98]. We keep at most the latest 14 dates. Empty, carry-forward, and post-cutoff candles are rejected. The prompt omits the outcome, final prices and volume, market status, realized terminal time, and the tags used for domain selection.

The archive records timestamped price history but not revisions to static question or resolution-rules text. We therefore treat those strings as time-invariant. This is a residual point-in-time limitation: the audit verifies that no timestamped or outcome field crosses the cutoff, but cannot determine whether an archived rule changed between the forecasting cutoff and resolution.

Family-disjoint temporal split. Before applying sample caps, every development row sharing an event, available series, or normalized template with train is removed; the same rule removes test rows matching train or development. Numeric values and month names are normalized in template keys. The final intersection is zero for each pair of splits. Sampling within a split is a stable hash of task identifier and protocol version, so row order in the export cannot change membership.

E.2 TRAINING AND HELD-OUT EVALUATION

Matched update budget. Training comprises 1,736 unique markets arranged in a deterministic, balanced schedule of 4,800 presentations, matching Experiment 3’s 300-round update budget. Every market appears twice and 1,328 appear a third time. Schedule order and task files are checksum-verified. With batch size 16, this gives exactly 300 rollout-and-update rounds (2,400 AdamW steps). The manifests and tables distinguish unique markets from presentations. The longest Qwen3-4B wrapped prompt contains 1,012 tokens against the 1,792-token training limit. Data, schedule, reward, optimizer, update count, seeds, and scoring rule are held constant across models; each uses its native chat template. Shared optimizer, batching, and runtime details appear in Section F.

² 🗄️ [Full Polymarket archive](#). Link anonymized for review.

Table 69: Experiment 4 training configuration. Held-out data are absent from training and evaluated once per model after all three training seeds.

Quantity	Value
Model roster / adapter	Qwen3-4B-Instruct-2507, Qwen3-8B, and Llama-3.1-8B-Instruct; LoRA rank 32, $\alpha = 64$
Algorithm	Dr. GRPO; $\beta = 0$; AdamW, lr 10^{-6}
Unique train / presentations	1,736 / 4,800 balanced repetitions
Rounds / seeds	300; 42, 43, 44
Rollout	16 prompts, 8 samples each; temperature 1.3
Output	one JSON probability; 128 new-token limit
Reward	$1 - (\hat{p} - y)^2$ only
Chat formatting	Qwen3-4B uses its frozen ChatML wrapper; Qwen3-8B uses its non-thinking template; Llama uses its native instruct template.
Development	512 tasks; temperature-zero evaluation every 25 rounds
Held-out evaluation	1,024 tasks; temperature zero; one evaluation per model

E.3 HELD-OUT RESULTS

Decision rule. The primary learned contrast is the mean loss over three trained seeds minus base loss. Invalid JSON receives Brier loss one. Intervals use 5,000 bootstrap resamples of connected event/template components.

All Qwen3-8B test forecasts parse. Its three-seed trained mean obtains Brier .11642 versus .12867 for the base, a difference of $-.01225$ with 95% interval $[-.01948, -.00602]$. Every training seed improves on the base. The interval excludes zero, supporting transfer to held-out events for Qwen3-8B.

Table 70: **Qwen3-8B on the 1,024-market held-out set.** The mean-row interval resamples connected event/template families; lower is better.

Forecaster	Brier $\downarrow$	Δ base	Log loss $\downarrow$
Untrained Qwen3-8B	.1287	—	.5003
Trained, seed 42	.1164	-.0123	.4058
Trained, seed 43	.1168	-.0119	.4069
Trained, seed 44	.1161	-.0126	.4146
Trained mean	.1164	-.0123 $[-.01948, -.00602]$	.4091

Table 71: **Qwen3-4B and Llama held-out comparisons.** The Qwen3-8B result is immediately above. Δ columns are trained mean minus comparator, so negative values favor training. Parse is the minimum coverage across base and trained endpoints; Llama’s 12 unparsed base outputs receive Brier loss one.

Model	Base	Seed 42	Seed 43	Seed 44	Trained mean	Δ base [95% CI]	Δ market [95% CI]	Δ Platt [95% CI]	Parse
Qwen3-4B	.1316	.1142	.1140	.1194	.1159	-.0157 $[-.02528, -.00667]$	.0022 $[-.00264, .00721]$	.0036 $[-.00125, .00874]$	1.000
Llama-3.1-8B	.1520	.1151	.1149	.1151	.1150	-.0370 $[-.05376, -.02258]$	.0013 $[-.00004, .00371]$	.0027 $[-.00067, .00539]$	.988

Qwen3-4B replicates the base-relative improvement: its trained mean is .11589 versus .13160 for base, a difference of $-.01572$ $[-.02528, -.00667]$. Llama-8B improves from .15203 to .11502, a difference of $-.03701$ $[-.05376, -.02258]$; all three trained Llama endpoints parse.

The latest pre-cutoff market price obtains Brier .11369, and logistic recalibration of its logit fit only on the training split obtains .11232. Neither comparator informs language-model training. The trained-minus-market interval spans zero for every model. Qwen3-4B is also statistically indistinguishable from the calibrated price, but Qwen3-8B and Llama-8B have higher error by $+.00409$ $[+.00098, +.00785]$ and $+.00269$ $[+.00067, +.00539]$, respectively. Thus the complete roster shows base-relative improvement rather than superiority to market-based forecasts.

E.4 SCALE AND ARCHITECTURE COMPARISON

The scale comparison uses all 318 eligible 2026-window tasks whose identifiers and event, series, and normalized question-template families are absent from the training, development, and 1,024-market evaluation sets. The holdout contains 227 event families and 288 normalized templates; selection does not use outcomes or market probabilities.

Qwen3 checkpoints at 4B, 8B, 14B, and 32B use three seeds, a common reward, a 4,800-presentation schedule, and the same development rule. The prospective 32B extension was registered before the shared holdout was opened; it retained the 14B scientific settings and changed only memory distribution across four A6000 GPUs. Each untrained base and trained adapter receives five non-thinking draws per task at temperature .7, top- p .8, and top- k 20. All 25,440 Qwen draws satisfy the strict JSON contract.

Table 72: **Exact values behind the main-text scale figure.** Lower Brier is better. Trained means average seeds 42–44. Missing five-draw task groups receive Brier loss one.

Checkpoint	Base	Trained mean	Trained — base	Trained parse
Qwen3-1.7B	0.2253	0.1248	-0.1005	100.0%
Qwen3-4B-Instruct-2507	0.1522	0.1366	-0.0156	100.0%
Qwen3-8B	0.1369	0.1286	-0.0083	100.0%
Qwen3-14B	0.1281	0.1270	-0.0012	100.0%
Qwen3-32B	0.1266	0.1265	-0.0001	100.0%
Llama-3.2-3B	0.2189	0.1367	-0.0822	100.0%
Llama-3.1-8B	0.1642	0.1264	-0.0377	100.0%

The trained means are .13664, .12859, .12695, and .12652 at 4B, 8B, 14B, and 32B; the corresponding bases are .15224, .13686, .12812, and .12659. The 32B seed Briers are .12653, .12651, and .12651. Trained-minus-base contracts from $-.01559$ at 4B to $-.00827$ at 8B, $-.00117$ at 14B, and $-.00007$ [$-.00045, +.00029$] at 32B. Thus the larger untrained checkpoints approach the same absolute performance reached by the trained endpoints under the shared recipe. The registered 14B-minus-8B trained contrast remains $-.00164$ [$-.00332, -.00024$].

On this holdout, contemporaneous Polymarket prices have Brier .12655. Trained 14B minus market is $+.00040$ [$-.00028, +.00139$]; trained 32B minus market is $-.00003$ [$-.00009, +.00002$]. Both intervals cross zero. The scale profile therefore shows convergence to the contemporaneous market frontier, not a resolved advantage over it. Because data volume, update count, and optimizer settings are held fixed, the comparison does not identify separately tuned optima for each model size.

Llama-3.1-8B uses the same holdout and five-draw endpoint. All 4,770 trained draws parse; nine of the 1,590 base draws do not, causing seven base task groups to receive Brier loss one. Trained seed Briers are .12638, .12629, and .12662, for a mean of .12643 versus .16417 for base. Trained minus base is $-.03774$ [$-.05860, -.01946$]; trained minus Polymarket is $-.00012$ [$-.00046, +.00012$]. The latter interval crosses zero.




Figure 16: **Llama scale replication.** At 3B and 8B, open/filled squares show base/trained means, small squares show seeds, and the dashed line shows Polymarket. The parameter axis is logarithmic. The 3B base is descriptive because its strict-JSON recovery was specified after observing syntax failures; all other endpoints follow the common protocol.

Qwen3-1.7B and Llama-3.2-3B preserve the same training/development data, three seeds, 4,800-presentation schedule, five stochastic draws, and holdout. Qwen3-1.7B seed Briers are .12413, .12610, and .12415 (mean .12479) versus .22531 for base; trained minus Polymarket is $-.00175$ [$-.00644, +.00265$]. Llama-3.2-3B seed Briers are .13197, .13666, and .14149 (mean .13671); trained minus Polymarket is $+.01016$ [$+.00121, +.02060$]. The pinned public BF16 Llama mirror is recorded because the gated upstream repository was unavailable to the workspace account. Its unconstrained base initially produced only 4/318 complete five-draw groups at the 128-token cap; an output-cap-only rerun reached 21/318. A disclosed, holdout-informed syntax recovery therefore repeated the full base endpoint with identical prompts and sampling settings while constraining output to a JSON object containing one numeric probability. All 1,590 recovered draws parse. Base Brier is .21887, and trained minus base is $-.08216$ [$-.11062, -.05367$]. Because the syntax constraint was introduced after observing the failures and differs from the trained endpoints, this base-relative contrast is descriptive. Both trained models also parse every draw.

Hosted frontier references. The API panel evaluates the same 318-task holdout. It is categorical rather than a parameter-scale curve: provider parameter counts are unavailable or not comparable, and the APIs do not expose identical decoding controls. DeepSeek-V4-Pro uses the common temperature .7/top- p .8 stochastic request. GPT-5.6-Sol uses provider-native low reasoning without user-set temperature or top- p ; Claude Opus 4.8 likewise uses provider-native sampling. Each endpoint receives five independent calls per task and uses the same strict parser, five-draw mean, Brier scoring, and family-clustered 5,000-repetition interval procedure as the local panel.

DeepSeek parses all 1,590 draws and has Brier .12660, differing from the crowd by $+.00005$ [$-.00105, +.00102$]. Claude parses all draws and has Brier .12927, a difference of $+.00272$ [$-.00049, +.00614$]. GPT receives all 1,590 responses but six are malformed; two task groups therefore receive Brier loss one, and task-level coverage is 99.37%. Its Brier is .13335, a difference of $+.00680$ [$+.00016, +.01635$].

Kimi uses a 1,024-token output budget. Thirteen transport failures are retried without replacing the 1,577 successful responses; all 1,590 final draws parse. Its Brier is .12682, differing from the crowd by $+.00027$ [$-.00046, +.00109$]. Hosted results provide performance context, not evidence for a common scale law or an identical-decoder comparison.

E.5 QWEN3-4B EVIDENCE-UPDATING EVALUATION

This evaluation applies the Experiment 1 update prompt and all nine fixed packets for each of its 100 markets (three pro- H_1 , three anti- H_1 , and three orthogonal). The packet and initial-context SHA-256 digests are 6741be0a...8796d36 and c550cf05...0c098. Every condition receives the same canonical first structured Qwen 2.5-7B forecast for each market, so this paired analysis isolates conditional revision rather than web retrieval or initial event-model construction. Search and thinking

are disabled; decoding is greedy with a 1,500-token output cap. The longest prompt is 1,480 tokens against a 6,144-token context limit. These contrasts characterize revision behavior and are separate from the held-out forecasting evaluation.

The paired evaluation produces all 900 packet-conditioned revisions for the base and each final adapter, with 100% parse coverage and ICS equal to one. Base EHC/HFC is .909 on 591 eligible directional updates. Seeds 42, 43, and 44 obtain .917, .922, and .918 (592, 588, and 591 eligible updates), for a three-seed mean of .919. The paired market-bootstrap trained-minus-base EHC difference is +.0101 with 95% interval $[+.0004, +.0207]$; HFC has the same estimate and interval $[+.0009, +.0210]$. Training reduces mean absolute directional revision by .52 percentage points $[-.70, -.36]$ and orthogonal revision by .22 points $[-.38, -.06]$; the directional-minus-orthogonal gap narrows by .30 points $[-.52, -.09]$. The sensitivity ratio changes from $4.64\times$ to $4.83\times$ $(+.19 [-.07, +.48])$. These components show improved directional consistency alongside smaller revisions overall; they do not indicate an increase in absolute packet selectivity.

Table 73: **Paired evidence-updating contrasts.** Movement rows are mean absolute probability revisions in percentage points; intervals resample the 100 markets.

Measure	Trained — base	95% interval
Directional consistency (EHC)	+.0101	$[+.0004, +.0207]$
Headline-forecast consistency (HFC)	+.0101	$[+.0009, +.0210]$
Directional movement	-.52 pp	$[-.70, -.36]$
Orthogonal movement	-.22 pp	$[-.38, -.06]$
Directional — orthogonal movement	-.30 pp	$[-.52, -.09]$
Sensitivity ratio	$+.19\times$	$[-.07, +.48]$

F EXPERIMENTS 3–4: SHARED TRAINING CONFIGURATION AND RUNTIME

Tables 74–76 record the effective OAT/Dr. GRPO (Liu et al., 2024) arguments for the reported Coin City and Polymarket studies, including framework defaults omitted from the command line. Execution manifests hash the data, trainer, launcher, evaluator, prompt renderer, and each cell’s environment.

Step accounting. A “step” in the figures and earlier tables is one rollout-and-update round rather than one AdamW call. Each full run consumes 16 prompts per round and samples eight responses per prompt (128 trajectories). With one learner GPU, per-device microbatch one, one PPO epoch, and gradient accumulation 16, OAT performs eight AdamW parameter updates per round. Thus 4,800 prompt presentations give 300 reported rounds and 2,400 AdamW updates. The command line supplied a .03 warmup ratio, but OAT selected Transformers’ `constant` scheduler, which does not use the warmup argument. The effective learning rate was therefore 10^{-6} from the first optimizer update. We use “round” below to avoid counting the eight within-round parameter steps as independent data exposures.

Table 74: **Shared Dr. GRPO implementation parameters.** These values apply to the Coin City and Polymarket studies; the study-specific tables below record exceptions.

Quantity	Effective value
Trainable parameters	Frozen base model; LoRA rank 32, $\alpha = 64$, dropout 0, no bias; targets q.proj, k.proj, v.proj, o.proj, down.proj, up.proj, and gate.proj; no quantization.
Algorithm	OAT Dr. GRPO (critic.type=drgrpo); eight-response group advantage $r_i - \bar{r}$ with no group-standard-deviation division; policy loss is token-summed with the generation limit as a fixed normalizer (no response-length normalization).
Policy objective	PPO ratio clipping 0.2; truncated importance-sampling cap 2.0; one PPO epoch; KL coefficient $\beta = 0$; no critic/value model; reward whitening, advantage whitening, and REINFORCE mode all off.
Optimizer	PyTorch AdamW (Loshchilov & Hutter, 2019) (substituted for DeepSpeed fused/CPU Adam); learning rate 10^{-6} ; $\beta_1 = .9$, $\beta_2 = .95$; weight decay 0; maximum gradient norm 1.0.
Schedule	Constant 10^{-6} learning rate from the first optimizer update; the supplied .03 warmup ratio is ignored by Transformers' constant scheduler; one prompt epoch; 4,800 prompt presentations = 300 rollout-and-update rounds = 2,400 AdamW steps.
Batching	16 prompts per round; 8 sampled responses per prompt; 128 trajectories; training microbatch 1; gradient accumulation 16; one learner GPU.
Training sampling	Temperature 1.3; top- $p = 1$; top- $k = -1$; one independently sampled completion at a time, repeated eight times per prompt.
Truncation and stopping	Model EOS only (no external stop string); a completion ending at the length limit receives reward zero. Missing or invalid required JSON also receives reward zero.
Precision and memory	bfloat16; FlashAttention 2; gradient checkpointing; DeepSpeed ZeRO-2 (Rajbhandari et al., 2020); reference-model offload; vLLM prefix caching (Kwon et al., 2023); dropout disabled.
Seeds and shuffling	Training seeds 42, 43, and 44. Each dataset is deterministically shuffled once, then OAT shuffles prompt batches independently under the training seed.
Checkpointing and logging	Evaluation includes step zero and the final round; intermediate evaluation at the setting-specific frequency; one final LoRA adapter retained; metrics logged each round; Weights & Biases disabled.

F.1 COIN CITY A-TO-B FACTORIAL

Table 75 records the original three-seed configuration for Experiment 3. The current Qwen3-8B and Qwen3-4B analyses extend both reward arms to eight seeds (Section D.1.6), retaining the five-draw stochastic round-300 endpoint. The structureless diagnostic and Llama-8B architecture comparator retain their original three seeds. The table supplements the shared optimizer details in Table 74; where the two tables differ, this setting-specific table governs. The execution manifest stores the environment; the protocol archive hashes the preflight manifest, data files, generator, prompt renderer, scorer, and trainer entry point.

Table 75: **Original three-seed configuration for Coin City hidden-structure transfer.** The later Qwen extensions are documented in Section D.1.6.

Quantity	Effective value
Model roster	Qwen/Qwen3-4B-Instruct-2507, Qwen/Qwen3-8B, and meta-llama/llama-3.1-8B-Instruct.
Chat formatting	Qwen3-4B uses the explicit ChatML wrapper; Qwen3-8B and Llama-3.1-8B use their tokenizer templates, with Qwen thinking disabled. All include an assistant-generation prefix.
Original grid	3 models $\times$ 2 reward arms (episode-matched reward, stored under the legacy artifact key <code>causal</code> , and <code>population.prior</code>) $\times$ 3 seeds (42, 43, 44) = 18 full training runs.
Diagnostics	Qwen3-4B <code>structureless</code> at the same three seeds, plus one untrained endpoint for each model. These six diagnostics are excluded from the primary episode-matched-versus-prior estimand.
Training data	Coin City and Structure A only; 4,800 prompts, balanced across $k \in \{0, 2, 4, 8\}$ and weak/strong target match. Two reference trajectories of length 14 and one target trajectory of maximum length 8.
Comparator target construction	Episode-matched reward uses the episode truth. Population-prior reward targets are a fixed no-self-match derangement within each k ; its target-vector multiset is exactly identical to episode-matched training. Prompt strings are byte-identical across these two arms.
Held-out data	2 domains $\times$ 2 target structures $\times$ 3 hint arms $\times$ 4 evidence depths $\times$ 30 episode seeds = 1,440 rows. The complete held-out dataset is byte-identical across all training arms.
Mechanism ranges	Evaluation $g \in [.22, .92]$; Structure B $\rho \in [.48, .74]$, $\phi \in [.18, .42]$, and $\lambda \in [.68, .94]$. Training reference gains are $[\cdot18, \cdot46]$ and $[\cdot68, \cdot98]$.
Domain ranges	Coin City: baseline 50, support $[0, 100]$, noise 1.6, driver pool $\{-12, -8, -4, 0, 4, 8, 12\}$. Coin Harbor: baseline 180, support $[40, 320]$, noise 4.5, driver pool $\{-24, -16, -8, 0, 8, 16, 24\}$.
Interventions	Five shocks at horizons one and three; Coin City $\{-12, -6, 0, 6, 12\}$, Coin Harbor $\{-24, -12, 0, 12, 24\}$. Ten joint forecasts produce eight nonzero-minus-zero response contrasts.
Reward	$\cdot4 r_{\text{level}} + \cdot6 r_{\text{response}}$. Each component averages $\max(0, 1 - e /S)$; level tolerance is 10% of the domain’s observable range (10 Coin City points; 28 Coin Harbor units), and response tolerance is half one shock unit (3 and 6, respectively). Invalid or length-truncated output receives zero online reward.
Output contract	Exactly one JSON object with sole key <code>forecasts</code> and exactly ten finite JSON numbers. The same strict parser is used online and offline; no prose, extra/duplicate key, Boolean, non-finite value, or wrong-length list is accepted.
Structured decoding	Syntax-only XGrammar GBNF in training, online evaluation, endpoint evaluation, and base evaluation. It fixes object shape and arity only; numeric values and their relationship to targets remain unconstrained.
Optimization	Frozen base plus LoRA rank 32, $\alpha = 64$; OAT Dr. GRPO; PyTorch AdamW learning rate 10^{-6} with an effective constant schedule and no warmup; $\beta = 0$; one PPO epoch; gradient norm 1; ZeRO-2 and reference offload.
Batch and update accounting	16 prompts per round; eight rollouts per prompt; 128 trajectories; microbatch 1; gradient accumulation 16; 4,800 presentations = 300 rollout-and-update rounds = 2,400 AdamW steps.
Training generation	Temperature 1.3, top- $p = 1$, top- $k = -1$, model EOS only; 192 new-token limit; prompt limit 2,304; total model context 3,072; vLLM prefix caching enabled.
Precision / memory	BF16, FlashAttention 2, gradient checkpointing, dropout off, no quantization; vLLM GPU fraction .38; expandable-segments CUDA allocator.
Online evaluation	All 1,440 rows at temperature 0, one draw, top- $p = 1$, 192-token limit, batch 120, at step zero and every 50 rounds including round 300.
Reported endpoint evaluation	Five draws per prompt at temperature .7 and top- $p = .95$; trained decode seed is training seed +20,260,818 and the base seed is 20,260,819. Temperature-zero checkpoints are retained only as online monitoring artifacts and are not reported as endpoints.
Offline scoring	Level MAE, response MAE, response error normalized by the blind 50/50 pathway mixture, response correlation, strict parse coverage, and nearest behavioral pathway. A parse failure receives the full domain range as both level and response error.
Variance and estimands	Five-draw outputs first average within episode. Original hierarchical intervals resample three training seeds, then paired episodes; the precision extension resamples eight seeds and adds the small-cluster checks in Section D.1.6. Primary: prior-minus-episode-matched response MAE for correct-hint Structure B in Coin Harbor, equal over k . Domain, structure, joint, hint, and hint-by- k contrasts form the factorial analysis.
Preflight validation	A 10-round run for every model and evaluated arm must save an adapter, produce the full temperature-zero monitoring and five-draw stochastic evaluation sets, achieve at least .95 strict parse rate under each decoder, and exhibit nonconstant finite reward. The 17-check data preflight must also pass before estimating the factorial.
Compute	Training and trained-checkpoint evaluation use two NVIDIA A6000 48-GB GPUs (one actor, one learner), 8 CPU cores, and 56–100 GB allocated host memory. Untrained base evaluation uses one A6000, 8 CPU cores, and 24 GB host memory. All runs load models from the same pinned local cache.

Table 76: **Runtime environment for the reported training runs.** Package versions come from the job execution environment.

Component	Version / configuration
Python / CUDA build	Python 3.10.20; PyTorch 2.6.0 built for CUDA 12.4.
Training stack	OAT 0.2.4 at Git revision 869706620a35ec5304fa80a0c74f0c566cc2bc57; DeepSpeed 0.16.8; PEFT 0.15.2.
Generation stack	Transformers 4.51.3; vLLM 0.8.4; FlashAttention 2.7.4.post1.
Data stack	Datasets 2.16.1; Accelerate 1.14.0.
Allocator / logging	PYTORCH_CUDA_ALLOC_CONF=expandable_segments:True; Weights & Biases disabled; stdout/stderr and per-round metrics retained with each run.

G EXPERIMENTS 3–4: REPRODUCIBILITY AND ARTIFACT MAP

The artifact roots `exp3_training_transfer/coin_city_structural/` and `exp3_training_transfer/polymarket/scripts/` contain the generators, prompt builders, trainers, evaluators, and analysis code for the Coin City and historical-market studies. Execution ledgers record effective hyperparameters, Git revisions, adapter identities, and SHA-256 hashes of data and executable protocol files; result renderers verify these fields before aggregation. For Coin City, `make_qwen3_8b_endpoint_outputs.py` validates the seven original five-draw files (50,400 scored draws). The later eight-seed reports in Section D.1.6 supply Figure 9; `render_exp3_paper_macros.py` renders the current prose values from those reports and the recorded structure-selection and mechanism-family endpoints. `make_paper_outputs.py` validates the original 24-endpoint roster and all 48 recorded endpoint files before writing the machine-readable result and the stochastic appendix comparator table. Temperature-zero files are checked as execution provenance but are excluded from manuscript estimates. For the historical study, `analyze_exp1_reapplication.py` validates four matched 900-record files and computes the paired-market direction, sensitivity, and movement-component contrasts. `generate_exp4_scale_figure.py` verifies the frozen 32B protocol, summary, and raw-output hashes before adding the point to Figure 10 and Table 72; it also writes the separate Llama replication in Figure 16. The Polymarket manifests preserve sample membership, prompt and adapter hashes, parse coverage, and bootstrap settings.